

LAKSH'S DAILY BRIEF

Laksh's Morning Brief — August 25, 2026

Tuesday, August 25, 2026

World · Markets & Money · AI & Infrastructure · Models & Research



BBC World — Two unvaccinated people die of measles in Pennsylvania



Al Jazeera — Trump's latest wave of Iran sanctions: Which 60 entities are targeted?



Al Jazeera — US military transport plane arrives in Moscow in rare occurrence

1. World & Geopolitics

US shifts from bombs to bank accounts in Iran — oil drops 3% as war risk premium evaporates

President Trump unleashed a new wave of sanctions on 60 Iranian entities, including banks, oil traders, and cyber operators, in what he calls an "economic war" to replace the shooting war that began in February ([Al Jazeera](#)). The administration is also calling on allies to join the isolation effort ([NPR](#)). Iran's president responded by saying a US memorandum is the "best path out" of the stalled war ([PBS](#)), though analysts warn the regime is in crisis ([Iran International](#)). The key market reaction: **Brent crude fell 3%** to around \$72, as investors concluded the shift from airstrikes to economic pressure reduces the risk of a Strait of Hormuz disruption ([CNBC](#)). Background: The US-Iran war began in February 2026 after a series of escalations. A ceasefire was reached in April, but it has been fragile. Now the US is trying to strangle Iran's economy instead of bombing it. The Strait of Hormuz — a chokepoint : a narrow passage through which about 20% of the world's oil passes — was the main fear during the war. If Iran had blocked it, oil could have spiked to \$150+. That fear is now receding.

Canada-US trade war: talks collapse, Ottawa prepares retaliation

Canadian negotiators left Washington last week without a deal to stop Trump's new tariffs, and Canada is now preparing to retaliate with its own tariffs on US goods ([CNBC](#)). The Canadian dollar fell to multi-year lows as investors priced in a full-blown trade war ([CNBC](#)). The US is demanding Canada drop its supply management system for dairy and accept a "Buy American" clause for government procurement. Canada says those demands are too severe. The tariffs — 25% on Canadian steel and aluminum, 10% on autos — are already in effect, and Trump has threatened to raise them to 50% on cars by January 2027. This is a serious escalation from the previous brief where talks were still "alive."

Ukraine proposes ceasefire plan with Donbas economic zone; Russia rejects

Ukraine has drafted a one-page ceasefire proposal with US and European officials, including a "Donbas Free Economic Zone" that would give the region special economic status while keeping it under Ukrainian sovereignty ([SBS Korea](#)). Russia rejected it, saying "Donbas is Russian territory." President Zelenskyy said he sees a window for US mediation until summer 2027 ([Reuters](#)). More than 50 countries have urged Russia to accept an immediate ceasefire ([Mezha](#)). Meanwhile, Russian drone operators are being accused of deliberately targeting Ukrainian civilians in "human safaris" — a war crime according to a UN commission ([France 24](#)).

Taiwan charges 9 people — including Nvidia and Super Micro staff — for illegally exporting AI servers to China

Taiwanese prosecutors charged nine people, including one Nvidia employee and two from Super Micro, for smuggling "high-end AI servers" to mainland China ([SecurityWeek](#)). This is a major enforcement action in the US-China tech war. The servers likely contained advanced Nvidia GPUs that are banned for export to China. The charges show that the "Taiwan loophole" — where Chinese companies used intermediaries to buy banned hardware — is being closed.

China slows exports of critical materials to Taiwan

China has strategically slowed exports of germanium and quartz to Taiwan, threatening the supply chain for semiconductors, optics, and robotics ([Tom's Hardware](#)). Germanium is used in infrared optics and fiber optics; quartz is essential for silicon wafer manufacturing. This is a quiet but potent weapon in the tech war — China controls about 60% of global germanium production and 80% of quartz.

Other world news: - **US Supreme Court** sided with Trump on mail-in voting restrictions, allowing states to require stricter ID for absentee ballots ([Time](#), [Guardian](#)). Midterm elections are in November. - **Colombia's new far-right president** Abelardo de la Espriella announced an immigration crackdown, ordering deportations of illegal immigrants ([Guardian](#)). He is a Trump ally. - **Kiribati detained a Chinese fishing vessel** for illegal shark finning in the Pacific ([Firstpost](#)). This is a test of sovereignty in the Pacific, where China has been building influence. - **A US military C-17 transport plane landed in Moscow** — a rare occurrence, likely for a prisoner exchange or diplomatic backchannel ([Al Jazeera](#)). - **Measles outbreak in the US** is at its highest since the virus was eliminated; two unvaccinated people died in Pennsylvania ([BBC](#)). - **A tornado in southern France** injured dozens and left 1,400 homes without power ([France 24](#)). - **Haiti** saw at least 47 killed in a gang attack near Port-au-Prince ([France 24](#)). - **Bangladesh** is considering joining the Mecca pact — a Saudi-led security alliance — testing its foreign policy balance ([KELO-AM](#)). - **Trump plans the largest mass visa cancellation in history**, potentially affecting up to 200,000 people, mostly asylum seekers ([Al Jazeera](#)).

2. Markets, Money & Deals

Stocks open mixed; oil drop boosts airlines, hurts energy

US futures were flat to slightly higher as the oil selloff (Brent -3%) eased inflation fears. The S&P 500 and Nasdaq were near flat. The "debasement trade" — buying gold, bitcoin, and other hard assets — is back in fashion as investors worry about US government debt and Treasury Secretary Bessent's bond-buying plans ([CNBC](#)). Gold had a massive options trade: someone took profit on a \$200 million position in the GLD ETF just 20 minutes after the open ([CNBC](#)).

Dick's Sporting Goods crashes 25% — worst day ever

The retailer missed earnings badly, citing a "challenging" footwear market where customers won't buy sneakers without deep discounts ([CNBC](#), [MarketWatch](#)). This is a red flag for consumer spending. If even Dick's — a discount-priced retailer — can't sell sneakers, it suggests the US consumer is pulling back. The company slashed its full-year outlook.

Citigroup's \$30 billion buyback sparks takeover speculation

Citigroup announced a \$30 billion share repurchase (buying back its own shares), which is a huge amount — equal to 60-129% of the market value of some regional banks ([ts2.tech](#)). This reignited speculation that Citi might acquire a regional bank, as US bank merger reviews become less strict. No deal has been announced.

McKesson buys Precision Medicine Group for \$2.25 billion

The healthcare distributor is acquiring a company that specializes in personalized medicine — tailoring treatments based on a patient's genetics ([Yahoo Finance](#)). This is a bet on the growing field of personalized medicine, which uses genetic data to prescribe drugs.

Paramount-Warner merger hits another roadblock

California's attorney general canceled a settlement meeting with Paramount over the \$110 billion merger with Warner Bros. Discovery ([WION](#)). California is suing to block the merger on antitrust grounds. The deal is already controversial because it would concentrate too much media power.

European stocks rise as oil drops; Siemens Energy up

European markets advanced, led by energy companies that benefit from lower oil prices (since they buy fuel) ([Bloomberg](#)). Siemens Energy was a standout.

Jackson Hole Symposium this week — all eyes on Fed Chair Warsh

The annual central bank conference in Jackson Hole, Wyoming, starts Friday. Fed Chair Kevin Warsh will speak. The market is pricing in a 1% chance of a rate cut in September ([ET](#), [Intellectia](#)). The Fed is expected to hold rates steady at 5.5-5.75%. The key question: will Warsh signal any easing if the economy slows?

Bessent considers tapping \$1 trillion Treasury account to buy back bonds

Treasury Secretary Bessent is reportedly considering using the Treasury General Account — the government's cash account at the Fed — to buy back long-term bonds ([CNBC](#)). This would be a massive intervention to lower long-term interest rates on government debt. It's a controversial move that could be seen as the Treasury "monetizing" debt — printing money to buy bonds, which is inflationary.

Xpeng shares sink on weak delivery forecast

Chinese EV maker Xpeng fell after giving weak delivery guidance, even as its robotics business was valued at \$6.3 billion — nearly matching the core EV business ([CNBC](#)). The robotics spin-off is getting a huge valuation, but the main business is struggling.

Crypto: XRP is up on news that Ripple's RLUSD credit fund is attracting institutional investors ([MEXC](#)). South Korea is speeding up a Digital Asset Framework Act to regulate crypto ([Crypto News](#)). World Liberty Financial's CEO says the success of its stablecoin proves the Trump family is not just profiting from cronyism ([CNBC](#)).

3. AI & Infrastructure

Nvidia earnings this Wednesday — the most important event of the week

Nvidia reports earnings after the bell on Wednesday. The stock is up 1.3% in pre-market. The key numbers to watch: revenue (expected ~\$45B?), data center revenue, and guidance. The big story is Nvidia's **dependence on hyperscalers** (large cloud providers like Microsoft, Amazon, Google) — they account for about 50% of Nvidia's data center revenue ([CNBC](#)). Nvidia is trying to expand its customer base by offering financing options. There's also a report that Nvidia is **raising prices by 15%** on its H100/H200 GPUs ([BigGo Finance](#)). This could boost revenue but risk alienating customers. The stock is down 5.5% over the past week — investors are nervous.

AMI orders 9,000 Nvidia Rubin GPUs, plans \$8B capex for 200 MW data center

AMI (Applied Materials? Actually AMI is a neocloud provider) has ordered 9,000 of Nvidia's next-generation Rubin GPUs and plans \$8 billion in capital expenditure (capex) for a 200 megawatt (MW) data center ([Business Standard](#)). This is a massive bet on the next generation of AI hardware. Rubin is Nvidia's successor to Blackwell, expected in 2027. The fact that a neocloud is ordering this early suggests strong demand for future AI compute.

Apple launches M5 Ultra and M6 chips — designed for local AI inference

Apple announced new chips: the M6 (for Mac Mini) and the M5 Ultra (for Mac Studio), along with updated Mac Mini and Mac Studio models ([Apple](#), [Tom's Hardware](#)). The M5 Ultra supports up to 512GB of unified memory, allowing developers to run large AI models locally. The Mac Mini with M6 starts at \$899 with 16GB RAM and 256GB storage

— a \$100 price hike from the previous model. The M5 Ultra Mac Studio is the most powerful Mac ever. Apple is explicitly positioning these for AI development: "With these frameworks and new chips, developers can run and fine-tune large AI models locally on their Mac" ([CNBC](#)). This is a direct challenge to Nvidia's GPU dominance for AI inference — Apple's unified memory architecture allows running models that would otherwise need expensive GPUs.

OpenAI's Jalapeño chip: built for fast inference, beats competition on benchmarks

OpenAI unveiled its own AI chip, Jalapeño, designed specifically for inference (processing new data, not training). On the InferenceX benchmark, Jalapeño achieved more tokens per user and more throughput per kilowatt than current state-of-the-art chips ([TechCrunch](#), [The Verge](#)). OpenAI says it offers "lower latency and higher throughput" — meaning faster responses and more users served simultaneously. This is a major step in the AI chip race: OpenAI is now both a model developer and a chip designer, reducing its dependence on Nvidia. The chip is likely made on TSMC's 3nm process.

Thomson Reuters launches its own frontier model

Thomson Reuters — the legal and news information company — has launched its own frontier AI model, trained on its proprietary legal, tax, and news data ([Thomson Reuters](#)). This is part of the "Sovereign AI" trend: companies with valuable proprietary data are building their own models rather than relying on OpenAI or Google. Thomson Reuters has a huge dataset of legal documents, case law, and regulatory filings that no other company has.

Google Cloud partners with Verizon to scale enterprise AI

Google Cloud announced a strategic partnership with Verizon to bring AI to enterprises, leveraging Verizon's 5G network and edge computing ([HPCwire](#)). This is about making AI inference fast and local — using Verizon's network to reduce latency.

Taiwan export charges and China material slowdown — supply chain risk

The illegal AI server exports to China (section 1) and China's slowdown of germanium/quartz exports to Taiwan are both threats to the AI supply chain. Taiwan makes over 90% of the world's advanced chips. If China restricts materials needed for chipmaking, it could slow production. The GaAs quantum-dot laser news (source 7) is a separate but related story: China's export controls on indium phosphide (InP) are pushing companies to develop alternatives like gallium arsenide (GaAs) lasers for AI data centers.

Data center power: UC San Diego to test innovative power tech

UC San Diego is testing a new power technology for AI data centers that could reduce the strain on the grid ([HPCwire](#)). This is a response to the "power gap" — there's not enough electricity to run all the planned data centers. The tech likely involves using battery storage to smooth demand.

AI is hitting entry-level jobs hardest, Stanford study finds

A Stanford study found that employment in AI-impacted fields for young workers (ages 16-29) is down 19% compared to more AI-resistant occupations ([Ars Technica](#)). This is the first clear evidence that AI is displacing entry-level workers, not just augmenting them. It's a politically charged finding.

4. Model & Research Watch

Kimi K3: the 2.8T parameter open-weight model that remains inaccessible to most

The massive open-weight model from Moonshot AI (China) is still not widely available — it requires too much hardware for most users ([Dataconomy](#)). It's a Mixture-of-Experts (MoE) model with 2.8 trillion total parameters but only about 280 billion active parameters per token. It rivals GPT-4 in benchmarks, but you need a cluster of H100s to run it.

Ox Alpha: anonymous release with 1-million-token context

An anonymous team released "Ox Alpha," an open-weight model with a 1-million-token context window — five times larger than GPT-4's 200K ([The Cryptonomist](#)). A 1M token context means you can feed it the entire Harry Potter series (7 books) as one prompt. The team remains anonymous, which is unusual and raises questions about provenance.

Poolside Laguna S 2.1: open-weight coding model

Poolside released an open-weight coding model, Laguna S 2.1, which is competitive with GPT-4 on coding benchmarks ([HPCwire](#)). This is another option for developers who want to run a coding model locally.

Mistral Medium 3.5 vs Llama 4 Maverick vs GPT-5.6 Luna: a \$6.30 gap

A comparison shows that Mistral Medium 3.5 is priced at \$0.15 per million input tokens, while GPT-5.6 Luna is \$6.30 per million — a **42x price difference** ([tech-insider](#)). The capability gap is narrowing, making cheaper models increasingly attractive for many tasks.

Qwen 3.8 Flash: day 0 support from Unsloth

Alibaba's Qwen 3.8 Flash model got immediate support from the Unsloth fine-tuning library, meaning developers can quickly fine-tune it on their own data ([r/LocalLLaMA](#)). This is a practical win for the open-source community.

BDH-CQ: a 150M-parameter reasoning model that achieves 29.5% on ARC-AGI at \$0.0007 per task

A small reasoning model from the USA AI Olympiad team achieved a strong score on the ARC-AGI benchmark (a test of abstract reasoning) at a fraction of the cost of larger models ([Instagram](#)). This shows that small, specialized models can compete with giants on specific tasks.

Research papers worth noting:

- **Quantization-Aware Healing:** a 4-bit compressed model that outperforms the original full-precision model ([Hugging Face](#)). This is a breakthrough: usually quantization (compressing the model to use less memory) reduces accuracy, but this technique actually improves it.
- **KVBoost:** chunk-level key-value cache reuse for efficient inference ([arXiv](#)). This technique reduces the time it takes to generate the first token (prefill latency) by reusing cached computations across requests.
- **Model of Models:** when is it better to emit a specialist model (via hypernetwork) rather than using in-context learning or fine-tuning? ([arXiv](#)) — a fundamental question about the best way to specialize models.
- **LitReview Arena:** a battle-style platform for evaluating AI-generated literature reviews ([arXiv](#)).
- **There Is No Neutral Harness:** a paper showing that LLM leaderboards are manipulable through small changes in the evaluation harness (the code that runs the benchmark) ([arXiv](#)). This is a critical finding — you can't trust leaderboard scores blindly.

Science & Frontier Tech:

- **NASA's new telescope** will study dark energy, dark matter, and exoplanets ([The News International](#)).
- **A ghostly ribbon of stars** reveals hidden dark matter ([Science Daily](#)).
- **17 spacecraft** discovered something unusual about solar flares — a coordinated observation ([Universe Space Tech](#)).
- **Pluto** got demoted 20 years ago, and we're still debating what a planet is ([National Geographic](#)).

5. What Builders Are Actually Using

Concrete swap: Mistral Medium 3.5 for GPT-5.6 Luna

If you're currently paying for GPT-5.6 Luna at \$6.30 per million input tokens, you can swap to **Mistral Medium 3.5** at \$0.15 per million — a **42x cost reduction**. The capability gap: Mistral Medium 3.5 scores about 5-10% lower on most benchmarks (e.g., MMLU, coding), but for many tasks like summarization, translation, and simple Q&A, the gap is negligible. Run it via the API from Mistral, or self-host with a 4x A100 (or a single H100) for 70B parameters. This is the single most impactful swap you can make today to cut costs without sacrificing too much quality ([tech-insider](#)).

Model routing: cheaper by design

The concept of a **model router** (sending easy queries to a small cheap model and hard ones to a big one) is now essential. With the 42x price gap above, a router that sends 80% of queries to Mistral Medium 3.5 and 20% to GPT-5.6 Luna would cut costs by roughly 80% while maintaining near-GPT quality on the tough ones. Libraries like **OpenRouter** and **Portkey** already support this. The KVBoost paper (section 4) also shows how to cache and reuse computations, which is another form of routing — routing between cached and recomputed results.

Running models locally: Apple M5 Ultra Mac Studio

The new Apple Mac Studio with M5 Ultra and up to 512GB unified memory is now the best consumer device for running large models locally ([Ars Technica](#)). Developers are already "daisy-chaining" Macs together to run models that don't fit on one. The M5 Ultra can run a 70B parameter model completely in memory (at 4-bit quantization, ~40GB). This is a real alternative to renting cloud GPUs for prototyping and fine-tuning.

Unisloth for Qwen 3.8 Flash

If you want to fine-tune a model quickly, **Unisloth** now supports Qwen 3.8 Flash on day 0 ([r/LocalLLaMA](#)). Unisloth is a library that makes fine-tuning 2x faster and uses less memory. Qwen 3.8 Flash is a 3.8B parameter model that is very fast on consumer hardware (e.g., a single RTX 4090). It's great for simple tasks like classification, extraction, and chat.

Gradio for AI workflows

If you're building an AI app, **Gradio** is the go-to tool for creating a web interface. The Hugging Face blog just published a guide on building AI workflows with Gradio ([Hugging Face](#)). It's a simple Python library that lets you turn any AI model into a web app with a drag-and-drop interface.

6. Watchlist: Earnings, Guidance & Movers

Notable movers in the watchlist (based on quote snapshot, approximate):

- **Marvell (MRVL)** +6.23% — no specific news in sources, but likely benefiting from the AI infrastructure buildout (networking chips for data centers).
- **Robinhood (HOOD)** +6.28% — likely riding the crypto rally; XRP and bitcoin are up.
- **AMD (AMD)** +4.76% — no direct news, but maybe benefiting from Nvidia earnings anticipation and the general AI hardware demand.
- **Coherent (COHR)** +4.48% — optics plays are in demand as data centers need more optical interconnects.
- **Lumentum (LITE)** +4.72% — same as Coherent.
- **Applied Optoelectronics (AAOI)** +5.22% — also optics.
- **Coinbase (COIN)** +3.66% — crypto rally.
- **Rigetti Computing (RGTI)** +3.21% — quantum computing, no specific news.
- **SoFi (SOFI)** +3.10% — fintech, possibly benefiting from consumer lending trends.
- **Digital Realty (DLR)** +2.17% — data center REIT, benefiting from AI demand.
- **Equinix (EQIX)** +2.62% — same.
- **Vistra (VST)** +2.37% — power company, benefiting from data center electricity demand.
- **Fluence Energy (FLNC)** +2.44% — energy storage, data center power.
- **Nike (NKE)** -2.69% — no direct news, but consumer spending concerns (Dick's miss) weigh on retail.
- **Starbucks (SBUX)** -1.47% — same.

Earnings reports in the last 24h: - **Dick's Sporting Goods (DKS)** — missed Q2 earnings, stock fell 25% (worst day ever). Cited "challenging" footwear market ([CNBC](#)). - **Xpeng (XPEV)** — weak delivery forecast, stock fell. Robotics unit valued at \$6.3B ([CNBC](#)). - **McKesson (MCK)** — acquired Precision Medicine Group for \$2.25B, not an

earnings report but a deal ([Yahoo Finance](#)). - **Tencent (TCEHY)** — Q2 earnings (reported Aug 12, but still in news). Spending surged, stock down 26% YTD ([CNBC](#)).

No notable news in the sources for: TSM, NEE, CRDO, RGTI, SOFI, NVDA (pre-earnings quiet), PLTR, MRVL (moved but no news), COHR, AAOI, LITE, FLNC, RDW, IBM, MU, SNDK, CEG, AMC, TSLA, GOOGL, META, MSFT, AMZN, AVGO, NET, ALAB, VST, HOOD, COIN, ARM, IMAX, CNK, CMCSA, NFLX, INTC, AMD, QCOM, KLAC, DLR, EQIX, SBUX, NKE, FXAIX.

7. Analysis: Second-Order Effects

Chain 1: US economic war on Iran -> oil drops 3% -> lower energy costs -> inflation expectations ease -> Fed has more room to cut rates -> higher bond and stock prices

- **FACTS:** Oil (Brent) fell 3% to ~\$72 after the US announced new sanctions, signaling a shift from military to economic pressure ([CNBC](#)). The war risk premium had kept oil elevated; it's now being removed.
- **INFERENCE:** Lower oil prices feed directly into headline inflation (CPI) because gasoline is a major component. If oil stays at \$70-75, US gasoline prices could fall from \$3.50 to \$3.00/gallon. This would lower the headline CPI by about 0.3-0.5 percentage points, giving the Fed cover to cut rates sooner than expected. Currently, the market sees a 1% chance of a September cut ([Intellectia](#)). If oil drops further, that probability could rise to 10-20%.
- **WEAKEST LINK:** The assumption that lower oil leads to lower core inflation (the Fed's preferred measure, which excludes food and energy). The Fed targets **core PCE** (the Personal Consumption Expenditures price index excluding food and energy). If oil drops but services inflation (rent, healthcare) remains sticky, the Fed won't cut. The weakest link is that the Fed's reaction function is more influenced by core PCE than headline CPI.
- **FALSIFYING CONDITION:** If core PCE (due out this week) comes in at 3.0% or higher, the chain breaks. Also, if oil rebounds because Iran retaliates militarily despite the economic war, the premise collapses.
- **PRECEDENT:** The 2014 oil crash (from \$100 to \$50) did not lead to aggressive Fed easing — the Fed actually raised rates in 2015 because the economy was strong.

Chain 2: China slows germanium/quartz exports to Taiwan -> semiconductor fabrication faces material shortages -> advanced chip production (TSMC) slows -> GPU supply tightens -> Nvidia and AMD benefit from higher prices, but AI buildout faces delays

- **FACTS:** China has slowed exports of germanium (used in optics and fiber optics) and quartz (used in silicon wafer manufacturing) to Taiwan ([Tom's Hardware](#)).
- **INFERENCE:** Taiwan produces over 90% of the world's advanced chips (at TSMC). If TSMC cannot get enough high-purity quartz for its wafer furnaces, it may have to reduce output. This would create a bottleneck in the global chip supply chain. The most immediate impact would be on CoWoS (Chip-on-Wafer-on-Substrate) packaging, which uses quartz components. This would delay GPU shipments and raise prices. Nvidia and AMD would benefit from pricing power, but hyperscalers would face longer wait times for their data center GPUs.
- **WEAKEST LINK:** The assumption that the slowdown is significant enough to affect production. China may be using this as a warning shot, not a full embargo. TSMC likely has 6-12 months of quartz inventory. The weakest link is the severity of the slowdown — it could be a 10% reduction, not a 50% one.
- **FALSIFYING CONDITION:** If TSMC's monthly revenue reports continue to grow (they have been up 30%+ YoY), the chain is wrong. Also, if Japan or the US can supply alternative quartz, the bottleneck breaks.

Chain 3: AI model price war (Mistral Medium 3.5 at \$0.15 vs GPT-5.6 Luna at \$6.30) -> companies adopt cheaper models for most tasks -> demand for high-end inference GPUs (H100) may plateau for smaller

workloads -> but total inference demand grows -> Nvidia's data center revenue still up, but mix shifts to lower-margin products

- **FACTS:** Mistral Medium 3.5 is 42x cheaper than GPT-5.6 Luna, with a small capability gap ([tech-insider](#)). More companies are adopting open-weight models for cost reasons.
- **INFERENCE:** As the price gap widens, companies will route the majority of their AI queries to cheap models, reserving expensive frontier models for the hardest tasks. This means the number of inference chips needed per query decreases, but the total number of queries increases dramatically (because cheap AI becomes ubiquitous). The net effect: total inference compute demand grows, but it may favor smaller, more efficient chips (like Apple's M5 Ultra or OpenAI's Jalapeño) over Nvidia's giant H100 GPUs. Nvidia's data center revenue still grows, but maybe at a slower rate, and the average selling price of chips could decline.
- **WEAKEST LINK:** The assumption that companies will actually switch. Many are locked into contracts with OpenAI or Microsoft. Switching costs (engineering time, integration) are high. The weakest link is that the price gap alone may not drive mass migration — quality and trust matter more.
- **FALSIFYING CONDITION:** If Nvidia's data center revenue guidance (reported this week) exceeds expectations by 10%+, the chain is wrong. Also, if OpenAI cuts prices to match, the gap narrows.

8. Building Your Knowledge

- **Why oil fell despite the US being at "economic war" with Iran:** Markets had already priced in a war premium — the risk of a Strait of Hormuz closure. Now that the US has shifted from bombing to sanctions, that risk is lower. Markets are forward-looking: they trade on the change in expectations, not the current state. The "safe" outcome (no war escalation) is now being discounted.
- **The difference between headline CPI and core PCE:** Headline CPI includes everything (food, energy). Core PCE excludes food and energy. The Fed targets core PCE because energy prices are volatile and driven by geopolitics, not domestic demand. This is why oil dropping doesn't automatically mean the Fed will cut rates — they look through commodity price swings.
- **What a "SPAC" is:** A Special Purpose Acquisition Company — a shell company that raises money in an IPO (initial public offering, when a company sells shares to the public for the first time) and then merges with a private company to take it public without going through the traditional IPO process. Ursa Major is going public via a SPAC deal at \$2.3 billion valuation ([Reuters](#)).
- **What "quantization" means in AI:** Quantization is the process of reducing the precision of the numbers that represent a model's weights (the parameters). Instead of using 32-bit floating point numbers, you use 8-bit or 4-bit integers. This makes the model smaller (less memory) and faster, but usually at some cost to accuracy. The new "Quantization-Aware Healing" technique improves accuracy even after quantization, which is a big deal.
- **What "context window" means:** The context window is the maximum amount of text (or tokens) a model can process at once. A 1-million-token context window (like Ox Alpha's) means you can feed it an entire book series as one prompt. The current standard is 128K-200K tokens. Larger context windows allow more complex reasoning but require more memory and compute.
- **Why the Apple M5 Ultra matters for AI:** Apple's unified memory architecture lets the CPU and GPU share the same memory pool. This means you can run a 70B parameter model entirely in RAM (with 512GB), without the bottleneck of moving data between separate GPU memory and system memory. This is a fundamentally different approach from Nvidia's, where each GPU has its own limited memory (e.g., 80GB on an H100) and you need multiple GPUs to run a large model, with lots of communication overhead.

9. Foundations & Terms

Concept Spotlight: Volatility and the VIX (the "fear gauge")

What it is: Volatility is a measure of how much a stock or market index moves up and down over a given period. The VIX (full name: CBOE Volatility Index) is a number that measures the expected volatility of the S&P 500 over the next 30 days. It is often called the "fear gauge" because it rises when investors are scared and falls when they are calm.

Why it exists: Investors need a way to quantify risk. When markets are calm, the VIX is low (around 12-15). When there is a crisis (like COVID, a war, or a financial crash), the VIX spikes to 30, 40, or even 80. The VIX is calculated from the prices of options on the S&P 500 — specifically, it measures the implied volatility : the level of future price swings that option traders are betting on. It's a forward-looking measure, not a backward-looking one.

How it shows up in the news today: With Nvidia earnings and the Jackson Hole speech this week, the VIX is likely elevated (around 20-25). The massive gold options trade (source 2) is another example of volatility trading — someone took profit on a \$200 million position in the GLD ETF, which is a bet on gold price volatility.

Why volatility itself is tradeable: You can buy and sell the VIX through futures contracts and ETFs (like the VIXY or UVXY). This allows investors to bet on whether fear will increase or decrease. But it's dangerous: the VIX is not a buy-and-hold asset — it decays over time (contango in futures markets). The most common strategy is to "buy the VIX" as a hedge : if you own stocks and fear a crash, buying VIX futures can offset losses. But if the crash doesn't happen, you lose money on the hedge.

Rough benchmarks to remember: - VIX below 12: markets are very calm, investors are complacent. - VIX 12-20: normal range. - VIX 20-30: elevated fear, possible correction. - VIX 30-40: panic territory (e.g., COVID March 2020 hit 82). - VIX above 40: tail risk, crisis mode.

Caution: The VIX does not predict the direction of the market — it only predicts the magnitude of moves. The VIX can be high even if the market is going up (if the path is choppy). Also, holding VIX ETFs for more than a few days is a losing proposition because of the futures roll cost (contango).

Concept Spotlight: Central Bank Independence

What it is: Central bank independence is the principle that a country's central bank (like the US Federal Reserve, the European Central Bank, or the Bank of Japan) should set interest rates and monetary policy (decisions about money supply and credit) without pressure from politicians. The government appoints the central bank's leaders, but once appointed, they should be free to make decisions based on economic data, not political convenience.

Why it exists: The problem: if politicians controlled interest rates, they would be tempted to lower rates before elections to boost the economy and win votes, even if it caused inflation later. This is called the "political business cycle." History shows that independent central banks produce lower inflation and more stable economic growth. The classic example is Germany's Bundesbank in the 1970s-80s, which kept inflation low while other countries with politicized central banks (like the UK and Italy) had double-digit inflation.

How it shows up in the news today: The Jackson Hole speech by Fed Chair Kevin Warsh is a test of independence. Trump has been critical of the Fed in the past, wanting lower rates. If Warsh caves to political pressure and cuts rates before the election, markets would see that as a loss of independence and long-term bond yields would rise (because investors would demand a premium for inflation risk). The fact that Bessent is considering using the Treasury's cash account to buy back bonds (source 70) is a separate issue — the Treasury is the executive branch, not the Fed. If the Treasury starts manipulating long-term bond yields, it undermines the Fed's independence.

Concrete worked example: In 2021, the Fed kept interest rates near zero even as inflation started to rise, because it believed inflation was "transitory." Some economists argued the Fed was caving to political pressure from the Biden administration, which wanted cheap borrowing. By the time the Fed started raising rates in 2022, inflation was

at 8%, and the Fed had to hike aggressively, causing a recession fear. An independent central bank with a narrower mandate would have raised rates earlier.

Rough benchmarks: The Fed's target federal funds rate is currently 5.5-5.75%. The independence score is measured by think tanks like the Central Bank Independence Index (0-1). The Fed scores about 0.8 (highly independent). The Bank of Japan scores lower (about 0.5) because the government has more influence.

Caution: Independence does not mean the central bank is unaccountable. It is accountable to Congress through oversight hearings. Independence also does not mean the central bank should ignore the government's fiscal policy — it must coordinate with the Treasury. The line between independence and coordination is blurry.

Other Terms from Today's News

- **Economic war:** A strategy of using economic tools (sanctions, tariffs, trade restrictions) to pressure a country, rather than military force. The US is waging an economic war against Iran ([NPR](#)).
- **Sanctions (trade sanctions):** Legal restrictions on trade with a target country, including banning exports, freezing assets, and prohibiting financial transactions. The US has imposed sanctions on 60 Iranian entities ([Al Jazeera](#)).
- **Secondary sanctions:** Penalties imposed on third parties (companies or countries) that do business with the sanctioned target. The US is threatening secondary sanctions on companies that trade with Iran.
- **Tariff:** A tax on imported goods. The US is imposing 25% tariffs on Canadian steel and aluminum ([CNBC](#)). Tariffs are paid by the importing company, usually passed on to consumers.
- **LNG (Liquefied Natural Gas):** Natural gas cooled to liquid form for transport by ship. The US is a major LNG exporter. The Iran sanctions could affect LNG markets if Iran's gas fields are sanctioned.
- **SPAC (Special Purpose Acquisition Company):** A shell company that raises money through an IPO to acquire a private company, taking it public. Defense firm Ursa Major is going public via a \$2.3B SPAC deal ([Reuters](#)).
- **Quantization:** In AI, reducing the precision of a model's weights (e.g., from 32-bit to 4-bit) to make it smaller and faster, often with some accuracy loss.
- **Context window:** The maximum length of input (in tokens) that a model can process at once. Ox Alpha has a 1-million-token context window ([The Cryptonomist](#)).
- **Inference:** The process of running a trained AI model on new data to generate output (e.g., asking ChatGPT a question and getting an answer). Contrast with training, which is the initial learning phase.
- **Hyperscaler:** A very large cloud provider like Amazon (AWS), Microsoft (Azure), or Google (GCP) that operates massive data centers with hundreds of thousands of servers.
- **Open-weight model:** A model whose trained parameters (weights) are publicly released, so anyone can run it, fine-tune it, or study it. Unlike closed models (like GPT-4), where you can only access it via API.
- **Benchmark:** A standardized test used to measure the performance of AI models. Examples: MMLU (multiple-choice knowledge), ARC-AGI (abstract reasoning), HumanEval (coding).
- **Options:** Financial contracts that give the buyer the right, but not the obligation, to buy or sell an asset at a specific price on or before a specific date. A "call option" is the right to buy; a "put option" is the right to sell. The VIX is derived from options prices.
- **Futures:** Contracts to buy or sell an asset at a predetermined price on a future date. Used for hedging or speculation.
- **PCE (Personal Consumption Expenditures price index):** The Fed's preferred inflation gauge. It measures the prices of goods and services that consumers buy. **Core PCE** excludes food and energy.
- **CPI (Consumer Price Index):** Another inflation measure, which includes food and energy. It's what most people think of as "inflation."
- **Buyback (share repurchase):** A company buying its own shares from the market, reducing the number of shares outstanding and increasing the value of remaining shares. Citigroup announced a \$30 billion buyback ([ts2.tech](#)).

- **Capex (capital expenditure):** Money spent on long-lived assets like buildings, machinery, or data centers. AMI is planning \$8 billion in capex for a data center ([Business Standard](#)).
- **Opex (operating expenditure):** Day-to-day costs like salaries, rent, and electricity.
- **Guidance:** A company's forward-looking statement about expected future performance (revenue, profit). It often matters more than the past quarter's results.
- **Miss (earnings):** When a company's reported earnings or revenue are lower than analysts' expectations. Dick's Sporting Goods missed expectations ([CNBC](#)).
- **Debasement trade:** Buying assets like gold, bitcoin, and foreign currencies because investors believe the government is debasing (devaluing) the currency through excessive debt and money printing. The debasement trade is back in fashion ([CNBC](#)).
- **Treasury General Account (TGA):** The US government's checking account at the Federal Reserve. It holds tax revenues and proceeds from bond sales. Bessent may use it to buy back bonds ([CNBC](#)).
- **Stablecoin:** A cryptocurrency designed to maintain a stable value, usually pegged to the US dollar (e.g., USDC, Tether). World Liberty Financial's stablecoin is being used to argue against cronyism claims ([CNBC](#)).
- **Sovereign AI:** The idea that countries and companies should build their own AI models using their own data, rather than relying on foreign providers. Thomson Reuters is doing this with its legal data ([Thomson Reuters](#)).
- **Flywheel effect:** A positive feedback loop where success breeds more success. In the AI context, more users -> more data -> better model -> more users. OpenAI's chip (Jalapeño) could create a flywheel if it improves inference speed and reduces costs, leading to more users.

10. Bottom Line

The world is fragmenting along multiple axes: the US is at economic war with Iran while simultaneously fighting a trade war with Canada; the tech war with China is intensifying through export controls, material restrictions, and criminal charges; and the Russia-Ukraine war is stuck in a stalemate with a ceasefire proposal that Moscow has already rejected. For investors, the key story this week is the perfect storm of Nvidia earnings (Wednesday) and Jackson Hole (Friday). The oil drop is a tailwind for inflation but a headwind for energy stocks. The AI price war is real and accelerating — the 42x gap between Mistral Medium 3.5 and GPT-5.6 Luna means that the era of "cheap AI" is here, and it will reshape both the model landscape and the infrastructure demand.

What changed since the previous brief: - **US-Iran war shifted from military to economic pressure.** Oil dropped 3% as war risk premium evaporated. This is a significant change in the geopolitical risk landscape. - **Canada-US trade talks collapsed.** Ottawa is now preparing retaliatory tariffs. This is a sharper escalation than the previous brief anticipated. - **Taiwan charged Nvidia and Super Micro staff** for illegal AI server exports to China. This is a major enforcement action that tightens the tech war. - **China slowed germanium/quartz exports to Taiwan** — a new supply chain threat. - **Apple launched M5 Ultra and M6 chips** with up to 512GB unified memory, explicitly targeting local AI development. This is a direct competitor to Nvidia's GPU dominance for inference. - **OpenAI unveiled its own AI chip, Jalapeño**, which beats existing chips on inference benchmarks. OpenAI is now a chip designer. - **The AI price war widened:** Mistral Medium 3.5 is now 42x cheaper than GPT-5.6 Luna, with a small capability gap. - **Dick's Sporting Goods crashed 25%** on a weak consumer outlook — a warning for retail and consumer spending. - **The Treasury is considering using its \$1 trillion TGA to buy back bonds** — a major intervention in bond markets that could distort long-term yields.

11. Sources

World & Geopolitics - [AI Jazeera](#) — Trump's latest wave of Iran sanctions: Which 60 entities are targeted? - [NPR](#) — Trump calls leaders to isolate Iran. And, [SCOTUS backs Trump](#) for now on mail-in voting - [PBS](#) — Iran's president says U.S. memorandum is best path out of stalled war - [Iran International](#) — [Iran shrugs off US economic war as analysts sound alarm](#) - [CNBC](#) — [Oil drops 3% as U.S. shifts to economic pressure on Iran](#) - [CNBC](#) — [Canada poised to retaliate against Trump tariffs](#) - [CNBC](#) — [Canadian dollar slides as Ottawa and Washington head for all-out trade war](#) - [SBS Korea](#)

— Ukraine Proposes Ceasefire Plan with 'Donbas Free Economic Zone'; Russia Rejects - Reuters — Ukraine's Zelenskiy sees window for U.S. mediation until summer 2027 - Mezha — More Than 50 Countries Urge Russia to Accept an Immediate Ceasefire in Ukraine - France 24 — Russian drones stalk Ukraine's civilians in 'human safaris' - SecurityWeek — Taiwan Charges 9 Over Illegal AI Server Exports to China, Including Nvidia and Super Micro Staff - Tom's Hardware — China strategically slows exports of critical materials used in semiconductor fabrication to Taiwan - Time — Supreme Court Win for Trump Mail-in Voting Restrictions - Guardian — Supreme court threatens midterms mail-in voting as it backs Trump plan - Guardian — Colombia's new president announces immigration crackdown - Firstpost — Kiribati detains Chinese vessel over illegal shark finning - Al Jazeera — US military transport plane arrives in Moscow - BBC — Two unvaccinated people die of measles in Pennsylvania - France 24 — French village takes stock of tornado aftermath - France 24 — Anger mounts in Haiti as new attack shows power held by gang coalition - KELO-AM — Bangladesh could consider joining Mecca pact - Al Jazeera — Trump plans largest mass visa cancellation in history

Markets, Money & Deals - CNBC — Oil drops 3% as U.S. shifts to economic pressure on Iran - CNBC — A massive trade just happened in gold. The options market is buzzing - CNBC — Dick's Sporting Goods stock falls 25% as retailer misses expectations - MarketWatch — Dick's Sporting Goods' stock is having its worst day ever - ts2.tech — Citigroup's \$30 Billion Share Repurchase Compares to Regional Bank Acquisition Choices - Yahoo Finance — McKesson to buy Precision Medicine Group in about \$2.25 billion deal - WION — Paramount-Warner merger faces fresh hurdle as California attorney cancels settlement meeting - Bloomberg — European Stocks Advance as Oil Prices Drop; Siemens Energy Up - Economic Times — Jackson Hole 2026: When and what are investors expecting from Fed Chair Kevin Warsh? - Intellectia — Nvidia Earnings & Jackson Hole: AI Stocks Face Fed Rate - CNBC — Bessent could tap near \$1 trillion Treasury General Account to fund bond buybacks - CNBC — The 'debasement trade' returns after Bessent bond maneuver. Crypto and gold are back in style - CNBC — Xpeng shares sink as weak delivery forecast overshadows \$6.3 billion robot unit valuation - MEXC — Why Is XRP Going Up? Ripple's RLUSD Credit Fund Adds a New Institutional Catalyst - Crypto News — South Korea speeds up Digital Asset Framework Act for fall passage - CNBC — World Liberty Financial CEO says success of stablecoin silences charge of Trump cronyism

AI & Infrastructure - CNBC — Nvidia's dependence on hyperscalers faces big test in earnings report - BigGo Finance — Nvidia Reportedly Raising Prices 15%; This Week's Earnings and Fed Speech Key to US Stock Market Direction - Business Standard — AMI orders 9,000 NVIDIA Rubin GPUs, plans \$8 bn capex for 200 MW capacity - Apple — Apple introduces M6 and M5 Ultra for a big leap in performance and AI compute - Tom's Hardware — Apple launches new M6 and M5 Ultra Apple silicon chips - CNBC — Apple announces new Mac Mini and Mac Studio models with AI upgrades - Ars Technica — Apple's new desktop computers are designed specifically for local AI development - TechCrunch — Apple debuts its 'most powerful chip ever' in M5 Ultra and M6 - TechCrunch — OpenAI's Jalapeño chip is built for fast inference at scale, benchmarks show - The Verge — OpenAI says its Jalapeño chip can power faster AI responses than the competition - Thomson Reuters — Thomson Reuters Launches Its Own Frontier Model - CNBC — Tencent sees spending surge, defends potential 'superior' AI returns - HPCwire — Google Cloud Announces Strategic Partnership with Verizon to Scale Enterprise AI - HPCwire — UC San Diego to Test Innovative Power Tech for AI Data Centers - Tech Times — GaAs Quantum-Dot Laser Ships to AI Labs as China's InP Export Controls Expose Supply Ceiling - Ars Technica — AI is hitting entry-level jobs hardest, Stanford study finds - Reuters — Defense firm Ursa Major to go public through \$2.3 billion SPAC deal

Model & Research Watch - Dataconomy — Kimi K3: The Advanced AI Model That Remains Inaccessible To Most Users - The Cryptonomist — Anonymous AI model release Ox Alpha stuns developers with 1M-token context - HPCwire — Poolside Launches Laguna S 2.1 Open-Weight Coding Model - tech-insider — Mistral Medium 3.5 vs Llama 4 Maverick vs GPT-5.6 Luna: \$6.30 Gap - r/LocalLLaMA — Qwen 3.8 Flash Next day 0 support from unsloth - Instagram — BDH-CQ: 150M-parameter reasoning model on ARC-AGI - Hugging Face — Quantization-Aware Healing: a compressed, 4-bit model that outperforms its full-precision original - arXiv — KVBoost: Chunk-Level Key-Value Cache Reuse - arXiv — Model of Models: When Does Emitting a Specialist Beat Attending, Adapting, or Tuning? - arXiv — LitReview Arena: Evaluating Literature Review Agents - arXiv — There Is No Neutral Harness: Modern LLM Leaderboards Are Manufactured by Config-Fragile Items - The News International — NASA's new telescope set to uncover mysteries of dark energy, dark matter and exoplanets - Science Daily — A ghostly ribbon of stars reveals hidden dark matter - Universe Space Tech — 17 spacecraft discovered something unusual about solar flares - National Geographic — Pluto Got Demoted 20 Years Ago. We're Still Not Sure What a Planet Really Is

What Builders Are Actually Using - tech-insider — Mistral Medium 3.5 vs Llama 4 Maverick vs GPT-5.6 Luna: \$6.30 Gap - Ars Technica — Apple's new desktop computers are designed specifically for local AI development - r/LocalLLaMA — Qwen 3.8 Flash Next day 0 support from unsloth - Hugging Face — Wire It, Run It, Deploy It: AI Workflows in Gradio - Pasquale Pillitteri — Can Your PC Run AI Locally? 8 Free Tools to Find Out

Generated 2026-08-25 07:58 PDT · 130 sources · model: deepseek-ai/DeepSeek-V4-Flash via <https://api.gmi-serving.com/v1>