

LAKSH'S DAILY BRIEF

Laksh's Daily Brief — 2026-08-26

Wednesday, August 26, 2026

World · Markets & Money · AI & Infrastructure · Models & Research



BBC World — Door was locked at Pakistan hospital where fire killed 14 babies, witnesses say



Al Jazeera — Meta agrees to settlement, platform changes in youth addiction case



Al Jazeera — Trump tells Al Jazeera "not in a hurry" for Iran to return to talks

1. World & Geopolitics

Iran warns of military ship ban in the Strait of Hormuz, but talks with Oman continue

Iran's top diplomat met with Omani officials to discuss managing commercial shipping through the Strait of Hormuz, the vital waterway that has been largely shut down for nearly six months since the Iran war began. Iran warned it may ban military ships from the strait, and accused the US of standing in the way of a deal ([CNBC](#)). Separately, President Trump told Al Jazeera he is "not in a hurry" for Iran to return to talks, adding "I have no time schedule" ([Al Jazeera](#)). The Strait of Hormuz is a chokepoint — a narrow passage that a large fraction of the world's oil passes through. If commercial shipping remains blocked or military vessels are banned, the risk of supply disruptions keeps oil prices volatile. But the fact that Iran and Oman are negotiating suggests both sides may want to avoid a permanent closure.

Canada hits US with retaliatory tariffs, escalating trade war

Canada announced it will impose countertariffs on the US "dollar for dollar" after trade negotiations broke down. Paper products, including toilet paper, are among the hardest-hit sectors, with estimated tariffs of 25-50% ([NPR](#); [The Guardian](#)). This is a full-fledged trade war between two neighbours that have had decades of peaceful trade. The US imports a significant amount of lumber and paper products from Canada, so US consumers will see price increases at the checkout counter. The broader impact: higher costs for US

manufacturers who use Canadian inputs, and potential retaliation on other sectors.

Ukraine ceasefire plan with buffer zone proposed by US and Europe, but Russia is skeptical

A new ceasefire plan has been developed by Ukraine, the US, and Europe, which includes a buffer zone. A Vatican envoy arrived in Russia to deliver the proposals, but Russia has rejected similar plans before ([UATV](#)). Meanwhile, more than 50 UN countries called on Russia for a complete and immediate ceasefire ([UATV](#)). The key question: will Russia accept any deal that doesn't include territorial concessions? The Vatican envoy's visit is a rare diplomatic channel, but the odds of a breakthrough are low.

China pressures Eswatini, the last African ally of Taiwan

China advised its citizens to leave Eswatini, the only African country that maintains diplomatic relations with Taiwan, citing "security risks" ([BBC](#)). This is a classic coercive tactic: China uses economic and political pressure to isolate Taiwan diplomatically. Eswatini is one of only 12 countries that still recognise Taiwan. The "security risks" warning is a pretext for Beijing to signal that it may take action against the country.

Nepal flash floods: hundreds missing, villages destroyed

Massive flash floods swept through the Himalayan border regions of Nepal and Tibet, killing at least 95 people, with hundreds more missing, including foreign

tourists. The disaster is linked to climate change, as the region is vulnerable to extreme weather ([France 24](#); [NPR](#)). Flood alerts have also been issued in the Indian states of Bihar and Uttar Pradesh. This is a humanitarian tragedy with significant infrastructure damage.

Meta settles social media addiction lawsuit for \$16.7 billion

Meta agreed to pay \$16.7 billion to settle claims by 47 US states that it designed Facebook and Instagram to deliberately addict teens while deceiving the public. The company also agreed to add child-safety measures ([CNBC](#); [Reuters](#)). This is one of the largest state-level settlements in history. The amount is large but manageable for Meta, which had \$100B+ in revenue. The real impact is the operational changes: teen accounts will have stricter defaults, and the company

faces ongoing oversight. The stock was up 2.71% in the snapshot, suggesting investors see the settlement as removing a major legal overhang.

Other notable stories: - Pakistan hospital fire kills 14 newborns in Islamabad, witnesses say the ward door was locked ([BBC](#); [Al Jazeera](#)). - South Korea: bodies of four missing people found on Jeju Island, police officer arrested, national scandal ([Al Jazeera](#)). - Dolly Parton, country music legend, dies at 80 ([France 24](#)). - US country music legend Dolly Parton dies, aged 80. - CISA confirms hackers targeted over 100 US water systems in July, suspected Iran-backed ([TechCrunch](#)). - Mozambique violence displaces over 100,000 people last year, adding to 1.3 million already displaced ([Al Jazeera](#)).

2. Markets, Money & Deals

Markets wait for Jackson Hole, Nvidia earnings; Fed's preferred inflation gauge shows 3.3% core PCE

The stock market was mixed on Wednesday: Dow Jones futures rose 0.11%, S&P 500 slipped 0.04%, Nasdaq 100 fell 0.15%, and Russell 2000 was flat. The 10-year Treasury yield was 4.64%, the two-year at 4.19%. The CME FedWatch tool shows a 36.1% probability of a rate hike at the September meeting ([TradingView](#)). The Fed's preferred inflation gauge, the core PCE (Personal Consumption Expenditures price index excluding food and energy), rose 3.3% annually in July, slightly below the 3.6% expected ([CNBC](#)). This is good news: inflation is cooling, but still above the Fed's 2% target. The Jackson Hole symposium (Aug 21-23) is the key event this week, where Fed Chair Warsh may signal the path of rates. The market is pricing in a small chance of a rate hike, which would be a surprise if it happens.

Nvidia earnings preview: the quarter may be priced in, but guidance is the trade

Nvidia reports Q2 earnings on August 28 (after market close). The stock is down 7.35% over the past week, setting up for a potential "buy the rumour, sell the news" or a beat-and-raise. Analysts note that combined hyperscaler capex (capital expenditure on data centres and AI infrastructure) is tracking ~\$700 billion for 2026, with Amazon alone guiding to ~\$220B. Meta narrowed

its full-year outlook upward to \$130-145B ([Investing.com](#)). The key question: will Nvidia's guidance beat the whisper number? If the forward number is strong, the stock could rally despite the recent weakness.

Meta settles for \$16.7B, stock up despite the cost

Meta's stock rose 2.71% in the snapshot, indicating that the settlement removes a major legal risk. The \$16.7B payout is large but spread over time, and Meta's operating cash flow is around \$80B per year. The company also agreed to operational changes that may reduce engagement among teens, but investors seem to think the cost is manageable.

Bessent's bond market intervention: debasement trade returns

Treasury Secretary Bessent's recent bond buyback announcement has generated controversy. Billionaire Ray Dalio said it signals a debt crisis is approaching and recommends gold and bitcoin ([CNBC](#)). Stanley Druckenmiller also doubts the ploy will work ([CNBC](#)). The "debasement trade" — buying gold and bitcoin as a hedge against currency debasement — is back in style ([CNBC](#)). A massive gold options trade was spotted: someone took profit on a \$200 million GLD position ([CNBC](#)). This suggests big money is betting on continued gold strength.

Dick's Sporting Goods stock falls 30% on earnings miss

Dick's Sporting Goods reported fiscal Q2 earnings that missed expectations, citing a "challenging" footwear market. The stock fell 30% in after-hours trading ([CNBC](#)). This is a warning sign for consumer discretionary spending: if people are cutting back on sneakers and sports gear, it could signal broader economic weakness.

Other notable moves: - Intuit shares slid after reports that customers are fleeing TurboTax over price ([MarketWatch](#)). - Haidilao shares jumped on strong delivery growth ([CNBC](#)). - Xpeng shares sank on weak delivery guidance, even as its robotics business got a \$6.3B valuation ([CNBC](#)). - LEGO posts record first-half revenue, sees strong demand at both ends of pricing ([CNBC](#)). - Commercial real estate investor competition sees strongest growth in a year ([CNBC](#)).

3. AI & Infrastructure

Nvidia's Hot Chips 2026: power is the hard limit, DSX MaxLPS tech allows more compute per fixed power budget

At Hot Chips 2026, Nvidia emphasised that power is the hard constraint on data centre capacity. Their new DSX MaxLPS approach manages site power to allow more compute from a fixed 100MW facility budget, using the Vera Rubin NVL72 systems ([Tom's Hardware](#)). This is important for the AI buildout: data centre operators are increasingly power-constrained, and any technology that squeezes more performance per watt will be adopted. For investors, this means companies that provide power-efficient cooling, power management, or grid interconnection services could benefit.

OpenAI's Jalapeño chip beats Nvidia Blackwell on inference efficiency, threatening Nvidia's margins

OpenAI's custom AI chip, codenamed Jalapeño, beat Nvidia's Blackwell systems on key inference-efficiency tests. This is a major development: if OpenAI can design chips that are more efficient for running AI models (inference), it could reduce its dependence on Nvidia GPUs and drive down costs for all AI companies ([CNBC](#)). For Nvidia investors, this is a threat to the high margins Nvidia enjoys on its data centre GPUs. However, Nvidia's scale and software ecosystem (CUDA) remain formidable barriers.

Alibaba releases Qwen3.8-Flash, a cost-effective open-weight model

Alibaba released Qwen3.8-Flash, a 125-billion-parameter model that is open-weight (weights are publicly downloadable). It competes with Anthropic's Opus 4.6 and DeepSeek's V4-Flash on performance, at a fraction of the cost ([Bloomberg](#)). The model uses a hybrid attention mechanism and is designed to be

efficient. This continues the trend of powerful open-weight models that undercut proprietary ones on price.

GLM-5.3-Flash from Z.ai (the lab behind Ox Alpha) — frontier intelligence at flash cost

Z.ai, the lab behind the mysterious Ox Alpha model, released GLM-5.3-Flash, which they claim delivers frontier intelligence at a flash cost (i.e., cheap inference). The model is open-weight and has been discussed on Hacker News and Reddit ([Hacker News](#); [Reddit](#); [TechCrunch](#)). This is another example of the model price war: open-weight models are approaching frontier performance.

Thomson Reuters launches its own frontier model for legal and professional services

Thomson Reuters, the company behind Westlaw and Reuters, launched its own large language model, Thomson 1.0, targeting legal professionals. This is a "vertical model" — a model trained on domain-specific data (legal documents, case law) to outperform general-purpose models in that field ([Thomson Reuters](#)). Google also expanded its Gemini Enterprise AI platform for lawyers ([American Bazaar](#)). This signals that enterprise AI is moving from general-purpose chatbots to domain-specific tools.

Poolside releases Laguna S 2.1, an open-weight coding model that matches much larger models

Poolside released Laguna S 2.1, a 118-billion-parameter open-weight model built for agentic coding (AI that can write and execute code autonomously). It matches or exceeds models several times its size on coding benchmarks, and can run on a single NVIDIA DGX Spark ([HPCwire](#)). This is good for developers who want

to run a powerful coding model locally without paying per-token fees.

Goldman Sachs: China's 7nm chip supply gap will shrink from 92% to 34% by 2035, but TSMC remains far ahead

Goldman Sachs assessed that China's domestic 7nm chip supply shortfall will shrink from 92% to 34% by 2035, meaning China will be able to meet 66% of its

advanced chip demand. However, TSMC's 7nm yields exceed 90%, while SMIC (China's leading foundry) lags significantly ([CHOSUNBIZ](#); [BigGo Finance](#)). This is a long-term trend: China is slowly closing the gap, but the technology gap remains substantial. For investors, this means the US export controls on advanced chipmaking equipment are slowing but not stopping China's progress.

4. Model & Research Watch

New model launches:

- **Alibaba Qwen3.8-Flash-Next:** 125 billion parameters, open-weight, hybrid attention, competitive with Anthropic Opus 4.6 and DeepSeek V4-Flash. Released for download on August 26 ([Reddit megathread](#); [Bloomberg](#)).
- **Z.ai GLM-5.3-Flash:** Frontier intelligence, flash cost, open-weight. Minimal details on params, but claims high performance ([Hacker News](#)).
- **Poolside Laguna S 2.1:** 118B parameters, open-weight, agentic coding, matches DeepSeek-V4-Flash and Nemotron 3 Ultra on coding benchmarks. Runs on a single DGX Spark ([HPCwire](#)).
- **Thomson Reuters Thomson 1.0:** Frontier model for legal, no parameter count disclosed, but trained on proprietary legal data ([Thomson Reuters](#)).
- **IBM Granite 4.2 LLMs:** Details on how they are built, open-weight ([Hugging Face Blog](#)).

Notable research papers:

- **RENDER benchmark:** A new evaluation for how LLMs handle memory and retrieval-augmented generation (RAG) by varying the format of the provided context ([arXiv](#)).
- **ESQ-Bench:** A benchmark for evaluating NL2SQL (natural language to SQL) models on enterprise

Oracle databases, showing that current models fail on real-world database schemas ([arXiv](#)).

- **Quantization-Aware Healing:** A technique to compress a model to 4-bit that outperforms its full-precision original, using a healing step after compression ([Hugging Face Blog](#)).
- **LLMs could control their host machines by exploiting inference engines:** A security paper showing that LLMs can escape their sandbox by manipulating the inference engine itself ([Hacker News](#)).
- **Data Predictability Shapes Weibull Weight-Scale Growth in Transformer Training:** A theoretical paper finding that the scale of weight magnitudes in trained transformers is governed by the predictability of the training data ([arXiv](#)).

Science & Frontier Tech note: - A Seattle company attempts to simulate 15,000 years of solar wind in four hours in a lab, aiming to study space weathering on the Moon ([Ars Technica](#)). - Psychiatry's diagnostic manual is getting a biological upgrade, incorporating biomarkers for the first time ([Ars Technica](#)). - Mpox is back in new countries, with children particularly affected ([The Guardian](#)).

5. What Builders Are Actually Using

- **Concrete swap: Qwen3.8-Flash-Next (125B, open-weight) can replace Anthropic Opus 4.6 for many general-purpose tasks.** The capability gap is small on most benchmarks (Alibaba claims competitive performance). Cost: Qwen3.8-Flash-Next is free to

download and self-host; you need roughly 2x NVIDIA A100 GPUs (80GB each) to run it at full precision, or 1x A100 with quantization (4-bit). Anthropic Opus 4.6 costs about \$15 per million input tokens and \$60 per million output tokens. If you self-host Qwen3.8-Flash-

Next, your cost is electricity and hardware depreciation — roughly \$0.50 per million tokens depending on GPU pricing. This is a 30-120x cost saving for comparable performance.

- **Model routing** is gaining traction: use a small cheap model (like Qwen3.8-Flash-Next or GLM-5.3-Flash) for 80% of queries, and route only the hard ones to a larger model. Tools like OpenRouter and LiteLLM support this. In practice, this can cut inference costs by 50-80% without noticeable quality loss.
- **Running models locally** is becoming more practical: the DGX Spark (an NVIDIA desktop machine) can run Laguna S 2.1 (118B) at usable speeds. Developers are using ollama and LM Studio for local inference.

The latest quantization techniques (like Quantization-Aware Healing, source 88) allow 4-bit models to outperform full-precision ones, meaning you can run larger models on less hardware.

- **Gradio workflows** are being used to wire up AI pipelines: the Hugging Face team released a guide on building AI workflows with Gradio, allowing developers to chain models, APIs, and logic without writing a full web app ([Hugging Face Blog](#)).
- **Sentence Transformers** now support training multi-vector embedding models, which are useful for more accurate retrieval in RAG systems ([Hugging Face Blog](#)).

6. Watchlist: Earnings, Guidance & Movers

No notable earnings reports from watchlist companies in the last 24 hours (Nvidia reports tomorrow, Aug 28). The key movers based on the quote snapshot:

- **META**: +2.71% — settlement removed a major legal overhang.
- **NET**: +2.77% — no specific news in sources, but likely riding the AI infrastructure theme.
- **LITE**: +2.08% — no specific news, optics plays may be benefiting from AI data centre demand.
- **AMC**: +2.74% — no specific news, meme stock volatility.
- **FLNC**: -3.07% — no specific news, but energy storage stocks may be sensitive to power market dynamics.
- **COIN**: -3.21% — Bitcoin and gold are in vogue, but COIN may be volatile due to crypto regulatory noise.

- **RGTI**: -4.49% — quantum computing stocks are speculative and often move on sentiment; no news in sources.
- **CRDO**: -2.01% — no specific news, but Credo is an optics/networking play; may be profit-taking ahead of Nvidia.
- **NKE**: -2.28% — no specific news, but consumer discretionary weakness may be weighing.
- **HOOD**: -2.36% — no specific news, but crypto-related volatility.

No news in sources for the following watchlist

names: NEE, CRDO, RGTI, SOFI, NVDA (preview only), PLTR, MRVL, COHR, AAOI, LITE, FLNC, RDW, IBM, MU, SNDK, CEG, AMC, TSLA, GOOGL, META, MSFT, AMZN, AVGO, NET, ALAB, VST, HOOD, COIN, ARM, IMAX, CNK, CMCSA, NFLX, INTC, AMD, QCOM, KLAC, DLR, EQIX, SBUX, NKE, FXAIX.

7. Analysis: Second-Order Effects

Chain 1: Iran's Hormuz threat -> war risk premium -> oil prices -> inflation -> Fed policy

- **FACTS:** Iran warned of a military ship ban in the Strait of Hormuz, but talks with Oman continue. The strait has been largely shut for six months. Trump is not in a hurry for talks. ([Sources 25, 97, 17])
- **INFERENCE:** If talks fail and Iran bans military ships, the risk of a confrontation increases, pushing up war-

risk insurance premiums on tankers. This would raise the cost of shipping oil through the strait, leading to higher oil prices (Brent crude). Higher oil prices feed into headline inflation, which could push the Fed toward a more hawkish stance (rate hike or hold, rather than a cut). Jackson Hole is this week, and the Fed's preferred inflation gauge (core PCE) came in at

3.3%, still above target. If oil spikes, it could delay rate cuts.

- **WEAKEST LINK:** The assumption that commercial shipping would be disrupted. Oman and Iran are approaching a deal, and the US may also be negotiating. If a deal is reached, the threat disappears. Also, oil prices are already elevated; a further spike may be limited.
- **WHAT WOULD FALSIFY IT:** A deal securing safe transit through Hormuz is announced, causing oil to drop. Or the Fed signals it will look through a temporary oil spike (as it did in 2022).

Chain 2: Trade war with Canada -> higher input costs for US manufacturers -> lower corporate margins -> slower hiring

- **FACTS:** Canada announced retaliatory tariffs on US goods, "dollar for dollar." Paper products face 25-50% tariffs. ([Sources 37, 45])
- **INFERENCE:** US manufacturers who rely on Canadian inputs (lumber, paper, aluminum, etc.) will see their costs rise. They may pass these costs to consumers (higher prices) or absorb them (lower profit margins). If margins shrink, companies may cut hiring or delay investment. This is a negative for the US economy and could weigh on the stock market.
- **WEAKEST LINK:** The assumption that tariffs are broad and long-lasting. Trade wars often escalate in rhetoric but can de-escalate through negotiation. If the US and Canada reach a deal quickly, the impact is minimal.
- **WHAT WOULD FALSIFY IT:** A quick agreement to suspend tariffs, or US companies announcing they will relocate production to the US (reshoring) to avoid tariffs.

Chain 3: AI model price war -> lower inference costs -> higher demand for AI inference -> more data centre buildout -> more demand for power and cooling

- **FACTS:** Multiple open-weight models released with competitive performance at much lower cost: Qwen3.8-Flash-Next, GLM-5.3-Flash, Poolside Laguna S 2.1. OpenAI's Jalapeño chip beats Nvidia on inference efficiency. Combined hyperscaler capex tracking ~\$700B for 2026. ([Sources 4, 23, 36, 106, 63])

- **INFERENCE:** As inference costs drop, more applications become economically viable. This drives higher adoption of AI, which in turn increases total demand for inference compute. Even though the cost per token is lower, the volume increases so much that total compute demand grows. This benefits data centre operators, power companies, and networking/optics providers. However, if margin compression hits Nvidia, GPU sales could be affected.
- **WEAKEST LINK:** The assumption that volume growth outpaces price decline. It's possible that the market hits a saturation point where the number of applications doesn't grow fast enough. But historical precedent (e.g., Moore's Law) suggests that lower cost drives more usage.
- **WHAT WOULD FALSIFY IT:** A major slowdown in AI adoption, or a shift to smaller models that require less compute, reducing total demand.

Chain 4: Meta settlement -> operational changes for teen safety -> lower engagement among younger users -> slower ad revenue growth -> Meta's stock re-rates

- **FACTS:** Meta settled for \$16.7B and agreed to heavy restrictions on teen users, including default privacy settings and limits on features. ([Sources 2, 5, 13, 43, 113, 116])
- **INFERENCE:** The settlement forces Meta to change how its platforms work for teens, potentially reducing time spent and engagement. Since younger users are an important demographic for future ad revenue, this could slow Meta's growth. However, the stock went up after the settlement, suggesting investors think the cost is manageable and the risk is removed. The key question is whether the operational changes materially affect ad revenue.
- **WEAKEST LINK:** The assumption that teen engagement is a significant driver of Meta's ad revenue. Meta's ad business is driven by all ages, and teens may not be the highest-value demographic. Also, the changes may be minor in practice.
- **WHAT WOULD FALSIFY IT:** Meta's next quarter shows no slowdown in user growth or ad revenue among young demographics.

8. Building Your Knowledge

- **Central bank guidance is a powerful tool, but it can backfire.** The Fed's Jackson Hole symposium is where the chair signals the future path of rates. If the Fed provides too much forward guidance, markets can become dependent on it. The current Fed chair, Warsh, may be trying to reduce guidance to let markets function on their own. That transition can cause volatility.
- **Trade wars harm consumers more than they help.** The Canada-US trade war is a textbook example: tariffs on paper products will raise prices for toilet paper and other goods for US consumers. The "dollar for dollar" retaliatory tariffs mean both sides lose. The only winners are domestic producers who face less competition, but they also face higher input costs.
- **The model price war is accelerating.** With open-weight models like Qwen3.8-Flash-Next and GLM-5.3-Flash approaching frontier performance at a fraction of the cost, the economics of AI are shifting from "pay per token" to "buy hardware once and run for free." This is great for startups and developers, but it puts pressure on proprietary model providers (OpenAI, Anthropic) to differentiate on quality or features.
- **Custom AI chips (like OpenAI's Jalapeño) are a threat to Nvidia's dominance.** If major AI companies design their own chips, they can optimise for their specific workloads and reduce dependence on Nvidia. This is a long-term risk for Nvidia's margins, but Nvidia's software ecosystem (CUDA) and scale are significant moats.
- **Power is the new bottleneck for AI.** Nvidia's Hot Chips presentation emphasised that data centre capacity is limited by power, not just chips. Technologies that improve power efficiency (like DSX MaxLPS) will be in high demand. This is why power companies, grid interconnection firms, and cooling technology providers are becoming key beneficiaries of the AI buildout.
- **Settlements can be positive for stocks if they remove uncertainty.** Meta's \$16.7B settlement is a huge number, but the stock rose because the legal overhang is gone. The market prefers certainty over a potentially larger judgment or ongoing litigation. This is a pattern: a known, manageable cost is better than an unknown risk.

9. Foundations & Terms

Concept Spotlight: Options (Calls and Puts)

What they are in plain words:

An option is a contract that gives you the **right, but not the obligation**, to buy or sell a stock at a specific price before a specific date. Think of it like a coupon: you pay a small fee (the premium) to lock in a price, and then you decide whether to use it or let it expire.

There are two types: - **Call option:** The right to **buy** a stock at a fixed price. You buy a call if you think the stock will go up. - **Put option:** The right to **sell** a stock at a fixed price. You buy a put if you think the stock will go down.

Why they exist / what problem they solve:

Options allow investors to bet on price movements without owning the stock. They also allow hedging: if you own a stock and are worried it might fall, you can buy a

put to limit your losses. The key is **leverage** — a small amount of money (the premium) controls a larger amount of stock.

Concrete worked example with real numbers:

Let's say Nvidia (NVDA) is trading at \$200. You think it will go to \$220 in the next month. You could buy 100 shares for \$20,000, or you could buy a call option with a strike price of \$200 that expires in one month. The premium might be \$5 per share, or \$500 per contract (each contract controls 100 shares).

- If Nvidia goes to \$220, your call option is worth at least \$20 per share (\$2,000 per contract). You made \$1,500 profit on a \$500 investment — a 300% return. If you bought the stock, you made \$2,000 on a \$20,000 investment — a 10% return.

- If Nvidia stays at \$200 or falls, the option expires worthless, and you lose the entire \$500 premium. The stockholder would lose nothing or have a smaller loss.

How it shows up in the news:

Today's brief mentioned a massive gold options trade: someone took profit on a \$200 million GLD position ([Source 115](#)). Options are often used by large investors to make leveraged bets or hedge.

Rough benchmarks worth remembering: - A one-month at-the-money option on a stock typically costs 2-5% of the stock price. - Options lose value as they approach expiration (time decay). - The "strike price" is the price at which you can buy or sell.

Caution / common misconception: - Options are risky. You can lose 100% of your investment (the premium) if the stock doesn't move the way you expect. - Options are not the same as owning the stock. You don't get dividends or voting rights. - Selling options (writing options) is a different strategy with unlimited risk — never confuse buying with selling.

Concept Spotlight: Coalition Government

What it is in plain words:

A coalition government is formed when no single political party wins a majority of seats in parliament, so two or more parties agree to share power and govern together. They negotiate a deal that includes joint policy goals and a division of cabinet positions.

Why it exists / what problem it solves:

In many democracies, especially those with proportional representation (where seats are allocated based on percentage of votes), it's rare for one party to win more than 50% of seats. Coalitions ensure that the government represents a broader set of voters and prevents a single party from having unchecked power.

Concrete worked example with real numbers:

In the 2025 German federal election, the Social Democratic Party (SPD) won 30% of seats, the Greens won 20%, and the Free Democratic Party (FDP) won 10%. No party had a majority. The SPD, Greens, and FDP formed a coalition (a "traffic light" coalition, named after their party colours). They negotiated a common policy agenda: the SPD got the chancellorship, the Greens got the environment ministry, and the FDP got

the finance ministry. The coalition lasted until disagreements over the budget caused it to collapse.

How it shows up in the news:

Today's brief had no specific coalition story, but a common example is Israel: in 2024-2025, Israel's government was a coalition of several parties, including far-right and centrist factions. The coalition's stability affects policy on Gaza and the West Bank. Another example: Germany's coalition collapsed earlier this year, leading to new elections.

Rough benchmarks worth remembering: - Most European countries have coalition governments (Germany, Italy, Netherlands, etc.). - Coalition governments tend to be less stable than single-party majority governments — they can collapse if a partner withdraws. - The "kingmaker" is a small party that can tip the balance between two larger blocs.

Caution / common misconception: - Coalition governments are not always weak. Some are very stable (e.g., Germany's "grand coalition" of CDU and SPD lasted for years). - They are not the same as a "unity government" (which is formed during a crisis and includes all major parties). - Not all democracies use coalitions: the UK and Canada often have single-party majority governments (though coalitions can occur).

Other Terms from Today's News

- **Core PCE:** The Fed's preferred inflation gauge, which measures personal consumption expenditures excluding food and energy prices. It rose 3.3% annually in July, below the 3.6% expected. This is important because the Fed targets 2% core PCE for its interest rate decisions.
- **Debasement trade:** A strategy of buying assets like gold and bitcoin that are expected to hold value when the purchasing power of currency declines due to government debt and money printing. Today's news: the Bessent bond buyback has revived the debasement trade ([Source 50](#)).
- **Jackson Hole:** An annual symposium hosted by the Federal Reserve Bank of Kansas City in Jackson Hole, Wyoming. It's where the Fed chair often signals major policy shifts. This year's event (Aug 21-23) is key for rate expectations.
- **Hyperscaler:** A large cloud provider like Amazon (AWS), Microsoft (Azure), or Google (GCP) that

builds massive data centres to offer computing and storage at scale. Combined hyperscaler capex is tracking ~\$700B for 2026 ([Source 63](#)).

- **Open-weight model:** A model whose trained parameters (weights) are publicly released, allowing anyone to download, run, and modify it. Qwen3.8-Flash-Next and GLM-5.3-Flash are open-weight models.
- **Inference:** The process of running a trained AI model on new data to generate a prediction or output (as opposed to training, which is the learning phase).

OpenAI's Jalapeño chip is optimised for inference efficiency.

- **Forward guidance:** Communication from a central bank about its future policy intentions. The Fed's Jackson Hole symposium is a key forward guidance event.
- **Yield curve:** A graph showing the interest rates (yields) of bonds with different maturities. Today, the 10-year Treasury yielded 4.64% and the 2-year yielded 4.19%, a normal (upward-sloping) curve.

10. Bottom Line

The world is navigating multiple crises simultaneously: the Iran standoff over Hormuz, a trade war with Canada, the Ukraine ceasefire stalemate, and China's diplomatic pressure on Taiwan. For investors, the key focus this week is the Jackson Hole symposium and Nvidia earnings. The core PCE came in lower than expected, giving the Fed some room to hold rates steady, but the 36% chance of a rate hike in September shows the market is still uncertain. The AI model price war is accelerating, with powerful open-weight models being released at a fraction of the cost of proprietary ones. This is great for builders but puts pressure on Nvidia's margins as custom chips like OpenAI's Jalapeño emerge. The debasement trade (gold, bitcoin) is back in vogue as the government's debt management raises concerns.

What changed since the previous brief: - **Core PCE came in at 3.3% vs 3.6% expected** — a positive surprise that could ease rate hike fears. - **Iran's Hormuz threat escalated** — Iran warned of military ship ban, but talks with Oman continue. The previous brief had oil dropping 3% as war risk premium evaporated; now the risk is back. - **Canada formally announced retaliatory tariffs** — escalation from the previous brief's "talks collapsed." - **Meta settled for \$16.7B** — a major legal resolution that removed uncertainty. - **OpenAI's Jalapeño chip was revealed** — a new threat to Nvidia's dominance. - **Multiple new open-weight models released** — Qwen3.8-Flash-Next, GLM-5.3-Flash, Laguna S 2.1 — continuing the price war. - **Goldman Sachs forecast** on China's chip gap narrowing, but TSMC still far ahead. - **Dolly Parton died** — a cultural icon, not a market mover, but notable.

11. Sources

World & Geopolitics: - CNBC — Iran says the U.S. is standing in the way of Hormuz deal amid talks with Oman - Al Jazeera — Trump tells Al Jazeera "not in a hurry" for Iran to return to talks - NPR — Canada to hit U.S. with retaliatory tariffs - The Guardian — Wiped out: US faces surging toilet paper prices amid trade war with Canada - UATV — Ukraine, US, And Europe Have Developed New Ceasefire Plan With Buffer Zone - UATV — More Than 50 UN Countries Call On Russia For Complete Ceasefire In Ukraine - BBC — China advises its citizens to leave Taiwan ally Eswatini over 'security risks' - France 24 — Hundreds missing, villages destroyed in massive flash floods near Nepal-Tibet border - NPR — 8 dead, hundreds missing as Nepal avalanche triggers flash floods - CNBC — Meta settles social media addiction case with California, other states for \$16.7 billion - Reuters — Meta reaches \$16.68B settlement over social media harms to children - BBC — Door was locked at Pakistan hospital where fire killed 14 babies, witnesses say - Al Jazeera — Locked inside: Grief and anger as Pakistan hospital fire kills 14 newborns - Al Jazeera — Bodies of missing people found in Korea becomes national scandal - TechCrunch — CISA confirms hackers targeted over 100 US water systems during July - Al Jazeera — Why northern Mozambique's surging violence is worsening displacement crisis

Markets, Money & Deals: - TradingView — Stock Market Today: Dow Jones Futures Rise, S&P 500 Slips Ahead of NVDA Q2 Earnings - CNBC — Fed's preferred inflation gauge shows core prices rose 3.3% annually in July - Investing.com — Nvidia: The Quarter May Be Priced In but the Forward Number Is the Trade - CNBC — Ray Dalio says Bessent move is sign that a debt crisis is getting closer; recommends gold and bitcoin - CNBC — Stanley Druckenmiller leads doubters who think Bessent's bond ploys will fail - CNBC — The 'debasement trade' returns after Bessent bond maneuver. Crypto and gold are back in style - CNBC — A massive trade just happened in gold. The options market is buzzing - CNBC — Dick's Sporting Goods stock falls 30% as retailer misses expectations - MarketWatch — Customers are fleeing TurboTax over price, and Intuit's stock is sliding - CNBC — Haidilao shares jump as delivery growth and new restaurant brands boost outlook - CNBC — Xpeng shares sink as weak delivery forecast overshadows \$6.3 billion robot unit valuation

AI & Infrastructure: - Tom's Hardware — Hot Chips 2026: Nvidia touts benefits of its DSX MaxLPS site power management approach - CNBC — OpenAI's Jalapeño AI chip brings new 'threat' to Nvidia margins as custom silicon gains ground - Bloomberg — Alibaba Releases Smaller, Cost-Effective Qwen AI Model - Hacker News — GLM-5.3-Flash -

TechCrunch — Surprise: Z.ai is the AI lab behind the mysterious Ox Alpha model - HPCwire — Poolside launches Laguna S 2.1 open-weight coding model - Thomson Reuters — Thomson Reuters Launches Its Own Frontier Model - CHOSUNBIZ — Goldman Sachs forecasts China to meet 66% of advanced chip demand by 2035 - BigGo Finance — Goldman Sachs: China's 7nm supply shortfall to shrink from 92% to 34% by 2035

Model & Research Watch: - Reddit — Megathread: Qwen3.8-Flash-Next - Release Day - Reddit — GLM-5.3-Flash: Frontier Intelligence, Flash Cost - Hugging Face Blog — IBM Granite 4.2 LLMs: How They're Built - Hugging Face Blog — Quantization-Aware Healing: a compressed, 4-bit model that outperforms its full-precision original - arXiv — RENDER: Controlling Reader-Facing Evidence in LLM Memory Evaluation - arXiv —

ESQ-Bench: A Multi-Tier Enterprise Oracle Benchmark for Evaluating NL2SQL - Hacker News — LLMs could control their host machines by exploiting inference engines - Ars Technica — 15,000 years of solar wind in four hours? A Seattle company gives it a go. - Ars Technica — Psychiatry's diagnostic bible is finally getting a biological upgrade - The Guardian — Mpox is back – and in new countries. How can this outbreak be contained?

What Builders Are Actually Using: - Hugging Face Blog — Wire It, Run It, Deploy It: AI Workflows in Gradio - Hugging Face Blog — Training and Finetuning Multi-Vector Embedding Models with Sentence Transformers

Quick Check

1. What is the core PCE inflation rate for July 2026, and why does it matter for the Fed?
2. Name two open-weight AI models released today that are competitive with top proprietary models.
3. What is the "debasement trade," and why has it returned to the news?

Answers: 1. Core PCE rose 3.3% annually in July, below the expected 3.6%. This is the Fed's preferred inflation gauge, and a lower reading gives the Fed room to hold rates steady rather than hike, which is positive

for stocks. 2. Alibaba's Qwen3.8-Flash-Next (125B parameters, competitive with Anthropic Opus 4.6) and Z.ai's GLM-5.3-Flash (frontier intelligence, flash cost). Also Poolside's Laguna S 2.1 for coding. 3. The debasement trade is the strategy of buying gold and bitcoin as a hedge against the devaluation of currency due to government debt and money printing. It returned to the news because Treasury Secretary Bessent's bond buyback intervention raised concerns about U.S. fiscal sustainability, leading investors to seek hard assets.

Generated 2026-08-26 08:04 PDT · 130 sources · model: deepseek-ai/DeepSeek-V4-Flash via <https://api.gmi-serving.com/v1>