

LAKSH'S DAILY BRIEF

The Laksh Daily Brief — August 28, 2026

Friday, August 28, 2026

World · Markets & Money · AI & Infrastructure · Models & Research



Guardian World — Egyptian queen's 673-diamond necklace stolen in Vienna smash-and-grab



Guardian World — Nepal-Tibet flash flood: EU pledges €2m in aid as fears rise for missing amid rescue effort



Al Jazeera — Oneness or Hindu supremacy? RSS chief Bhagwat's secretive New York hosts

1. World & Geopolitics

Fed Chair Warsh warns on inflation, signals caution on rate cuts at Jackson Hole — the market's main event today

Federal Reserve Chairman Kevin Warsh used his keynote speech at the Jackson Hole Economic Symposium to express concern about inflation and advocate for a "quieter" central bank that avoids giving too much guidance to markets ([CNBC](#)). He said the Fed may "have work to do" on inflation, which investors interpreted as a sign that rate cuts are not coming soon. The 2-year Treasury yield (a proxy for short-term interest rate expectations) jumped immediately after the speech ([CNBC](#)). Warsh stopped short of signalling a rate hike, but his tone was more hawkish (favouring higher rates to fight inflation) than many hoped. The speech came alongside a call from Cleveland Fed President Beth Hammack, who said "now is the time to act" on raising rates ([CNBC](#)). The Jackson Hole Symposium is the Fed's annual summer retreat where it signals the direction of monetary policy (the central bank's management of interest rates and money supply). Warsh's cautious stance cools expectations for a September rate cut.

US shifts to economic warfare against Iran, slashing oil exports with a naval blockade

The US Navy has imposed a blockade on Iranian oil exports, cutting Tehran's ability to sell crude on global markets ([CNBC](#)). This is a shift from the earlier phase of the conflict, which involved a dozen waves of airstrikes in July that failed to force Iran to abandon its claim to the

Strait of Hormuz. Now the US is using economic warfare — cutting off Iran's main revenue source. Meanwhile, another tanker attack was reported in the Strait of Hormuz, and regional powers are seeking an "off-ramp" (a way to de-escalate) from the war ([Greenwich Time](#)). Iran dismissed the latest US sanctions and called for a global pushback against them, saying diplomacy "isn't impossible" ([CNBC](#)). The Strait of Hormuz is a narrow waterway between the Persian Gulf and the Gulf of Oman; about 20% of the world's oil passes through it. A blockade or attacks there can spike global oil prices.

Nepal-Tibet flash floods: over 500 dead, 1,500 missing, rescue efforts hampered by terrain

The death toll from Wednesday's catastrophic flash floods on the Nepal-Tibet border has reached 579 in Nepal and 5 in Tibet, with more than 1,500 people still missing ([NPR](#)). The floods swept away roads, bridges, and entire villages. Rescue teams are struggling with the remote, mountainous terrain and the risk of further flooding from debris-dammed lakes ([France 24](#)). The EU has pledged €2 million in aid ([The Guardian](#)). This is a humanitarian disaster unfolding in a geopolitically sensitive border region between China and India.

China calls for "overcoming obstacles" with US ahead of Trump-Xi summit

Chinese Foreign Minister Wang Yi met with US Ambassador to China David Perdue and said the two countries should "overcome obstacles" between them ahead of an expected summit between President Trump and President Xi Jinping next month ([San Francisco](#)

[Chronicle](#)). The summit is tentatively set for September. Meanwhile, the Shanghai Cooperation Organization (SCO) — a security bloc including China, Russia, India, Pakistan, and Central Asian states — will hold a summit in Kyrgyzstan, where Xi, Modi, and Putin are expected to meet ([Anadolu Ajansi](#)). The SCO is a platform for China and Russia to counter US influence in Asia.

Russia escalates war in Ukraine, but September talks possible

Russia is preparing to escalate attacks on Ukraine after concluding that peace talks have reached a dead end, according to sources close to the Kremlin ([Farm Progress](#)). However, Ukrainian President Zelensky's top aide said that talks between Ukrainian and Russian teams, possibly including US representatives, could be held in September ([New York Post](#)). The US has been pushing for a ceasefire, but Russia's demands remain maximalist.

Gaza: Israeli strikes kill 4 including 2 children, a week after US request to curb attacks

Israeli airstrikes in Gaza killed four people, including two children, according to Palestinian medics ([Arab News](#)). This comes a week after the US asked Israel to reduce

civilian casualties. The war continues with no end in sight.

UK has billions in contracts with firms tied to illegal Israeli settlements

An Al Jazeera investigation found that UK public money is flowing to companies linked to Israeli settlements in occupied Palestinian territory, which are illegal under international law ([Al Jazeera](#)). The report says the contracts are worth billions of pounds.

Other briefs: - Spain's government called for calm in Ceuta after anti-migrant protests turned violent ([France 24](#)). - Brussels is seeing a surge in shootings as drug gangs torment the capital ([BBC](#)). - Norway's king died, causing the Oslo mayor to call off her wedding ([BBC](#)). - Security researchers found surveillance implants in Chinese-made routers sold worldwide ([Tom's Hardware](#)). - Authorities arrested two alleged members of the prolific hacking group TeamPCP, which infected over 1,000 organizations ([Ars Technica](#)). - Civil rights groups made an urgent plea for voting rights in DC ahead of the midterm elections ([The Washington Post](#)). A judge temporarily blocked new Postal Service rules that would restrict mail-in voting ([CBS News](#)).

2. Markets, Money & Deals

Warsh speech hits bonds: 2-year yield jumps, stocks flat

The 2-year Treasury yield rose after Warsh's cautious tone, as investors reduced bets on near-term rate cuts ([CNBC](#)). S&P 500 futures were little changed, with the market in wait-and-see mode ([CNBC](#)). The yield on the 2-year note is sensitive to Fed policy expectations; a higher yield means the market expects rates to stay higher for longer.

Marvell (MRVL) shares tumble 6% as outlook underwhelms despite 37% revenue growth

Marvell Technology reported fiscal second-quarter revenue that beat expectations, but its fiscal 2028 outlook failed to meet investors' high hopes ([CNBC](#)). Revenue grew 37% year-over-year, driven by strong demand for its data center and networking chips (used in AI infrastructure). However, guidance (the company's forecast for future earnings) was "muted" — management's tone suggested growth may slow. The

stock fell 6% in premarket trading. This is a classic "sell the news" reaction: the company beat the quarter but the future outlook disappointed investors who had already priced in perfection.

PayPal (PYPL) takeover bid falls through, stock plunges

Advent International and Stripe have abandoned their \$50 billion pursuit of PayPal, according to reports ([Bloomberg](#)). The deal was rumored for months and had propped up PayPal's stock. Now the company may have to execute a turnaround on its own, which is a harder sell ([MarketWatch](#)). PayPal's core business is under pressure from competitors like Block (Square) and Apple Pay.

Nvidia (NVDA) adds \$400 billion in market value after blowout earnings

Nvidia shares soared on Thursday, adding over \$400 billion in market capitalization (the total value of all its

shares) after its quarterly earnings crushed expectations and revenue guidance reassured investors that AI demand remains unbelievably strong ([CNBC](#)). Revenue doubled to \$96.2 billion. The stock is now up ~170% for the year. Nvidia's results are the single most important event for the AI infrastructure trade.

Salesforce (CRM) surges on AI strength, Marc Benioff gets his mojo back

Salesforce's stock had its second-best day on record after earnings, as Wall Street showed renewed confidence in CEO Marc Benioff's AI story ([CNBC](#)). The company reported strong growth in its AI-powered products, particularly the Agentforce platform. This is a sign that enterprise AI adoption is becoming real for many companies beyond just the hyperscalers.

CrowdStrike (CRWD) and Okta (OKTA) surge on cybersecurity earnings

CrowdStrike posted its best day ever, and Okta's stock popped nearly 29% after both companies reported strong earnings driven by rising AI threat ([CNBC](#)). As AI adoption grows, so does the need for cybersecurity — both to protect AI systems and because AI helps attackers create more sophisticated threats. This is a "picks and shovels" play in the AI ecosystem.

IREN (IREN) falls 6% on \$639 million mining rig writedown

IREN (Iris Energy) reported a \$639 million writedown on its mining rigs, overshadowing growth in its AI cloud business ([24/7 Wall St.](#)). A writedown is a reduction in the reported value of an asset because its market price has fallen. IREN's Bitcoin mining rigs became less valuable as the crypto mining landscape changed.

IMAX (IMAX) says it's open to a sale, but no buyers appear

IMAX's stock and box office are breaking records, but studio conflicts and a soaring valuation complicate the list of potential buyers ([CNBC](#)). The company has put itself on the market, but no credible bidder has emerged. This is a classic "we're open to offers" statement that often tests the market.

Other deal news: - Astrum Space, a direct-to-device satellite venture, announced plans to merge with a SPAC (a blank-check company) in a \$1 billion deal ([SpaceNews](#)). - JERA, a Japanese energy company, is backing a startup developing solutions for grid-to-data center projects ([Rigzone](#)). - Coinbase (COIN) fell on reports that banks are planning stablecoin issuance ahead of the CLARITY Act, a bill that would regulate stablecoins ([Cryptonews](#)). - Bank of America warns that the bull market for stocks may face an "autumn reality check" from midterm elections and a potential end to the Iran war ([MarketWatch](#)).

3. AI & Infrastructure

Amazon to buy 2 million more Nvidia GPUs, tripling its orders

Amazon (AMZN) is planning to buy an additional 2 million Nvidia GPUs, according to reports ([Telecoms](#)). This triples its previous orders and is part of a plan to spend \$220 billion in capital expenditure (capex) on data centers and other long-lived assets over the next few years. The move underscores the hyperscaler's (a giant cloud computing company like Amazon, Microsoft, Google) commitment to AI infrastructure. For context, a single Nvidia H100 GPU costs about \$30,000; 2 million GPUs would cost roughly \$60 billion. This is a massive bet that AI demand will continue to grow.

a16z creates a \$1.1 billion "Machine Age" fund to accelerate the physical buildout of AI

Andreessen Horowitz (a16z), the iconic venture capital firm known for software investments, launched a \$1.1 billion fund dedicated to "the physical buildout of AI" ([TechCrunch](#)). This includes data centers, power infrastructure, cooling systems, and hardware. The fund signals that the next phase of AI investment is shifting from software to physical infrastructure — the "picks and shovels" of the AI era.

EPA says power for data centers can sidestep pollution laws

The Environmental Protection Agency (EPA) issued new permitting guidance that allows data centers to bypass certain pollution laws, as part of the Trump administration's agenda to promote AI infrastructure ([Hacker News](#)). This is a major win for the AI buildout, as data centers are enormous consumers of electricity (a

single large data center can use as much power as a small city). The guidance effectively eases environmental review for new data center construction, which could accelerate the pace of building.

SK Hynix to start HBM4E production at Indiana plant in 2029

SK Hynix, the world's second-largest memory chipmaker (after Samsung), announced that it will begin mass production of its next-generation HBM4E memory at its new advanced packaging plant in Indiana in 2029 ([The Elec](#)). HBM (High-Bandwidth Memory) is a type of memory that sits next to the GPU in AI accelerators, providing the fast data access needed for training and inference. The \$4 billion investment is SK Hynix's first U.S. advanced HBM packaging base. HBM4E is the next generation after HBM4, expected to offer even higher bandwidth and lower power consumption. This is important for the AI supply chain because it diversifies HBM manufacturing away from South Korea and into the U.S.

JERA backs startup for grid-to-data center projects

JERA, Japan's largest power generation company, is backing a startup that develops solutions for connecting data centers directly to the power grid ([Rigzone](#)). This is part of a broader trend: as AI data centers consume

more power, energy companies are looking for ways to provide reliable, dedicated electricity. "Grid interconnection" — the process of connecting a data center to the power grid — is a major bottleneck for new data center projects.

AI's "capex tsunami" is fueling a potential credit boom

A Yahoo Finance article asks whether the massive wave of capital expenditure (capex) from AI companies is the catalyst for the next big credit boom ([Yahoo Finance Canada](#)). The idea: companies are borrowing heavily to fund AI data center buildouts, which creates demand for credit. If that borrowing is done responsibly, it could stimulate economic growth. But if it's overheated, it could lead to a credit bubble. This is a second-order effect worth watching.

Anthropic gets its first court win over the Pentagon's supply chain risk label

Anthropic, the AI company behind Claude, won a court case challenging the Pentagon's designation of its technology as a supply chain risk ([TechCrunch](#)). This is a regulatory win for Anthropic, which is trying to sell its AI models to the U.S. government. The Pentagon had labeled Anthropic a risk, potentially blocking government contracts; the court ruled that the label was not justified.

4. Model & Research Watch

Zhipu AI's "Ox Alpha" revealed as GLM-5.3-Flash — open-weight model ran on 100,000 Chinese chips

The mysterious "Ox Alpha" model that went viral on OpenRouter last week has been confirmed as GLM-5.3-Flash, an open-weight model from Chinese AI company Zhipu AI ([CNET](#)). The model ran entirely on a cluster of 100,000 domestically produced Chinese chips, setting a usage record and demonstrating China's ability to scale AI on domestic hardware despite US export controls ([SCMP](#)). The claim is that the model served 62 trillion tokens (units of text output) at frontier scale ([Tech Times](#)). However, this claim remains unverified by independent analysts ([Tech Times](#)). The model is open-weight (weights are publicly available) and is a Flash variant, meaning it's optimized for inference speed. This is a significant milestone in China's push for "Sovereign

AI" (having independent AI capabilities not reliant on US technology).

Tencent drops Hy4-preview 770B-A49B — a massive Mixture-of-Experts model

Tencent released the weights for its Hy4-preview model, which has 770 billion total parameters but only 49 billion active parameters per token ([Reddit r/LocalLLaMA](#)). This is a Mixture-of-Experts (MoE) architecture: the model has many "expert" sub-networks, but only a few are activated for each input, making it efficient to run despite its massive size. 49B active parameters puts it in the same ballpark as GPT-4 class models. The model is open-weight, meaning anyone can download and run it (with enough hardware). Tencent claims it outperforms rivals on some benchmarks ([news24.com](#)).

Meta Muse Glimmer: 30B open model for consumer GPUs

Meta released a new open model called "Muse Glimmer," a 30-billion-parameter model designed to run on consumer GPUs (tech-insider.org). This is aimed at the local AI community: people who want to run powerful models on their own hardware without paying API fees. 30B parameters is a sweet spot — it's powerful enough for many tasks but can run on a single high-end GPU like an Nvidia RTX 4090 (with quantization). This is part of Meta's strategy of open-sourcing models to challenge OpenAI and Anthropic.

Qwen3.8-Flash matches DeepSeek V4 Pro on coding benchmarks at a quarter of the price

Alibaba's Qwen team released Qwen3.8-Flash, a small model (3.8B parameters) that matches DeepSeek V4 Pro on coding benchmarks but costs only a quarter as much to run (Intelligent Living). This is a big deal for cost-conscious developers: they can now get frontier-level coding ability at a fraction of the cost. The "Flash" variant is optimized for speed, likely using quantization (reducing the precision of the model's weights to make it faster and smaller) or other inference optimizations.

Claude Sonnet 5 vs GPT-5.6 Sol vs Kimi K3: a \$20 gap?

A comparison article looks at the three leading AI models — Anthropic's Claude Sonnet 5, OpenAI's GPT-5.6 Sol, and Moonshot AI's Kimi K3 — and notes that the price gap between them is about \$20 per million tokens (tech-insider.org). This shows the market is becoming commoditized, with multiple frontier models competing on price and capability. For users, this is good: it means lower prices and more choice.

Anthropic previews Model Hardware Standard — a proposal for AI hardware interoperability

Anthropic published a research preview on a "Model Hardware Standard" — a proposed specification that would allow AI models to run on different hardware platforms without modification (Anthropic blog). This is a response to the problem of hardware lock-in: currently, Nvidia's CUDA platform dominates AI, making it hard to run models on other chips. A hardware standard could open up the market to more competition, which is good for Anthropic's goal of making AI more accessible.

Other research papers from today: - "NeuronFuzz: Safety Neuron Guided Fuzzing for LLM Safety Evaluation" — a new method to test LLM safety by targeting specific neurons (arXiv). - "PICasso: An AI-Enabled Design Framework for Autonomous Optimization of Silicon Photonic Devices" — using AI to design photonic chips for data centers (arXiv). - "The Artificial Experimentalist: Discovery and Control of Self-Organizing Phenomena with Autotelic Reinforcement Learning" — an AI agent that autonomously explores complex systems like cellular automata (arXiv). - "CIFQA: A Deterministic Tool-Grounded Multi-Agent LLM Framework for Financial Query Answering" — an AI system for accurate financial calculations (arXiv).

Science & Frontier Tech: NASA's Roman Space Telescope is ready for launch

NASA's Nancy Grace Roman Space Telescope (formerly WFIRST) is set to launch soon. It will see 100 times more sky than the Hubble Space Telescope and has a 300-megapixel camera (Science Daily). The telescope was built using a repurposed spy satellite from the National Reconnaissance Office, which gives it a large mirror and wide field of view (ABC News). It will study dark energy, dark matter, and exoplanets. The telescope is named after Nancy Grace Roman, NASA's first chief astronomer and the "mother of Hubble."

5. What Builders Are Actually Using

Concrete swap: Qwen3.8-Flash can replace DeepSeek V4 Pro for coding tasks at a quarter of the cost

Alibaba's Qwen3.8-Flash (3.8 billion parameters) matches DeepSeek V4 Pro on coding benchmarks but costs about 75% less to run (Intelligent Living). Rough pricing: Qwen3.8-Flash via API is about \$0.15 per million

tokens, while DeepSeek V4 Pro is about \$0.60 per million tokens. For self-hosting, Qwen3.8-Flash can run on a single consumer GPU (like an RTX 4090 with quantization), while DeepSeek V4 Pro needs multiple high-end GPUs. The capability gap in coding is negligible for most practical tasks; the tradeoff is that Qwen3.8-Flash may be weaker on very complex

reasoning or long-context tasks. For a solo developer or small team, this is a clear upgrade on cost.

OpenRouter alternative: a Rust-native open-source model gateway

A new open-source project called "Experiential" is an open-source model gateway that allows developers to manage their own self-hosted, frontier, and open-source models in one place ([Hacker News](#)). It's written in Rust for performance, adds under 1 millisecond of latency for bring-your-own-key (BYOK) requests, and handles the quirks of different providers (streaming formats, tool calls, rate limits, error behavior). This is useful for developers who want to avoid vendor lock-in with OpenRouter or other centralized gateways.

The local AI stack for productive SLMs — a practical guide

A KDnuggets article outlines the "Local AI Stack" for productive Small Language Models (SLMs) — models that run on consumer hardware ([KDnuggets](#)). The stack includes quantized models, inference engines like llama.cpp, and tools like Ollama for easy deployment. This is a growing trend: as models get smaller and more efficient, running AI locally becomes practical for more users.

Meta's Muse Glimmer (30B) is now available for consumer GPUs

Meta released Muse Glimmer, a 30-billion-parameter open model that can run on consumer GPUs ([tech-insider.org](#)). With 4-bit quantization (reducing the model's precision to save memory), it can fit on a single 24GB GPU like the RTX 4090. This is a strong alternative to running GPT-4 or Claude via API — you get full control, privacy, and no per-token costs, but you need the hardware.

6. Watchlist: Earnings, Guidance & Movers

Ticker	Company	Notable News	Price Change (vs close)
MRVL	Marvell Technology	Revenue beat but guidance underwhelmed; stock fell 6%	-7.82%
NVDA	Nvidia	Added \$400B in market cap after earnings; stock up 7.6% on Thursday	-0.53% (pullback)
CRWD	CrowdStrike	Best day ever on earnings; stock up ~15%	Not in watchlist snapshot
OKTA	Okta	Stock popped 29% on earnings	Not in watchlist snapshot
CRM	Salesforce	Second-best day on record on AI strength	Not in watchlist snapshot
IREN	IREN	Fell 6% on \$639M mining rig writedown	Not in watchlist snapshot
COIN	Coinbase	Fell 2% on bank stablecoin plans	-2.01%
AMZN	Amazon	Up 3.38% on GPU order news; Amazon to buy 2M more Nvidia GPUs	+3.38%
IMAX	IMAX	Said open to sale; no buyers yet	-1.39%
PYPL	PayPal	Failed takeover bid, stock plunging	Not in watchlist snapshot
META	Meta	Up 2.71% on AI momentum	+2.71%
GOOGL	Alphabet	Up 1.79%	+1.79%
MSFT	Microsoft	Up 1.91%	+1.91%
AVGO	Broadcom	Up 0.57%	+0.57%
PLTR	Palantir	Up 0.36%	+0.36%
MU	Micron	Up 0.46%	+0.46%
TSM	TSMC	Down 0.07%	-0.07%
CRDO	Credo	Down 0.80%	-0.80%
ALAB	Astera Labs	Down 2.54%	-2.54%
COHR	Coherent	Down 2.65%	-2.65%
AAOI	Applied Optoelectronics	Down 3.85%	-3.85%
LITE	Lumentum	Down 3.36%	-3.36%
FLNC	Fluence Energy	Down 2.49%	-2.49%
VST	Vistra	Down 0.29%	-0.29%
CEG	Constellation Energy	Down 0.33%	-0.33%
DLR	Digital Realty	Down 0.82%	-0.82%
EQIX	Equinix	Down 0.78%	-0.78%
ARM	Arm Holdings	Down 2.45%	-2.45%
RGTI	Rigetti Computing	Down 3.28%	-3.28%
SOFI	SoFi	Down 2.01%	-2.01%
TSLA	Tesla	Down 0.20%	-0.20%
NFLX	Netflix	Up 3.02%	+3.02%
NKE	Nike	Up 1.86%	+1.86%
AMD	AMD	Up 0.33%	+0.33%
KLAC	KLA Corp	Down 1.74%	-1.74%
QCOM	Qualcomm	Down 0.22%	-0.22%

Ticker	Company	Notable News	Price Change (vs close)
INTC	Intel	Down 0.08%	-0.08%
NET	Cloudflare	Down 1.22%	-1.22%
HOOD	Robinhood	Down 0.40%	-0.40%
CMCSA	Comcast	Up 1.84%	+1.84%
CNK	Cinemark	Down 0.44%	-0.44%
IMAX	Imax	Down 1.39%	-1.39%
SBUX	Starbucks	Up 0.46%	+0.46%
FXAIX	Fidelity 500 Index	Up 0.72%	+0.72%

Companies with notable earnings or news in the last 24 hours:

- **Marvell (MRVL)**: Revenue beat expectations, but fiscal 2028 outlook was muted. Stock fell 6% ([CNBC](#)).
- **Nvidia (NVDA)**: Added over \$400 billion in market cap after blowout earnings; stock up 7.6% on Thursday, pulling back slightly today ([CNBC](#)).
- **CrowdStrike (CRWD)**: Fiscal Q2 earnings beat estimates; stock posted its best day ever, up ~15% ([CNBC](#)).
- **Okta (OKTA)**: Revenue beat and strong guidance; stock surged nearly 29% ([CNBC](#)).
- **Salesforce (CRM)**: Strong earnings on AI products; stock had its second-best day on record ([CNBC](#)).

- **IREN (IREN)**: Reported \$639M writedown on mining rigs; stock fell 6% ([24/7 Wall St.](#)).
- **PayPal (PYPL)**: Takeover bid from Advent and Stripe abandoned; stock plunging ([Bloomberg](#)).

No notable news in the sources for the following watchlist names today: SOFI, RGTI, FLNC, RDW, INFQ, SNDK, AMC, NET, HOOD, ARM, CNK, SBUX, NKE, FXAIX, SPXC, SPCX, FOTO, INFQ, AAOI (no earnings news, but stock moved on broader AI selloff). Also no news on IBM, KLAC, QCOM, INTC, AMD, DLR, EQIX, CEG, VST, TSLA, GOOGL, META, MSFT, AMZN (beyond general AI news), AVGO, PLTR, MU, TSM, CRDO, ALAB, COHR, LITE, NEE.

7. Analysis: Second-Order Effects

Chain 1: Iran oil blockade -> higher oil prices -> headline inflation -> Fed stays hawkish -> data center borrowing costs rise

- **FACTS**: The US Navy has imposed a blockade on Iranian oil exports, cutting off a major source of global supply ([CNBC](#)). Another tanker attack was reported in the Strait of Hormuz ([Greenwich Time](#)). Brent crude oil prices (the global benchmark) are likely to rise as supply tightens. Fed Chair Warsh expressed concern about inflation and said the Fed may "have work to do" ([CNBC](#)).
- **INFERENCE**: If oil prices spike, headline inflation (which includes energy costs) will rise. The Fed targets core PCE (which excludes food and energy), but headline inflation still affects consumer and business sentiment. Even if the Fed does not react directly, higher oil prices raise costs for the entire

economy — including for data center operators, who consume enormous amounts of electricity. Since many data centers run on diesel generators for backup, diesel costs also rise. This would increase the operating expenses (opex) of AI infrastructure.

- **WEAKEST LINK**: The Fed's actual policy response. The Fed targets core PCE, which excludes energy. So a spike in oil prices alone does not force the Fed to raise rates. However, if oil prices stay high for months, they can feed into core inflation through higher transportation costs and production inputs. The chain depends on oil prices staying elevated long enough to affect core inflation.
- **WHAT WOULD FALSIFY IT**: If oil prices spike but then quickly fall back (as happened after the 2019 Abqaiq attack on Saudi oil facilities, where Brent rose ~15% in a day then fully retraced within weeks), or if

the Fed explicitly says it will look through energy price spikes.

Chain 2: Amazon tripling GPU orders -> Nvidia supply tightens -> HBM demand surges -> SK Hynix and Micron benefit -> power demand for data centers explodes -> grid interconnection bottlenecks worsen

- **FACTS:** Amazon plans to buy 2 million more Nvidia GPUs ([Telecoms](#)). Nvidia's latest earnings showed revenue doubling and guidance that demand is still outstripping supply ([CNBC](#)). SK Hynix is building a \$4 billion HBM4E packaging plant in Indiana ([The Elec](#)). The EPA issued guidance to ease pollution rules for data center power ([Hacker News](#)).
- **INFERENCE:** Even more GPU orders mean Nvidia's production capacity will be stretched, keeping prices high. Every GPU needs HBM memory (High-Bandwidth Memory) — the more GPUs, the more HBM, which benefits SK Hynix, Samsung, and Micron. Data centers that house these GPUs need enormous amounts of electricity. The EPA guidance helps, but grid interconnection (the process of connecting to the power grid) remains a bottleneck that takes years. This could lead to a "power gap" where data center construction outpaces the ability to supply electricity.
- **WEAKEST LINK:** The grid interconnection timeline. The EPA guidance is a policy change, but actual interconnection requires new transmission lines, which face NIMBY opposition and regulatory hurdles. The chain assumes that easing pollution rules will accelerate interconnection, but it may not be enough.
- **WHAT WOULD FALSIFY IT:** If data center construction slows due to power shortages being worse than expected, or if alternative power sources (like small modular nuclear reactors) come online faster than grid connections.

Chain 3: China's domestic chip breakthrough (GLM-5.3-Flash on 100,000 Chinese chips) -> US export controls lose effectiveness -> US-China tech war enters new phase -> US may tighten controls further -> global chip supply chain disruption

- **FACTS:** Zhipu AI's GLM-5.3-Flash model ran on 100,000 domestically produced Chinese chips, serving 62 trillion tokens ([Tech Times](#)). The claim is that China's chip supply chain is becoming "safer"

through self-sufficiency ([SCMP](#)). However, the claim is unverified ([Tech Times](#)). The US has been tightening export controls on advanced chips to China, and a recent chip smuggling case in Taiwan showed that enforcement is improving ([FDD](#)).

- **INFERENCE:** If China can currently train and run frontier-scale AI models on domestic chips, the US export control strategy of cutting off advanced chips may be failing. This would push the US to adopt even more aggressive controls — perhaps targeting the Chinese chip manufacturing ecosystem itself (e.g., restricting equipment sales to SMIC, China's main chipmaker). It could also accelerate the US push for "Sovereign AI" in allied countries, and fuel more subsidies for domestic chip production (TSMC's Arizona plant, Intel's foundry plans).
- **WEAKEST LINK:** The claim that the chips are truly domestic. It's possible that the "100,000 Chinese chips" include some that rely on US technology or were smuggled. The claim is unverified. If the chips are actually repurposed designs that still depend on US intellectual property, the export controls are still somewhat effective.
- **WHAT WOULD FALSIFY IT:** Independent verification that the chips are indeed fully domestic and competitive with Nvidia's H100/B200 on performance. Or, conversely, if US export controls are tightened further and China's domestic progress stalls.

Chain 4 (SPECULATIVE): Warsh's hawkish tone + Hamas not backing down -> Israel-Iran escalation remains possible -> oil risk premium stays elevated -> suppliers of diesel generators and backup power benefit

- **FACTS:** Warsh expressed concern about inflation and the 2-year yield jumped ([CNBC](#)). The Strait of Hormuz remains a flashpoint with tanker attacks ([Greenwich Time](#)). Regional powers are seeking an off-ramp, but no deal is in place yet.
- **INFERENCE:** If the central bank is worried about inflation, it will keep rates higher for longer. Higher rates make it more expensive to finance data center construction. Meanwhile, the risk of further escalation in the Middle East keeps oil prices elevated, which directly impacts data center operating costs (via electricity and diesel). This could compress margins for AI infrastructure companies that haven't locked in

long-term power purchase agreements (PPAs). Companies that sell generators, backup power, and energy efficiency solutions (like Fluence, Generac, Eaton) could benefit from the dual pressures of rising power demand and energy price volatility.

- **WEAKEST LINK:** The oil price reaction. Markets may have already priced in the current level of risk in the

Strait of Hormuz. A further escalation is needed to move oil prices significantly. Without a new attack, the risk premium may stay flat.

- **WHAT WOULD FALSIFY IT:** A rapid de-escalation in the Iran conflict, or a ceasefire that reduces the risk premium. Also, if the Fed signals rate cuts despite oil prices, the whole chain breaks.

8. Building Your Knowledge

How the Jackson Hole speech moves markets: The Federal Reserve's annual symposium in Jackson Hole, Wyoming, is a key event where the Fed chair often signals major policy shifts. Markets watch every word for hints about future rate decisions. A hawkish tone (worried about inflation) pushes bond yields up and stocks down; a dovish tone (worried about the economy) does the opposite. This year, Warsh's speech was more hawkish than expected, which caused the 2-year Treasury yield to jump. The S&P 500 was flat, suggesting the stock market is less sensitive to rate moves than the bond market right now. This is a pattern: bonds react first, then stocks follow.

Why China's chip self-sufficiency is a geopolitical and investing milestone: The GLM-5.3-Flash model running on 100,000 domestic chips is a claimed breakthrough in China's ability to train and deploy AI without US hardware. If verified, it means the US export control strategy of cutting off advanced chips is not working as well as hoped. For investors, this has implications: Nvidia's monopoly on AI training could be challenged by Chinese competitors (like Huawei's Ascend chips), which would pressure Nvidia's margins. It also means the "TSMC loophole" (the idea that Taiwan makes all advanced chips) may be less critical if China can produce its own. However, the claim remains unverified, and Chinese chips are likely still less efficient than Nvidia's.

How Mixture-of-Experts (MoE) models work and why they matter for cost: Tencent's Hy4-preview has 770 billion total parameters but only 49 billion active per token. In an MoE model, the model has many "expert" subnetworks, but only a few experts are activated for each input. This means the model has huge knowledge capacity (700B+ parameters) but costs only as much to run as a much smaller model (49B parameters). This is

why MoE models like GPT-4, DeepSeek V4, and Tencent's Hy4 can offer frontier capability at reasonable inference costs. For builders, choosing an MoE model often gives you the best price-performance ratio.

The "sell the news" pattern in earnings: Marvell beat revenue expectations but gave a muted outlook, and the stock fell 6%. This is a classic "sell the news" reaction: when a company's earnings are very good but the future guidance is cautious, investors who bought the stock in anticipation of good news sell it after the news is out. The key is that the market had already priced in the good quarter; what matters is the future. Guidance matters more than the past.

What "economic warfare" means: The US is now using a naval blockade and sanctions to cut off Iran's oil exports, rather than airstrikes. This is economic warfare: using economic tools (blockades, sanctions, trade restrictions) instead of military force to achieve political goals. Economic warfare can be slower but is often less costly in lives and can be sustained for longer. The risk is that it can also hurt neutral countries and global markets.

How data center power is becoming a national security and investment issue: The EPA issued guidance to let data centers bypass some pollution laws, and a16z created a \$1.1B fund for AI physical infrastructure. This reflects a growing recognition that AI data centers need enormous amounts of power — and that the grid is not ready for it. The tension between "we need more power for AI" and "we need to reduce carbon emissions" is a major theme. Companies that provide power solutions (solar, nuclear, gas, batteries, grid interconnection services) are becoming essential to the AI buildout.

9. Foundations & Terms

Concept Spotlight: Liquidity and Market Makers

What it is in plain words: Liquidity is the ease with which you can buy or sell something without moving its price much. A "liquid" market is one where you can sell a large amount of something quickly at a fair price. An "illiquid" market is one where trying to sell a large amount causes the price to drop sharply.

Why it exists / what problem it solves: Markets need people who are willing to buy when others want to sell, and vice versa. Without liquidity, investors would be stuck with assets they can't sell. Market makers are financial firms (often banks or trading firms) that quote both a buy price (bid) and a sell price (ask) for a stock or other asset, committing to trade at those prices. They make money on the spread (the difference between bid and ask). They provide liquidity to the market.

A concrete worked example with real numbers: - You want to sell 10,000 shares of Apple (AAPL). The current bid is \$200.00, ask is \$200.05. The spread is 5 cents. - If you place a market order to sell 10,000 shares, you'll get filled at the bid price (\$200.00) for most of it, but the market maker may adjust the bid down as they absorb your shares. The price might drop to \$199.90 by the time your order is fully filled. You "moved the market" by 10 cents. - For a small stock with low liquidity, the same 10,000 shares might cause the price to drop from \$20.00 to \$18.00 — a much bigger move. You'd lose a lot more.

How it shows up in the news: When you see a stock "gapping" (opening much higher or lower than the previous close), it's often because of a liquidity gap — not enough buyers or sellers at the open. The VIX (fear gauge) measures volatility, which is related to liquidity: when liquidity dries up, prices move more. The Fed's policies affect liquidity: when the Fed tightens, it reduces the amount of money in the system, which can reduce market liquidity.

Rough benchmarks or figures worth remembering: - The S&P 500 is extremely liquid: you can trade billions of dollars with minimal price impact. - Small-cap stocks (market cap under \$2 billion) are much less liquid. A \$1 million trade can move the price 1-2%. - The bid-ask spread for a highly liquid stock like Apple is often 1-2 cents. For a penny stock, it can be 10 cents or more. - Market makers are required to provide liquidity in

exchange for certain privileges. On the NYSE, designated market makers (DMMs) have obligations to maintain fair and orderly markets.

A caution or common misconception: "Liquidity" is not the same as "volume." A stock can have high trading volume (many shares traded) but still be illiquid if the order book is thin (only a few shares at each price level). Conversely, a stock can have low volume but be liquid if the spreads are narrow. Also, liquidity can disappear in a crisis: during the 2008 financial crisis, even supposedly liquid assets like mortgage-backed securities became impossible to trade at any price.

Concept Spotlight: Food and Fertiliser Security

What it is in plain words: Food security means that a country has enough food to feed its population. Fertiliser security means that a country has reliable access to the chemical inputs (like nitrogen, phosphorus, potassium) needed to grow crops. Many countries are not self-sufficient in food or fertiliser; they rely on imports from a few large producers. When those imports are disrupted (by war, sanctions, weather, or shipping problems), it can cause food shortages, price spikes, and even political instability.

Why it exists / what problem it solves: The modern global food system is highly concentrated. A few countries (Ukraine, Russia, Brazil, US) produce most of the world's wheat, corn, and soybeans. A few countries (Russia, Belarus, China, Morocco) produce most of the world's fertiliser. This creates a vulnerability: if a major producer suffers a war or a bad harvest, or if shipping routes are blocked, prices can spike globally. The problem is that food is a basic necessity — people need to eat every day. When food prices rise sharply, it can cause protests, riots, and even government collapses (as in the 2011 Arab Spring, where wheat prices were a factor).

A concrete worked example with real numbers: - Ukraine and Russia together account for about 30% of global wheat exports. In 2022, after Russia invaded Ukraine, wheat prices doubled from \$7 to \$14 per bushel. - Russia is also the world's largest exporter of fertiliser (especially nitrogen, potash, and phosphates). When Russia's fertiliser exports were disrupted by sanctions, farmers in developing countries like Kenya

and Bangladesh could not afford to buy enough fertiliser, leading to lower crop yields and higher food prices. - The 2022 crisis led to food price inflation of 20-30% in many African countries, causing protests in Sri Lanka, Pakistan, and elsewhere.

How it shows up in the news: The Iran war is disrupting shipping in the Strait of Hormuz, which is a major route for fertiliser exports from the Gulf states. The Russia-Ukraine war directly affects wheat and fertiliser. Climate events like El Niño can cause crop failures in Australia, Brazil, or Southeast Asia. When you see "global food prices rise," it's often linked to fertiliser costs or shipping disruptions. The UN's Food and Agriculture Organization (FAO) tracks a Food Price Index that is a key indicator.

Rough benchmarks or figures worth remembering: - The world uses about 200 million tonnes of fertiliser per year. - Russia is the largest exporter of nitrogen fertiliser (about 15% of global trade) and the second-largest exporter of potash. - Ukraine is the fifth-largest exporter of wheat. - About 80% of the world's population lives in countries that are net importers of food. - The FAO Food Price Index doubled between 2020 and 2022.

A caution or common misconception: Food security is not just about calories — it's also about nutrition and affordability. Even if enough food exists globally, it may not reach the people who need it (due to poverty, logistics, or politics). Also, fertiliser security is often overlooked: people focus on wheat or grain prices, but fertiliser is the hidden input that makes modern agriculture possible. Without fertiliser, crop yields would fall dramatically.

Other Terms from Today's News

- **Economic warfare:** The use of economic tools (sanctions, blockades, tariffs, asset freezes) to weaken an adversary's economy and force it to change its behavior, as an alternative to or supplement to military force. Today's example: the US Navy blockade of Iranian oil exports ([CNBC](#)).
- **Off-ramp:** A diplomatic term for a way to de-escalate a conflict or crisis without losing face. Regional powers are seeking an off-ramp from the Iran war ([Greenwich Time](#)).
- **Writedown:** A reduction in the book value of an asset on a company's balance sheet because its market

value has fallen. IREN took a \$639 million writedown on its mining rigs ([24/7 Wall St.](#)).

- **Open-weight model:** An AI model whose weights (the trained parameters that define its behavior) are publicly released, allowing anyone to download, run, and modify it. Examples: Meta's Llama series, Zhipu's GLM-5.3-Flash, Tencent's Hy4-preview.
- **Mixture-of-Experts (MoE):** A neural network architecture where the model has many specialized sub-networks (experts) but only a few are activated for each input, allowing it to be very large (billions of parameters) while being efficient to run. Tencent's Hy4-preview has 770B total parameters but only 49B active per token ([Reddit](#)).
- **Active parameters:** In an MoE model, the subset of parameters that are actually used for a given input. The total number of parameters can be large, but the active parameters determine the cost of running the model.
- **Grid interconnection:** The process of connecting a new power consumer (like a data center) to the electrical grid. This involves building new transmission lines, installing transformers, and getting regulatory approvals. It is a major bottleneck for the AI buildout because it can take years ([Rigzone](#)).
- **Capex tsunami:** A term used to describe the massive wave of capital expenditure (investment in long-lived assets) that AI companies are undertaking. Amazon's \$220 billion plan is part of this. The concern is that this could lead to a credit boom or bubble ([Yahoo Finance Canada](#)).
- **SPAC (Special Purpose Acquisition Company):** A blank-check company that raises money through an IPO and then merges with a private company to take it public. Astrum Space is merging with a SPAC in a \$1 billion deal ([SpaceNews](#)).
- **Sell the news:** A market phenomenon where a stock rises in anticipation of good news (like earnings) and then falls after the news is announced, because the good news was already "priced in." Marvell's earnings beat was already expected, so the muted guidance caused a selloff ([CNBC](#)).
- **Guidance:** A company's forecast of its future financial performance. Often more important than the actual quarterly results because it shows management's confidence in the future. Marvell's guidance was

"muted," meaning the company hinted at slower growth ahead.

- **Tanker attack:** An attack on an oil tanker, often in the Strait of Hormuz, that threatens shipping and oil supply. Such attacks can cause war-risk insurance premiums to spike, raising shipping costs ([Greenwich Time](#)).
- **Export control:** Government regulations that restrict the export of certain technologies (like advanced

semiconductors) to other countries, usually for national security reasons. The US has export controls on AI chips to China. The Zhipu breakthrough challenges these controls ([FDD](#)).

- **Falsifiable:** A scientific term meaning that a claim can be proven false by observable evidence. In the analysis section, we used "falsifiable" to describe what would break a chain of reasoning. This is a key concept in epistemology: strong claims are falsifiable.

10. Bottom Line

The world is in a tense period: the Fed is worried about inflation and may not cut rates soon, the Iran war is shifting to economic warfare with a naval blockade, and the US-China tech war is at a critical juncture as China claims a breakthrough in domestic AI chips. For markets, the focus is on AI infrastructure spending, which is accelerating despite higher interest rates. Nvidia, Amazon, and a16z are all betting big on the physical buildout of AI, from GPUs to data centers to power. The practical AI landscape is also shifting: open-weight models from China (Tencent, Zhipu, Alibaba) are matching or challenging frontier models at a fraction of the cost, which is great for builders but could pressure margins at US AI companies.

What changed since the previous brief: - **Fed Chair Warsh gave a hawkish speech at Jackson Hole** — the previous brief had the speech as an upcoming event; now we have the content. He expressed concern about inflation and advocated for a quieter central bank, pushing the 2-year yield up. - **The US Navy imposed a blockade on Iranian oil exports** — a significant escalation from airstrikes to economic warfare. The

previous brief had the Iran war in a stalemate; now the US is trying a new strategy. - **Marvell reported earnings but gave a muted outlook** — the stock fell 6%. The previous brief had no Marvell earnings data. - **China's Zhipu AI confirmed GLM-5.3-Flash ran on domestic chips** — a major milestone that was previously just a rumor. The model is now open-weight. - **Tencent dropped an open-weight MoE model (770B total, 49B active)** — a new competitor in the open-source AI space. - **Amazon confirmed plans to buy 2 million more Nvidia GPUs** — a massive increase in AI infrastructure spending. - **a16z created a \$1.1B "Machine Age" fund** — a new fund specifically for physical AI infrastructure. - **The EPA issued guidance to ease pollution rules for data center power** — a policy win for the AI buildout. - **SK Hynix announced HBM4E production in Indiana in 2029** — a new step in the semiconductor supply chain. - **PayPal takeover bid fell through** — a major M&A story that was previously rumored; now it's dead. - **CrowdStrike and Okta surged on cybersecurity earnings** — a reminder that AI is also driving cybersecurity spending.

11. Sources

World & Geopolitics - CNBC — Fed Chairman Warsh expresses concern about inflation, advocates for 'quieter' central bank - CNBC — 2-year Treasury yield jumps as Warsh says Fed may 'have work to do' - CNBC — Fed's Hammack says 'now is the time to act' on raising interest rates - CNBC — U.S. Navy blockade slashes Iran oil exports as Trump administration shifts to economic warfare - CNBC — Iran calls for global pushback against U.S. sanctions, says diplomacy 'isn't impossible' - Greenwich Time — Regional powers seek off-ramp from Iran war, and other Middle East developments - NPR — Hundreds killed, over 1,000 missing after flash floods in Nepal and Tibet - France 24 — Rescuers grapple with mountain of challenges after the Nepal-Tibet floods - The Guardian — Nepal-Tibet flash flood: EU pledges €2m in aid as fears rise for missing - San Francisco Chronicle — Chinese foreign minister calls for US and China to 'overcome obstacles' ahead of presidential summit - Anadolu Ajansı — Xi, Modi, Putin to meet in Kyrgyzstan: What to watch at

SCO summit - Farm Progress — Putin moves to escalate war in Ukraine as talks at dead end - New York Post — Ukraine-Russia talks possibly including US may be held in September, Zelensky's top aide says - Arab News — Strikes in Gaza kill 4, including 2 children, a week after US asked Israel to curb attacks - Al Jazeera — UK has billions in contracts with firms tied to illegal Israeli settlements - France 24 — Spanish government urges calm after Ceuta anti-migrant protesters clash with police - BBC — Shootings surge as drug gangs torment Brussels - BBC — Oslo mayor calls off wedding ceremony after death of Norway's king - Tom's Hardware — Security researchers find surveillance implants in Chinese-made routers sold worldwide - Ars Technica — Authorities arrest 2 alleged members of prolific hacking group TeamPCP - The Washington Post — Civil rights groups make urgent plea for voting rights in DC, ahead of midterm elections - CBS News — Judge temporarily blocks new Postal Service rules for mail voting

Markets, Money & Deals - CNBC — S&P 500 futures are little changed as investors count down to Warsh's Jackson Hole address - CNBC — Marvell shares tumble 6% as outlook underwhelms despite 37% revenue growth - Bloomberg — Advent and Stripe Abandon \$50 Billion Pursuit of PayPal - MarketWatch — PayPal's stock plunges as a takeover bid is said to fall through - CNBC — Nvidia adds more than \$400 billion in value after blowout earnings boost AI confidence - CNBC — Marc Benioff is getting his mojo back as Salesforce's AI strength quiets skeptics - CNBC — CrowdStrike posts best day ever, Okta's stock pops nearly 29% as rising AI threat lifts earnings - CNBC — Imax says it's open to a sale. Why is no one buying? - 24/7 Wall St. — IREN Falls 6% as \$639M Mining Rig Writedown Overshadows AI Cloud Growth - SpaceNews — D2D venture Astrum Space joins SPAC resurgence with \$1 billion deal - Cryptonews — Coinbase Fall on Bank Stablecoin Plans Before CLARITY Act - MarketWatch — The bull market for stocks is defying everything, but Bank of America warns that an autumn reality check is coming

AI & Infrastructure - Telecoms — Amazon to buy another 2 million Nvidia GPUs - simplywall.st — How Does Amazon.com (AMZN) Tripling GPU Orders Reshape Its AI Buildout? - TechCrunch — a16z creates a \$1.1B 'Machine Age' fund to 'accelerate the physical buildout of AI' - Hacker News — EPA says power for data centers can sidestep pollution laws - The Elec — SK hynix to Start HBM4E Production at Indiana Plant in 2029 - Rigzone — JERA Backs Startup Developing Solutions for Grid-to-Data Center Projects - Yahoo Finance Canada — Is AI's 'capex tsunami' the

catalyst for the next big credit boom? - TechCrunch — Anthropic gets its first court win over the Pentagon's supply chain risk label

Model & Research Watch - CNET — The Powerful Stealth AI Model 'Ox Alpha' Is Now GLM-5.3-Flash, and You Can Use It - SCMP — Zhipu AI shares jump as viral Ox Alpha model revealed as GLM-5.3-Flash on Chinese chips - Tech Times — Sanctioned Chinese Chips Just Served 62 Trillion AI Tokens at Frontier Scale - Tech Times — Ox Alpha Was GLM-5.3-Flash: China Inference Chip Claim Stands Unverified - Reddit r/LocalLLaMA — Tencent/Hy4-preview 770B-A49B weight dropped - news24.com — Tencent touts new AI model it claims outperforms rivals - tech-insider.org — Meta Muse Glimmer: 30B Open AI Model for Consumer GPUs - Intelligent Living — Qwen3.8-Flash Matches DeepSeek V4 Pro on Coding Benchmarks at a Quarter of the Price - tech-insider.org — Claude Sonnet 5 vs GPT-5.6 Sol vs Kimi K3: \$20 Gap [2026] - Anthropic blog — Previewing the Model Hardware Standard - Science Daily — NASA's Roman telescope will see 100 times more sky than Hubble - ABC News — It was a spy satellite, now NASA's new telescope is set to head into space

What Builders Are Actually Using - Intelligent Living — Qwen3.8-Flash Matches DeepSeek V4 Pro on Coding Benchmarks at a Quarter of the Price - Hacker News — Show HN: We built open OpenRouter that turns usage into a better model (Experiential) - KDnuggets — The Local AI Stack for Productive SLMs - tech-insider.org — Meta Muse Glimmer: 30B Open AI Model for Consumer GPUs

Quick Check

- 1. Fed Chair Warsh's Jackson Hole speech was seen as hawkish. What was the immediate market reaction in the bond market?** Answer: The 2-year Treasury yield jumped, as investors reduced bets on near-term rate cuts.
- 2. What was the key claim about the GLM-5.3-Flash model (Ox Alpha) that makes it significant for the US-China tech war?** Answer: It ran on 100,000 domestically produced Chinese chips, suggesting

China may be able to train and run frontier AI models without US hardware.

- 3. Why did Marvell (MRVL) stock fall 6% despite beating revenue expectations?** Answer: Because its fiscal 2028 outlook (guidance) was muted, suggesting slower growth ahead. The market had already priced in the good quarter; the future mattered more (a "sell the news" reaction).

Generated 2026-08-28 08:00 PDT · 130 sources · model: deepseek-ai/DeepSeek-V4-Flash via <https://api.gmi-serving.com/v1>