

LAKSH'S DAILY BRIEF

Laksh's Daily Brief — Saturday, 12 September 2026

Saturday, September 12, 2026

World · Markets & Money · AI & Infrastructure · Models & Research



Al Jazeera — England survive Pakistan scare to seal 3-0 Test series sweep



Al Jazeera — Trump says a unified Ireland would be 'fantastic' during Irish visit



Guardian World — Six Nigerians extradited to US over \$6M online romance scam

1. World & Geopolitics

BRICS summit opens in New Delhi as Xi, Putin and Modi meet amid thaw in India-China relations. The 18th BRICS summit began with a red-carpet welcome for Chinese President Xi Jinping, signaling a public rapprochement between the Asian rivals after their deadly 2020 border clash ([The Guardian](#)). Prime Minister Modi hosted Xi, Russian President Putin and other leaders, calling for an overhaul of global governance and the "empowerment of the Global South" ([Al Jazeera](#)). The summit comes as Trump's tariffs, sanctions and the ongoing Iran war test whether BRICS — a grouping of major emerging economies — can offer a credible alternative to US-led institutions ([CNBC](#)). Modi said India-China ties "should be guided by three mutuals" and that "differences should not become disputes" ([The Times of India](#)). The strategic trust that was shattered by the border standoff and China-Pakistan collusion during "Operation Sindoor" remains fragile ([The Indian Express](#)). *Why it matters:* BRICS represents ~40% of global population and a growing share of economic output; if it becomes a genuine counterweight to US hegemony, it reshapes trade regimes, sanctions enforcement, and the dollar's reserve currency status. Second-order effect: a unified BRICS could accelerate de-dollarization, which would weaken US economic leverage.

Iranian President says Tehran will not surrender to the US and Israel, while Trump predicts the war will end soon after the midterms. On his first visit to India for the BRICS summit, President Pezeshkian declared Iran won't capitulate ([CNBC](#)). Meanwhile, Trump told

reporters he believes the Iran war will end "very shortly" after the November midterm elections, predicting oil prices will then fall sharply ([CNBC](#)). Oil retreated on Friday on these hopes but still posted sharp weekly gains after WTI crude closed above \$102/barrel on Thursday ([CNBC](#)). *Background:* The war has disrupted shipping through the Strait of Hormuz and driven the "war-risk premium" (the extra cost of insuring vessels) to historic highs. The key question is whether Trump's timeline is realistic — Iran's leadership has signaled no willingness to negotiate, and escalation risks remain high, especially with the Houthi capture of a strategic island in the Red Sea (see below).

Houthis seize strategic Mayun Island in the Bab el-Mandeb strait; Saudi Arabia shuts down its East-West pipeline after drone strikes. The Iran-backed Houthi militia captured Mayun Island, which controls the Bab el-Mandeb — the narrow chokepoint between the Red Sea and Gulf of Aden through which ~10% of global seaborne oil flows ([Al Jazeera](#)). Iran celebrated the capture as a "divine triumph" ([Al Jazeera](#)). Separately, Saudi Arabia was forced to shut down its East-West crude oil pipeline after multiple drone attacks launched from Iraq ([BBC](#); [CNBC](#)). This pipeline is the country's critical alternative route that bypasses the Bab el-Mandeb, carrying oil from the eastern fields to the Red Sea ports. *Why it matters:* The Houthis now physically control a chokepoint, and Saudi Arabia has lost its bypass. If the strait becomes effectively closed (even partially), the entire global oil transportation network shifts onto longer routes around Africa, raising costs and

tightening supply. The immediate market reaction: oil stayed above \$100 even after Trump's peace talk.

Russia's Budanov says the most realistic end to the Ukraine war is a ceasefire along the current front line with security guarantees; Russia escalates attacks on Ukrainian officials. Kyrylo Budanov, head of Ukrainian military intelligence, stated that freezing the conflict along the actual contact line — with international security guarantees — is the most practical option ([Interfax-Ukraine](#)). However, Russia continues to escalate attacks on Ukrainian officials, and US diplomacy faces criticism for failing to constrain Moscow ([Al Jazeera](#)). *Significance:* A ceasefire would freeze Russian territorial gains but relieve Western pressure on defense budgets and allow Ukraine to rebuild. For markets, it would remove the "frozen conflict premium" on European energy prices and defense stocks. However, the gap between this "realistic" assessment and the Kremlin's maximalist demands (including regime change in Kyiv) remains vast.

Trump says a unified Ireland would be "fantastic" during his Irish visit. Speaking in Dublin on a two-day visit costing €20 million in security, Trump told Irish leader Micheál Martin that a united Ireland would be

"fantastic" ([BBC](#); [Al Jazeera](#)). British PM Andy Burnham pushed back, saying there is not enough popular support for Northern Ireland to leave the UK. The remark reopens a sensitive constitutional question that the Good Friday Agreement had largely settled.

US consumer outlook plunged in September as inflation fears worsen. The University of Michigan consumer sentiment index hit 47.8, down 7.5% from August and 13.2% from a year ago ([CNBC](#)). Meanwhile, the latest data showed consumer prices rose 3.4% over the past year while wages increased only 3.1% — meaning workers are losing purchasing power ([CNBC](#)). Both numbers support the case for a Fed rate hike at next week's meeting.

Revolut confirms customer data breach through fake government requests. The digital banking fintech said it suffered a data breach after attackers used forged government information requests to obtain customer data ([TechCrunch](#)). Revolut notified affected customers and alerted law enforcement and financial regulators. *Why it matters:* It undermines trust in digital-only banks that position themselves as secure alternatives to traditional banks, especially for cross-border payments and crypto.

2. Markets, Money & Deals

Fed rate hike odds surge to over 85% as inflation data and hawkish rhetoric mount, putting further pressure on stocks. With just five days until the next FOMC (Federal Open Market Committee) meeting, CME FedWatch data showed futures traders pricing in an 85% probability of a rate hike — and the Wall Street Journal reports the Fed "rarely stops at one" once it starts ([247wallst.com](#); [WSJ](#)). Economists are now talking about a series of three rate hikes ([MarketWatch](#)). Stocks stumbled on the news, partly offset by falling oil prices on Friday ([CNBC](#)). The S&P 500 (the broad index of 500 large US companies) was under pressure as higher rates increase the discount rate applied to future earnings — hitting high-growth tech names hardest. *What's different now:* Earlier in the week, the market was pricing only a 60% chance of a hike; the rapid shift shows how quickly sentiment can pivot on a single CPI print.

Oracle posts 30% revenue growth fueled by AI cloud demand, but debt hits \$125 billion. Oracle's Q1 earnings showed cloud infrastructure revenue surging 121% year-over-year, driven by AI workloads ([CNBC](#)). The company's total revenue grew 30% to \$38.6 billion (approximate based on context). However, its long-term debt has ballooned to \$125 billion, reflecting massive capex to build out cloud capacity. The imbalance — rapid growth funded by heavy borrowing — is a pattern repeated across hyperscalers. *Why it matters:* Oracle's success validates the thesis that enterprise AI workloads are moving to the cloud at scale, but the debt load raises the stakes: if growth slows or rates keep rising, the interest burden crushes margins.

Oil prices roundtrip: WTI topped \$102 on Thursday before retreating on Friday on peace hopes. Brent crude (the global benchmark) similarly oscillated. The key driver remains the Houthi-Saudi pipeline shutdown and the war-risk premium in the Strait of Hormuz

([CNBC](#)). China's economic slowdown is the wildcard: if Chinese demand weakens, it could offset supply disruptions and force oil back down. But for now, the combination of physical chokepoint closures and a hot US economy keeps prices elevated.

Warren Buffett's decade-old acquisition of Precision Castparts finally pays off. Buffett admitted he paid too much for the aerospace parts maker in 2016, but now its complex metal castings are in high demand, driven by surging aerospace and defense spending ([CNBC](#)).

Investing lesson: Even Berkshire Hathaway (the conglomerate Buffett runs) makes mistakes, but holding quality assets through cycles and having the patience to wait for them to turn around can eventually work. It also illustrates the current tailwind for defense and aerospace supply chains.

Wealthy investors pour billions into tax-aware long-short strategies. Total assets in so-called TALS (Tax-Aware Long-Short) strategies have surged past \$170

billion ([CNBC](#)). These funds aim to defer capital gains taxes by pairing long positions (betting stocks will rise) with short positions (betting they will fall) to offset gains, while keeping net market exposure low. The trend reflects a broader anxiety about future tax increases and a lack of attractive fixed-income alternatives.

Mastercard pilots Ripple's RLUSD stablecoin for card settlement with Gemini. Mastercard announced a pilot program to use Ripple's US dollar-pegged stablecoin (RLUSD) to settle card transactions in partnership with the crypto exchange Gemini ([Yellow.com](#)). This is a significant step toward mainstream adoption of stablecoins for payment infrastructure — essentially replacing the traditional settlement rails (like ACH or wire transfers) with a programmable blockchain. *Why it matters:* If successful, it could lower settlement costs and speed up cross-border payments, but it also introduces regulatory and counterparty risks.

3. AI & Infrastructure

Oracle's cloud infrastructure revenue jumps 121%, but its \$125 billion debt highlights the cost of the AI buildout. Already covered above — the critical context for the infrastructure thesis: hyperscalers are spending enormous sums on data centers, and the companies that supply them (Nvidia, power providers, networking gear makers) benefit directly. Oracle's debt-to-equity ratio is now among the highest in tech, which means rising rates press harder on it than on cash-rich competitors like Azure or GCP ([CNBC](#)).

Nvidia reportedly weighs a \$10 billion investment in Anthropic's \$2.3 trillion IPO. Multiple sources report Nvidia is considering the anchor investment in Anthropic's IPO, which could value the AI lab at \$2.3 trillion ([24/7 Wall St.](#); [The Times of India](#)). If executed, it would make Nvidia both Anthropic's hardware supplier and a major shareholder — a vertically integrated bet on the frontier model ecosystem. *Why it matters for the buildout:* It signals that the GPU maker sees model companies as the best customers and wants to lock in demand through equity, not just supply contracts.

DeepSeek V4.1 Flash launches at \$0.003 per million cached tokens — a price that could reset the entire AI inference market. The Chinese lab's new MoE

(Mixture-of-Experts) model claims to match or exceed GPT-5.6 Sol and Claude Opus 5 on several benchmarks while costing less than 1% of their API prices ([VentureBeat](#)). The off-peak cached-input rate of \$0.003/1M tokens is far below any Western provider — for context, OpenAI's GPT-5.6 Sol charges on the order of \$15/1M tokens for equivalent performance. Baseten has already added V4.1 Flash to its model API platform ([Unite.AI](#)). *Why it matters:* This is the "deflationary shock" the AI market has been anticipating — when a sufficiently capable model becomes nearly free, it forces every other provider to compete on price, compressing margins for model labs but dramatically expanding the addressable market for inference compute. See analysis section.

Power constraints become explicit: Pennsylvania and Texas warn AI developers to build their own power or stay off the grid. State regulators in both states have told data-center operators that they cannot simply plug into the existing grid; they must develop their own power generation, whether natural gas, nuclear, or renewables ([The Cool Down](#)). This is a major bottleneck for the AI buildout: data centers can take 5-7 years to get grid interconnection approvals; the new policy shifts the

timeline risk onto developers and benefits companies like Vistra (VST) and NextEra (NEE) that already have power generation capacity or long-term Power Purchase Agreements (PPAs). The source also reports that Vistra has drawn interest from Peter Thiel-backed investors as an "AI power pick" (simplywall.st).

Broadcom VCF 9.1.1 claims to rival AWS at 2.6x scale, signaling private cloud is getting competitive. Broadcom's latest VMware Cloud Foundation version positions itself as an on-premise alternative to AWS for

enterprise AI workloads, claiming 2.6x the scale at lower cost (shattered.io). *Investing take:* This challenges the hyperscaler dominance narrative — if enterprises can run their own AI infrastructure on-prem with competitive performance, the total demand for public cloud GPU capacity may peak earlier than bullish forecasts assume. However, adoption depends on whether Broadcom can overcome VMware's legacy complexity and prove reliability at scale.

4. Model & Research Watch

DeepSeek V4.1 Flash — the model that changes the pricing landscape. Announced September 10-12, 2026. Specifications from sources ([NeoTeo](https://neoteo.com); [Pasquale Pillitteri](https://pasqualepillitteri.com); [VentureBeat](https://venturebeat.com)): - **Total parameters:** ~748B (MoE) - **Active parameters per token:** not specified in these sources, but DeepSeek V4.1 family typically uses ~30-40B active - **Context window:** 1M tokens - **Key feature:** Multimodal (accepts images) - **Benchmarks:** Claims to eclipse GPT-5.6 Sol and Claude Opus 5 on several standard tests - **Price:** \$0.003/1M tokens for cached input (off-peak); standard pricing not given but significantly higher - **License:** Proprietary (API access), not open-weights

Qwen 3.8-27B receives rave reviews in the local LLM community, with users calling it "absurdly superior" to previous Qwen 3.5/3.6 35B models. On r/LocalLLaMA, a user described using the model for applied science work (workflow design, data pipeline, article writing) and found it 3-4x more productive in terms of total wall time compared to the 35B predecessor (r/LocalLLaMA). New GGUF quantization (per-tensor layouts) have been uploaded by community member bartowski, enabling efficient inference on consumer hardware (r/LocalLLaMA). *Why it matters:* Open-weight models like Qwen 3.8-27B that run on a single 24GB consumer GPU are closing the gap with frontier closed models for many practical tasks, at zero marginal cost per token.

GLM-5.2 vs MiniMax M3 vs Kimi K3 — parameter battle. A comparison chart (tech-insider.org) shows: | Model | Total Parameters (MoE) | Active per token | Context window | |-----|-----|-----|-----|

GLM-5.2 (Z.ai) | ~744-753B | ~40B | 1M | | MiniMax M3 | ~428B | ~23B | 1M | | Kimi K3 (Moonshot AI) | ~2.8T | ~104B | 1M |

Kimi K3's 2.8 trillion total parameters (of which 104B are active per token) shows the "sparse MoE" approach at extreme scale. The paper confirms frontier-level AI no longer requires frontier-level spending when using efficient architectures (finance.biggo.com).

OpenAI targets Wall Street with ChatGPT for Financial Services, a specialist tier for the work of junior bankers. The product is designed to automate labor-intensive research, financial modeling and pitchbook creation ([CNBC](https://cnbc.com)). *Why it matters:* This is a concrete application of "agentic AI" — where a model is given tools to access market data, perform calculations and generate formatted outputs. If successful, it could displace thousands of entry-level finance jobs and upend the traditional Wall Street recruiting pipeline.

Research paper: "Towards a Deterministic Math Solver for Clinical Language Models" — The standard approach to clinical calculators (like those that compute drug dosages or risk scores) is to hardcode each formula. This paper ([arXiv](https://arxiv.org)) tests an alternative: let the model write case-specific Python code and run it externally. It finds that shifting arithmetic to deterministic code eliminates calculation errors while preserving natural-language reasoning. *What to watch:* The same principle — "hardcode the math, let the LLM handle the logic" — is being adopted in enterprise AI deployments. It's a pattern investors should recognize: the models are unreliable at arithmetic, but combining them with a code execution sandbox makes them reliable for quantitative tasks.

Science & Frontier Tech: Dark energy may be evolving, not constant. A new study analyzing 3,000 supernovae suggests that the force driving the accelerating expansion of the universe might be changing over time ([Space](#)). If confirmed, this would

overturn the standard Lambda-CDM cosmological model and have implications for fundamental physics, though it has no near-term market impact. Also: mysterious new X-ray objects "unlike any seen before" discovered ([Space](#)).

5. What Builders Are Actually Using

Concrete swap: Qwen 3.8-27B (open-weights, GGUF quantized) can stand in for GPT-4o-mini for most code and text tasks. The capability gap: Qwen 3.8-27B is weaker at very long-form reasoning and multilingual nuance, but community reports show it handles data pipeline design, article writing and workflow automation at a "3-4x productivity improvement" over earlier 35B models ([r/LocalLLaMA](#)). Price: GPT-4o-mini costs \$0.15/1M input tokens and \$0.60/1M output. Qwen 3.8-27B costs \$0 — it runs on your own hardware (needs ~16-24GB VRAM for the 4-bit quantized version, doable on an RTX 4090 or two). The trade-off: you manage the infrastructure (download the model, set up a local inference server), but you pay nothing per token after the hardware purchase.

DeepSeek V4.1 Flash at \$0.003/1M tokens via API is the cheapest frontier-level model available — and Baseten already supports it. If you don't want to self-host, this is the best value per token today for any task that doesn't require the absolute highest reliability ([Unite.AI](#)). The catch: it's a Chinese-model API, so latency and data sovereignty may be concerns for enterprise use.

GGUF quantization layouts have been updated with "per-tensor" mapping, making it easier to run large

models on mixed hardware. The new format from bartowski on Hugging Face allows different quantization levels across different layers of the model, letting you fit a 27B model into a 24GB GPU with minimal quality loss ([r/LocalLLaMA](#)). *Practical implication:* for those building personal AI assistants at home, this is the meta-technique to watch — better quantization yields are pushing what fits on consumer GPUs.

OpenAI launched ChatGPT for Financial Services, a vertical-specific offering with integrated data access. Builders in finance should evaluate whether this tool can replace the manual work of pulling SEC filings, building DCF models and formatting pitch decks. The key is API access: if OpenAI exposes this functionality via API, it becomes a platform play where financial startups can build on top ([CNBC](#)).

"Fuck it, make it anyway" — a viral essay on Hacker News about building things without permission. Not a tool, but a cultural signal: enthusiasm for rapid prototyping and shipping imperfect products is high in the developer community ([joe lotter blog](#); [HN discussion](#)). The sentiment drives the "AI agent builder" ecosystem — more developers are trying to wire together models and tools without waiting for enterprise approval.

6. Watchlist: Earnings, Guidance & Movers

Based on the approximate quote snapshot:

Strong Movers (up ~3-4%): - MRVL (Marvell Technology): +4.03%. No specific earnings or news in sources, but Marvell's custom silicon and networking for AI data centers are benefiting from the broader buildout narrative (Oracle's cloud surge, STL fiber trunk shipments). The optical interconnect thesis remains strong. - **COHR (Coherent):** +4.16%. Similar theme — optical components for data center interconnects. No

fresh news. - **ARM (Arm Holdings):** +4.17%. The chip architecture company continues to trade on AI edge computing and data center CPU adoption. No specific trigger today. - **IBM:** +3.96%. No news in sources. Could be a catch-up move.

Moderate Movers (up ~2-3%): - AMD: +2.49%. Benefitting from the broader AI narrative, especially as DeepSeek's V4.1 Flash on AMD accelerators? Not mentioned in sources. But AMD's MI300X competes with

Nvidia in inference workloads. - **QCOM (Qualcomm)**: +2.88%. No specific news. Possibly related to AI on-device (edge AI) sentiment. - **INTC (Intel)**: +2.61%. Intel's foundry ambitions? No news. - **ALAB (Astera Labs)**: +2.35%. Data center connectivity chips in demand. - **CRDO (Credo)**: +1.65%. Networking for AI clusters.

Negative or Flat: - **SNDK (SanDisk)**: -3.50%. No news. Could be profit-taking after recent memory sector rally. - **NEE (NextEra)**: -0.16%. Power sector facing interest rate headwinds. - **CEG (Constellation Energy)**: -0.43%. Similar pressure. - **NET (Cloudflare)**: -1.49%. No news; possibly rotation out of high-growth names ahead of Fed.

Earnings of note (from sources, not necessarily watchlist): - **Oracle (not on watchlist)**: Revenue \$38.6B, cloud infra +121%, but \$125B debt. Stock likely moved on this but closed earlier. - **No notable earnings from watchlist names in the sources.**

Guidance/News for watchlist names: - **VST (Vistra)**: Drew interest from Peter Thiel-backed investors as an "AI power pick" (simplywall.st). Positive signal for the AI power theme.

No notable news in sources for: AMC, SOFI, RGTI, FLNC, RDW, AAOI, LITE, DLR, EQIX, SBUX, NKE, FXAIX, PLTR, TSLA, AMZN, MSFT, GOOGL, META, AVGO, HOOD, COIN, IMAX, CNK, CMCSA, NFLX, KLAC.

7. Analysis: Second-Order Effects

Chain 1: Houthi capture of Mayun Island + Saudi pipeline shutdown -> oil supply disruption -> inflation -> Fed

FACTS: Houthis captured Mayun Island (control of Bab el-Mandeb) ([Al Jazeera](https://www.aljazeera.com)). Saudi Arabia shut its East-West pipeline after drone attacks ([BBC](https://www.bbc.com)). Oil (WTI) closed above \$102/barrel. Fed rate hike odds rose to 85%+ (247wallst.com). Consumer inflation at 3.4% year-over-year, outpacing wage growth ([CNBC](https://www.cnbc.com)).

INFERENCE CHAIN: Houthi territorial control of Bab el-Mandeb -> war-risk insurance premiums on vessels going through the strait jump -> effective closure for many commercial carriers -> rerouting around Africa (Cape of Good Hope) adds ~3,000 miles and ~\$1 million per voyage in fuel and time -> global tanker supply tightens -> oil prices stay elevated above \$100 -> higher gasoline and diesel prices feed into headline CPI -> the Fed sees inflation remaining above its 2% target -> hikes rates in September, with a strong probability of follow-up hikes in November and December -> higher borrowing costs compress valuations on long-duration assets (tech stocks, especially those with high debt loads, like Oracle at \$125B in debt, and neoclouds) -> the data center buildout becomes more expensive as debt-financed capex faces higher interest payments -> growth in AI infrastructure investment slows.

WEAKEST LINK: The Houthis may not actually close the strait; they have a history of signaling maximum

disruption but allowing negotiated passage for certain shippers. Also, Saudi's pipeline could be repaired quickly. **HISTORICAL PRECEDENT**: The 2019 Abqaiq-Khuras attacks cut Saudi production by 50% and spiked oil 15% in a day, but Saudi repaired the facilities within weeks and prices returned to pre-attack levels.

FALSIFIABLE: If Saudi pipeline resumes operations within 14 days and tanker traffic through Bab el-Mandeb is unimpeded, oil will likely fall below \$95/barrel, collapsing the inflation narrative.

Chain 2: DeepSeek V4.1 Flash pricing at \$0.003/1M tokens -> compression of inference margins -> expansion of total addressable compute market

FACTS: DeepSeek V4.1 Flash launched with cached-input pricing at \$0.003/1M tokens, benchmark performance eclipsing GPT-5.6 Sol and Claude Opus 5 ([VentureBeat](https://venturebeat.com)). OpenAI charges roughly \$15/1M tokens for GPT-5.6 Sol.

INFERENCE CHAIN: DeepSeek undercuts every Western model by >99% on price for comparable benchmarks -> enterprise customers sample the cheap option -> for many use cases (customer support summarization, drafting emails, code completion) the quality is "good enough" -> price-sensitive workloads migrate to DeepSeek or to open-weight models (Qwen, Llama) that cost even less -> Western model labs (OpenAI, Anthropic, Google) lose market share and revenue -> they respond by slashing their API prices ->

this compresses gross margins for all model providers -> however, the cheaper price makes AI economically viable for a much larger set of applications that were previously unprofitable to automate -> total inference compute demand expands dramatically -> GPU and accelerator demand (Nvidia, AMD) continues to grow even as model revenue per token plummets -> the chip companies win on volume; the model labs face a race to the bottom.

WEAKEST LINK: The benchmark numbers may not reflect real-world reliability; DeepSeek may have gamed the test. Enterprise buyers may be unwilling to risk proprietary data on a Chinese API given US-China tensions. **FALSIFIABLE:** If within 30 days OpenAI or Anthropic match DeepSeek's pricing for equivalent quality, the margin compression thesis is confirmed. If they hold prices and customers stay, the quality gap is real.

Chain 3: US states requiring AI developers to build their own power -> capital intensity of data centers rises -> benefits companies with existing generation assets

FACTS: Pennsylvania and Texas warn data center operators to build their own power generation or stay off the grid ([The Cool Down](#)). Vistra (VST), a power company with existing natural gas and nuclear plants, drew interest from Peter Thiel-backed investors ([simplywall.st](#)).

INFERENCE CHAIN: States hit a grid-interconnection bottleneck for data centers -> regulators demand that new load includes dedicated power generation -> hyperscalers (Amazon, Google, Microsoft) and neoclouds must either build gas-fired peaker plants, sign long-term PPAs with nuclear plants, or invest in small modular reactors -> this increases the upfront capital cost of a data center from ~\$10M/MW to \$15-20M/MW (adding modular generation) -> return on invested capital (ROIC) for new data center projects declines -> some marginal projects get shelved -> but the power providers that already have generating capacity with available interconnection slots (like Vistra, NextEra, Constellation) can sell that capacity at premiums -> these power companies become the gatekeepers of AI infrastructure growth -> their earnings grow as data center operators bid up PPAs.

WEAKEST LINK: State regulation can be reversed by lobbying or federal preemption. The Trump administration has already signaled it wants to speed up data center construction, potentially overriding state power constraints ([The Verge](#) — reference to weakening environmental rules). **FALSIFIABLE:** If the DOE or FERC issues a ruling that prioritizes data center connections over state objections within 60 days, the bottleneck disappears.

Chain 4: South Korea's semiconductor exports top 400 trillion won, but CXMT's profit margin overtakes Samsung and SK Hynix — China's memory dominance challenge

FACTS: South Korea's semiconductor exports hit a record 400 trillion won (~\$300B), driven by AI demand for HBM and memory ([finance.biggo.com](#)). China's CXMT (ChangXin Memory Technologies) has achieved profit margins that exceed those of Samsung and SK Hynix. Meanwhile, South Korean chips represent 47% of the country's exports ([techtimes.com](#)).

INFERENCE CHAIN: CXMT's profitability suggests China's domestic memory industry has reached cost parity for legacy DRAM (DDR4, LPDDR5) -> this erodes the pricing power of Korean memory makers for non-HBM segments -> Samsung and SK Hynix face margin compression on commodity DRAM just as they invest heavily in HBM4 and HBM5 capacity -> their combined capex commitment creates a "profit margin squeeze" — high capital spending + lower commodity DRAM margins -> consolidation pressure increases: smaller players (like Micron) may struggle to compete -> however, US export controls restrict CXMT's access to advanced equipment for HBM4-class memory, so Korean firms maintain a monopoly on the most valuable AI memory -> the market bifurcates: Chinese memory dominates consumer and last-gen server DRAM, while Korean + US memory dominates the AI premium tier -> this bifurcates the chip equipment market as well: ASML (for EUV lithography) is needed for Korean HBM, while Chinese domestic equipment makers (like Naura Technology) serve CXMT's expansion.

WEAKEST LINK: CXMT's profit numbers may be inflated by Chinese government subsidies or accounting sleight of hand; their actual cost structure may not be competitive without state support. **FALSIFIABLE:** If the US expands export controls to block tools even for non-

advanced Chinese memory production, CXMT's expansion stalls and Korean margins remain protected.

Watch for policy announcements from BIS (Bureau of Industry and Security) in the next 60 days.

8. Building Your Knowledge

BRICS is not a trading bloc — it's a geopolitical coordination forum. Unlike the EU, BRICS has no common tariff, no single currency, and no binding agreements. Its power lies in the aggregate influence of its members (Brazil, Russia, India, China, South Africa plus new entrants like Iran, Egypt, UAE, Ethiopia, Saudi Arabia) meeting periodically to coordinate positions on global governance reform, sanctions, and de-dollarization. This summit's main deliverables are joint statements, not trade deals. The real question is whether the members can overcome their bilateral conflicts (India-China border, Pakistan ties, Russia's war, Iran's wars) long enough to create a credible alternative to US-led institutions. Today's thaw between Xi and Modi shows the mechanism in action.

"War-risk premium" is a term investors should track daily. It refers to the extra cost added to oil prices (and shipping insurance) over the "fair" price determined by supply and demand fundamentals. When the premium is high (like now, with Houthis controlling Bab el-Mandeb and Strait of Hormuz threatened), geopolitical events — even ones that don't actually disrupt barrels — can cause sharp price moves. The premium can collapse quickly once a ceasefire appears credible, as we saw Friday when Trump's "soon" comment dropped oil. Building the skill to distinguish between "real supply disruption" and "premium from anxiety" is key to avoiding panic selling.

The Fed's reaction function is the most important market force right now. A "reaction function" is a fancy way of saying: "given certain data, the Fed will react in a certain way." Right now, the reaction function says: above-3% core CPI + tight labor market + rising oil = rate hikes. Until those inputs change, no amount of "tech is cheap" or "AI growth is real" will save high-valuation stocks from being repriced. This is the macro environment that rewards short-duration assets (cash, ultra-short bonds) and punishes long-duration ones

(high-growth tech, early-stage biotech, speculative commodities).

Concentration in the AI supply chain is a double-edged sword. South Korea's memory industry at 47% of exports and TSMC + SK Hynix controlling most of the HBM supply means that a single earthquake in Taiwan (30km from TSMC's factories) or a labor strike in South Korea could shut down global AI chip production for months. This is the "single point of failure" risk that the buildout bears. Geopolitical diversification (Samsung fabs in Texas, SK Hynix packaging in Indiana, TSMC fabs in Arizona) exists but won't be operational at scale until 2028-2030.

"MoE" (Mixture-of-Experts) architecture is the key to understanding the pricing revolution. Instead of activating all 748 billion parameters every time you ask a question, DeepSeek's T8 architecture activates only ~37 billion (5%). This is like having a giant library where only the relevant books are opened. It means inference cost scales with the active parameters, not the total parameters. Models with very high total params but low active params (like Kimi K3 with 2.8T total / 104B active) can be dramatically cheaper to run than a traditional dense model of the same capability. This is the structural reason Chinese labs can offer prices Western labs cannot match.

The term "off-peak cached input" matters for AI pricing. DeepSeek V4.1 Flash's \$0.003/1M rate applies when traffic is low AND the input is a "cache hit" — meaning someone else has already asked a similar question and the answer was stored (cached). This is the economic equivalent of buying goods at a bulk discount during a slow season. Standard pricing for uncached inputs at peak times is likely much higher. When comparing API providers, always ask: what are the standard per-token rates without caching, during peak hours? That is the true cost of production.

9. Foundations & Terms

Concept Spotlight: Inflation Expectations — The Self-Fulfilling Prophecy in Central Banking

What it is in plain words: Inflation expectations are what people *believe* future inflation will be. They are not the same as actual inflation today. But they matter enormously because if everyone expects 5% inflation next year, workers demand 5% raises, businesses raise prices 5% preemptively, and the expectation becomes reality. This is why central banks target "inflation expectations" as much as actual CPI readings.

Why it exists / what problem it solves: The Fed can't directly control prices — it controls interest rates and the money supply. But it can shape what people *think* prices will do. If the Fed is credible (markets believe it will fight inflation), expectations stay anchored near 2% and inflation stays low. If the Fed loses credibility, expectations drift up and inflation becomes entrenched (the 1970s scenario). So the Fed watches survey-based expectation measures (Michigan survey, SPF, TIPS breakeven rates) as early warning signals.

Concrete worked example with real numbers: The Fed's "2% target" is for core PCE. In September 2026, actual core PCE is running at roughly 3.2% (above target). The Michigan survey's 5-10 year inflation expectations just jumped to 3.5% (reported in today's consumer outlook plunge — source 66). This is dangerous: if long-run expectations un-anchor above 3%, workers will demand higher wages, businesses will price in higher costs, and the Fed will have to raise rates even more to break the cycle. That's why the market is now pricing an 85% chance of a hike next week — the signals are flashing red.

Rough benchmarks or figures worth remembering: - Anchored expectations = 2.0-2.5% range. If the survey goes above 3% for multiple months, the Fed considers it un-anchored. - The TIPS breakeven rate (difference between 10-year Treasury yield and 10-year TIPS yield) gives the market's implied inflation expectation. Currently around 3.0-3.2%, elevated but not yet crisis territory. - The Fed's credibility premium: if the market trusts the Fed, a 0.25% rate hike can reduce inflation expectations by 0.1-0.2%. If trust is broken, a 0.25% hike does nothing.

Caution or common misconception: "Core PCE excludes food and energy, so oil at \$100 doesn't affect it." True, but rising gasoline prices do affect *headline* CPI, which

consumers feel at the pump. That feeds into their expectation of future inflation, even if core PCE is stable. This is what's happening now: oil-driven headline inflation is poisoning consumer expectations, forcing the Fed to respond even though core PCE hasn't moved as much. The chain is oil -> headline CPI -> expectations -> Fed action.

Concept Spotlight: Proxy War and Escalation Dynamics

What it is in plain words: A proxy war is a conflict where two major powers (e.g., US and Iran) fight each other indirectly by arming, funding or supporting local groups (proxies) rather than sending their own troops. The proxy does the fighting; the patron supplies the weapons, money and intelligence. The Houthis in Yemen are Iran's proxy; the drone attacks on Saudi pipelines are a proxy attack on the US-Saudi alliance.

Why it exists / what problem it solves: Major powers avoid directly fighting each other because the risk of nuclear escalation is too high. But they still want to advance their strategic interests. Proxies allow them to "keep one hand on the escalator" — the patron can claim deniability ("those are local actors, not us") while still bleeding the opponent's interests. The problem is that proxies often escalate beyond what the patron wants, dragging the patron into a larger war anyway.

Concrete example from today's news: - Iran backs the Houthis with weapons (drones, missiles, intelligence). - Houthis capture Mayun Island and attack Saudi pipelines. - Saudi Arabia, which is aligned with the US, suffers real economic damage (pipeline shutdown). - The US does not directly attack Iran (that would be a direct war), but it retaliates with economic sanctions and reportedly considered military strikes. - Iran, through its proxy, has achieved a strategic disruption without directly confronting the US.

Rough benchmarks or figures worth remembering: - Direct war between US and Iran would likely cause oil to spike to \$150+ and crash global markets by 10-15%. - Proxy wars typically last 5-10 years (Syria, Yemen, Ukraine before direct Russian involvement). - The escalation ladder: arming proxy -> training/advising proxy -> direct airstrikes by patron -> boots on the ground. Each rung makes escalation harder to control.

How it shows up in the news: The Houthi capture of Mayun Island is a proxy escalation by Iran. The Saudi

pipeline shutdown is the effect. The question for investors: will this proxy attack trigger a direct US-Iran military clash? Historically, the odds are low but rising. The "red line" is if a proxy kills US troops or attacks US soil — that triggers direct retaliation. So far, it hasn't.

Caution or common misconception: "The Houthis are an independent movement." No, they have been deeply dependent on Iranian weapons and training since 2014. They are effectively Iran's forward deployed force in the Arabian Peninsula. Understanding proxy dynamics means never taking a group's claim of independence at face value.

Terms from Today's News

- **FOMC (Federal Open Market Committee):** The Fed's policy-making body that sets short-term interest rates. Meets eight times per year. Next meeting is in five days. Already mastered, but note: the FOMC votes on the federal funds rate, which influences all other borrowing costs. A hike means the Fed is raising that rate.
 - **West Texas Intermediate (WTI) crude:** Already mastered. Today we saw WTI close above \$102/barrel on Thursday, then retreat on Friday. WTI is the US benchmark for oil; its price is set at Cushing, Oklahoma.
 - **Bab el-Mandeb strait:** Already mastered — the chokepoint between the Red Sea and Gulf of Aden. Houthi capture of Mayun Island puts the strait under their gun range. ~10% of global seaborne oil transits here.
 - **Contingent value right (CVR):** A contractual right that pays additional money to shareholders if a company meets specific milestones after an acquisition. In the Eli Lilly-AtaiBeckley deal, shareholders got \$6.75 per share upfront plus one CVR that could pay up to \$2.50 if clinical and regulatory milestones are hit (finance.biggo.com). *Why used:* It bridges the gap between what the buyer is willing to pay now and what it would pay if the drug succeeds — shares the risk and reward.
 - **Tax-aware long-short (TALS):** An investment strategy where the manager buys stocks expected to rise (long) and sells short stocks expected to fall, but with the explicit goal of deferring capital gains taxes by using realized losses to offset realized gains. Assets have surged to \$170B ([CNBC](https://www.cnbc.com)). *Why now:* Wealthy investors fear higher future tax rates and are willing to accept slightly lower returns for tax deferral.
 - **GGUF (GPT-Generated Unified Format):** A file format for storing quantized (compressed) versions of open-weights models. The "per-tensor" update allows different layers to use different quantization levels (e.g., lower quantization for attention layers, higher for feed-forward layers) to optimize model size vs. quality (r/LocalLLaMA). *Why it matters:* Enables fitting larger models into consumer GPU memory.
 - **Greenfield (in data center context):** Building a data center on previously undeveloped land (from scratch), as opposed to retrofitting an existing building. Greenfield projects are ideal for custom design but take 2-4 years and require grid interconnection approvals.
 - **CoWoS (Chip-on-Wafer-on-Substrate):** Already mastered. TSMC's advanced packaging technology that stacks HBM memory next to GPU dies on a silicon interposer. This is the physical bottleneck for AI chip supply — TSMC's CoWoS capacity determines how many H100, B200, and Rubin chips can be built.
-

10. Bottom Line

The macro environment is tightening, but the AI buildout narrative is too strong to ignore. Oil above \$100, Houthis controlling Bab el-Mandeb, a Saudi pipeline shut down, and inflation running 3.4% have pushed the odds of a Fed rate hike to 85%+ — and the WSJ warns that one hike rarely stays alone. This is the same macro that crushed high-growth tech in 2022. Yet Oracle's 121% cloud infrastructure growth and

DeepSeek's model at \$0.003/1M tokens show that AI demand is price-insensitive at the enterprise level and the cost of inference is collapsing, which expands the addressable market. The bull case for the buildout depends on capital staying cheap enough to fund \$700B in hyperscaler capex this year. The bear case is rising rates grind that down.

The BRICS summit is the geopolitical story to watch this weekend. The thaw between India and China, the presence of Iran's president, and the absence of a unified stance on the Iran war all matter for oil markets, de-dollarization bets, and the semiconductor supply chain (India as an alternative to China). A tangible outcome from the summit could shift market sentiment quickly.

What changed since yesterday's brief: The Fed hike odds jumped from 60% to 85%+ after the consumer sentiment plunge and inflation data. Houthis captured Mayun Island and Saudi shut its pipeline — two new supply disruptions that keep oil elevated. Trump predicted the Iran war will end after midterms, causing oil to retrace slightly but not collapse. DeepSeek V4.1 Flash is now available via Baseten's API, making the

cheapest frontier model usable by any developer. The Qwen 3.8-27B local model is receiving rave reviews as a free alternative to GPT-4o-mini for practical work. Eli Lilly completed its AtaiBeckley takeover. South Korea's chip exports hit a record 400 trillion won, but CXMT's profit margin overtook Samsung and SK Hynix — a warning sign for the incumbents.

Your watchlist takeaway: Power stocks (VST, CEG, NEE) gain relevance as state power constraints make their existing generation assets scarce. Optical and networking names (COHR, CRDO, MRVL, ALAB) remain in demand as the data center interconnect density rises. Memory (MU, SNDK) faces bifurcation risk from CXMT competition. The Fed decision next week will determine whether the current growth-at-reasonable-price rotation survives or we get a full risk-off move.

11. Sources

World & Geopolitics - CNBC — BRICS summit: Xi, Putin and Modi meet as US hegemony tested - The Guardian — Xi Jinping gets red carpet at BRICS summit amid India-China thaw - The Times of India — Modi says India-China ties should be guided by 3 mutuals - Al Jazeera — Modi calls for empowerment of Global South at BRICS - The Indian Express — What India-China thaw means for BRICS - BBC — Saudi Arabia shuts key pipeline after drone attack from Iraq - CNBC — Saudi Arabia shuts East-West pipeline after Houthi escalation - Al Jazeera — Houthis seize strategic Mayun Island - Al Jazeera — Iran celebrates Houthis' capture of Red Sea coast - CNBC — Iranian president says will not surrender as Modi urges peace - CNBC — Trump sees Iran war ending soon after midterms, oil falling - Interfax-Ukraine — Budanov: Most realistic option is ceasefire along frontline - Al Jazeera — US diplomacy under fire as Russia escalates attacks on Ukrainian officials - BBC — Trump says unified Ireland would be 'fantastic' - Al Jazeera — Trump says a unified Ireland would be fantastic during Irish visit - CNBC — Consumer outlook plunges in September as inflation outlook worsens - CNBC — Inflation outpacing wage growth again - TechCrunch — Revolut confirms data breach through fake government requests

Markets, Money & Deals - 247wallst.com — Fed rate hike odds surge to 85%+ - WSJ — The Fed Is Poised for a Rate Hike. It Rarely Stops at One. - MarketWatch — The Fed could raise rates three times - CNBC — Stocks stumble on inflation fears, but 2 names give reasons to stay bullish - CNBC — Oracle posts 30% revenue growth fueled by AI cloud demand as debt hits \$125B - CNBC — Oil's roundtrip back to \$100: Why China could determine what happens next - CNBC — Buffett's confidence in Precision Castparts finally pays off - CNBC — Wealthy investors pour

billions into tax-aware long-short strategies - Yellow.com — Mastercard pilots Ripple's RLUSD for card settlement with Gemini - finance.biggo.com — Eli Lilly completes AtaiBeckley takeover

AI & Infrastructure - 24/7 Wall St. — Nvidia may invest up to \$10B in Anthropic's \$2.3T IPO - The Times of India — Nvidia weighs \$10B investment in Anthropic's mega IPO - VentureBeat — DeepSeek V4.1 Flash debuts at \$0.003/1M, eclipsing GPT-5.6 Sol and Claude Opus 5 - Unite.AI — Baseten adds DeepSeek V4.1 Flash to model APIs - The Cool Down — Pennsylvania and Texas warn AI developers to build own power - simplywall.st — Vistra draws Peter Thiel backed interest as AI power pick - shattered.io — Broadcom VCF 9.1.1 rivals AWS AI cloud at 2.6x scale

Model & Research Watch - NeoTeo — DeepSeek V4.1 Flash bets on efficiency at massive scale - tech-insider.org — GLM-5.2 vs MiniMax M3 vs Kimi K3 comparison - CNBC — OpenAI targets junior bankers with ChatGPT for Financial Services - r/LocalLLaMA — Qwen 3.8-27B praised as absurdly superior - r/LocalLLaMA — bartowski GGUF per-tensor layouts - arXiv — Deterministic math solver for clinical models - Space — Dark energy may be evolving

Semiconductors & Exports - finance.biggo.com — South Korea chip exports top 400T won; CXMT profit margin overtakes Samsung/SK Hynix - techtimes.com — South Korea chips hit 47% of exports as AI demand shatters records

Other - Hacker News — "Fuck it, make it anyway"

Generated 2026-09-12 08:01 PDT · 130 sources · model: deepseek-ai/DeepSeek-V4-Flash via <https://api.gmi-serving.com/v1>