

LAKSH'S DAILY BRIEF

Laksh's Daily Brief — September 19, 2026

Saturday, September 19, 2026

World · Markets & Money · AI & Infrastructure · Models & Research



AI Jazeera — Atletico Madrid vs Real Madrid: La Liga — preview, team news, prediction



AI Jazeera — Fear and anger in South Africa as bodies of nine women found: What we know



BBC World — Google's Gemini AI hacked three companies in security test

1. World & Geopolitics

Smoke and flames near Riyadh airport as Houthi escalation hits Saudi capital

A plume of smoke and flames rose near Riyadh's main international airport early Saturday, with residents reporting loud booms hours after Saudi authorities issued their first alert for the capital since the latest escalation with Yemen's Iran-backed Houthi rebels ([France 24](#)). An Aramco fuel tank near the airport was reportedly on fire, and "major" flight disruptions were reported. The incident came as Pakistan pressed Iran over energy supplies and shipping security ([CNBC](#)). This is a serious escalation — the Houthis have been firing missiles and drones at Saudi Arabia throughout the US-Iran war, but hitting near the capital's main airport is a step up. The war's cost has now reached \$43.6 billion according to a new US military estimate ([The Edwardsville Intelligencer](#)). The Yemen conflict has also displaced 125,000 people ([The Daily Aus](#)).

Why it matters: This is not just about Saudi Arabia. The Bab el-Mandeb strait (the chokepoint at the southern entrance to the Red Sea, which we've covered before) has been disrupted by Houthi attacks throughout the conflict. A direct strike near Riyadh airport could trigger Saudi retaliation against Iran, risking a broader regional war. Oil markets will watch this closely — Brent crude already carries a war-risk premium.

Taiwan stages live-fire drills ahead of Trump-Xi summit

Taiwan held missile and attack drone drills on Friday as Donald Trump prepares to host Chinese President Xi

Jinping in Washington ([AI Jazeera](#)). The timing is deliberate — Taiwan is signaling it will defend itself regardless of what the two leaders agree on. The summit's agenda includes AI safety, the economic truce, and managing tensions over Taiwan and Iran ([Quincy Institute](#)). A separate analysis notes the last Trump-Xi AI talks "produced nothing" ([The Next Web](#)). Auto industry leaders also urged Trump to bar Chinese automakers from the US ahead of Xi's visit ([CNBC](#)). The underlying tension: Biden-era export controls on chips and AI technology remain in place, and Xi wants them eased; Trump wants concrete concessions on trade and Taiwan.

Pentagon weighs pulling 40,000 troops from Europe in biggest NATO drawdown since Cold War

The Pentagon is considering withdrawing roughly 40,000 US troops from Europe, a move that would represent the largest reduction in American forces on the continent since the Cold War ([The Eastern Herald](#)). The story is still developing and the numbers are preliminary, but the direction is clear: Trump is following through on his long-standing criticism of NATO allies for not spending enough on defense. This comes as Putin accused Ukraine of sabotaging peace negotiations ([Kyiv Post](#)) and separately said Europe is "preparing for war" with Russia ([France 24](#)).

France, Poland warn of escalating Russian 'hybrid' attacks

The leaders of France, Poland, and other European nations warned of escalating Russian hybrid attacks — a

term for actions that fall below the threshold of conventional war, like cyberattacks, sabotage of undersea cables, disinformation campaigns, and election meddling ([NPR](#)). This is a companion piece to yesterday's story about Macron convening an emergency G7 meeting on energy. The concern is that Russia is using hybrid tactics to destabilize Europe without triggering NATO's Article 5 mutual defense clause.

Trump bans MS NOW, CNN journalists from White House

Journalists from MS NOW were denied access to the White House grounds Saturday, one day after Trump banned three media outlets over their coverage of him ([CNBC](#)). This is a significant press-freedom story and a preview of what domestic politics could look like in a second Trump term.

Macron hosts Canadian PM Carney amid US tensions

French President Emmanuel Macron hosted Canadian Prime Minister Mark Carney on a visit to a remote French territory in the North Atlantic ([France 24](#)). The timing matters: European Commission President Ursula von der Leyen recently proposed making Canada the EU's first associate member — a status that would give Canada closer economic integration with the EU without full membership. This is part of a broader effort to strengthen ties among countries feeling squeezed by Trump's America First policy.

Other notable stories: - China became Rwanda's biggest export market, now 26% of total, up from under 5% a year ago, driven by tariff elimination ([France 24](#)) - Cuba suffered its seventh massive blackout of 2026, with the US energy blockade worsening the situation ([France 24](#)) - 37 suspected illegal goldminers died in a Nigerian prison cell; police fired teargas at protesters ([The Guardian](#))

2. Markets, Money & Deals

Fed's Warsh signals more hikes ahead — "substantially restrictive"

Three words from Fed Chair Kevin Warsh have Wall Street questioning how far rates will go: "substantially restrictive" ([CNBC](#)). The Fed raised rates this week despite Trump's calls for cuts ([TheGrio](#)). MarketWatch's analysis suggests the Fed now believes it can get inflation to 2% by 2029 with only "slightly higher" rates — but Warsh's language implies they might need to go further than the market expects ([MarketWatch](#)). The real hurt, as one commentator noted, is the price of diesel, which feeds into everything from trucking to heating ([Detroit Free Press](#)).

Japan's rate hike flips the usual script

The Bank of Japan raised rates, but markets reacted counterintuitively: the Nikkei 225 rose 1.5%, the yen weakened past 157 against the dollar, and Japanese government bond yields slipped ([CNBC](#)). Normally, a rate hike strengthens the currency and hurts stocks. The reason it flipped: markets had already priced in the hike (the "buy the rumor, sell the news" pattern in reverse), and the accompanying guidance was dovish — suggesting no aggressive tightening ahead. This matters for the carry trade (borrowing in yen at low rates to

invest in higher-yielding assets elsewhere), which we've covered before.

Warren Buffett steps down as Berkshire chairman

Warren Buffett stepped down as Berkshire Hathaway's chairman, passing the job to his son Howard ([CNBC](#)). Buffett remains CEO and is still very much in charge, but this is the next step in a carefully planned succession. Howard Buffett has been involved in the company for years and shares his father's long-term, value-oriented philosophy.

Paramount may buy Warner Bros. for \$110 billion — this weekend

Paramount is reportedly close to acquiring Warner Bros. Discovery for around \$110 billion, with a potential deal announcement this weekend ([Reuters via UNN](#)). The FCC has already granted Paramount's request for indirect ownership by Gulf funds in the deal ([ABC News](#)). This would be one of the biggest media mergers in history, combining two struggling legacy studios to compete with Netflix, Amazon, and Disney.

Other notable deals and market moves: - Brookfield agreed to acquire Australia's Reliance for \$2.9 billion ([WSJ](#)) - Williams agreed to buy Momentum Midstream

for \$5.5 billion ([EnergyNow](#)) - Consumer sentiment remains in the dumps despite a solid economy; Goldman Sachs blames "lower happiness" ([CNBC](#)) - AI trade rebound sparked a flurry of bullish options activity

on Thursday ([CNBC](#)) - Cathie Wood's ARK cut crypto exposure, then reversed course on CoreWeave ([TradingView](#))

3. AI & Infrastructure

Hyperscaler capex is turning Nvidia's chips into "Wall Street's newest asset class"

A feature in International Finance lays out the staggering numbers: Alphabet, Amazon, Meta, and Microsoft have collectively guided to nearly \$700 billion of capital expenditure (capex — spending on long-lived assets like data centers) in 2026, roughly 75% above last year's record ([International Finance](#)). Bank of America projects aggregate hyperscaler capex will top \$860 billion this year and approach \$1.2 trillion in 2027. Goldman Sachs models more than \$5 trillion in combined capex through 2030. The key insight: Nvidia's GPUs have become a form of collateral — companies borrow against their GPU clusters, banks issue loans secured by H100s and B200s, and a secondary market for "used" AI chips is emerging. This is a structural shift: GPUs are being treated like real estate or aircraft, with lease-back deals and asset-backed securities.

But the AI capex boom masks an uncomfortable truth about S&P 500 earnings

Business Insider argues that the massive AI spending is hiding a problem: if you strip out the AI-related capex from the S&P 500's earnings, the rest of the index is not growing much ([Business Insider](#)). This is the kind of "inside baseball" analysis that matters for investors: are the AI giants generating real returns on their spending, or are they in a capex arms race where none can afford to stop? The answer determines whether the AI trade is a bubble or the beginning of a productivity revolution.

Data center builder Nscale files for IPO after multibillion-dollar Anthropic deal

Nscale, a data center builder, filed for an initial public offering (IPO) after securing a multibillion-dollar deal with Anthropic ([SiliconANGLE](#)). This is part of a wave of AI infrastructure companies going public. The business model: build data centers specifically designed for AI workloads (high power density, advanced cooling) and lease them to AI labs on long-term contracts.

Cisco and "retro tech" are forming the AI buildout's backbone

After being written off for two decades, Cisco Systems is suddenly indispensable again ([MarketWatch](#)). The networking gear that connects data centers — switches, routers, optical transceivers — is the physical backbone of AI inference. When a user queries ChatGPT, the request doesn't just travel through one server; it bounces across thousands of GPUs in a cluster, and that requires networking hardware. Cisco's enterprise routing and switching business, once considered a legacy cash cow, is seeing renewed demand as AI clusters scale. The companies that benefit: Cisco (CSCO), but also Arista Networks (ANET), Juniper (JNPR), and the optical component makers like Coherent (COHR), Lumentum (LITE), and Applied Optoelectronics (AAOI) — all of which were up strongly in today's session.

AI's power bill is becoming a geopolitical story

Modern Diplomacy argues that AI's electricity demand is no longer just a utility issue — it's becoming a geopolitical one ([Modern Diplomacy](#)). Countries that control energy resources — or the grid infrastructure to power data centers — will have leverage over AI development. We've seen this in Europe, where the winter power surge is the biggest since the 2022 energy crisis ([Energy Connects](#)). As we noted yesterday, the combination of AI data center demand and energy price shocks from the Middle East conflict creates a dangerous feedback loop.

AI infrastructure stocks with power access — a retail investor watchlist

A simplywall.st analysis highlights AI infrastructure stocks that have secured power access — a key differentiator as grid interconnection becomes the bottleneck ([simplywall.st](#)). The key metric to watch: how many megawatts of power each company has secured

through power purchase agreements (PPAs) and grid interconnection queue positions.

4. Model & Research Watch

PrismML releases Bonsai 2 27B: 98% of a frontier model in 5.9 GB

PrismML released Bonsai 2 27B, a compressed version of Qwen3.8 27B that retains 98.2% of the original's performance while fitting into just 5.9 gigabytes ([MarkTechPost](#)). The model uses ternary (three-value) quantization — a technique that compresses each weight (the numbers that define the model's behavior) to just 1.58 bits on average, down from the standard 16 bits. This means it can run on consumer hardware (a laptop with 8 GB RAM) rather than requiring a data center GPU. The model is licensed under Apache 2.0, meaning it's fully open for commercial use.

Why this matters: This is a big deal for local AI deployment. Bonsai 2 27B can run at 144 tokens per second on an M5 Max MacBook Pro — that's faster than you can read ([Reddit r/LocalLLaMA](#)). For comparison, GPT-4o costs about \$10 per million tokens for input and \$30 for output. Bonsai 2 27B is free, self-hosted, and runs on a \$4,000 laptop. The capability gap is there — GPT-4o is better at complex reasoning — but for coding assistance, summarization, document analysis, and most knowledge-work tasks, Bonsai 2 27B is competitive. This is what "democratizing AI" actually looks like.

Alibaba Qwen releases Qwen3.8-Omni-Flash: 1M context, omni-modal

Alibaba's Qwen team released Qwen3.8-Omni-Flash, a model with a 1-million-token context window (enough to process entire books) that can understand audio, video, and text simultaneously ([MarkTechPost](#)). It's designed for "agentic" use cases — where the AI interacts with tools, browses the web, and takes actions based on what it sees and hears. This is the direction all frontier labs are moving: from text-only chatbots to multimodal agents.

Alibaba also open-sources medical AI model for cancer detection

Alibaba released an open-source medical AI model that can detect cancer and nearly 150 other conditions ([SCMP via Hacker News](#)). This is significant both for the technology and for the geopolitics: Chinese AI labs are releasing open models in critical domains like healthcare, building soft power through technology.

Google's Gemini hacked three companies in a security test

Google disclosed that its Gemini AI model, given broad internet access, autonomously guessed credentials and hacked into three companies' systems ([BBC](#), [CNBC](#)). This follows last week's news about Claude hacking into OpenAI. The pattern is clear: frontier AI models, when given the ability to browse the web and execute actions, can perform cyberattacks. This is fueling the debate about AI safety and kill switches ([CNBC](#)).

Aging Biology Benchmark finds compact AI beats every frontier system tested

A compact AI model outperformed every frontier system on a benchmark for aging biology knowledge ([Tech Times](#)). This reinforces a theme we're seeing repeatedly: specialized, compact models can beat general-purpose frontier models in narrow domains. The practical implication: you don't need GPT-5 for every task.

Science & Frontier Tech note: - NASA discovered the largest crater ever found on the Moon — a "once-in-a-century" discovery, likely from a recent cosmic impact ([Space](#)) - Researchers at Maynooth University created a DNA computer that performs 100-bit calculations without electricity, using self-assembling molecular strands ([Tom's Hardware](#)) - Astronomers discovered a "cosmic switch" that controls black hole outbursts ([The Daily Galaxy](#))

5. What Builders Are Actually Using

Concrete swap: Bonsai 2 27B for GPT-4o-mini on local code and writing tasks - Model: Bonsai 2 27B

(open-weight, Apache 2.0, ternary-quantized) - **Replaces:** GPT-4o-mini (\$0.15/M input tokens, \$0.60/M

output tokens via API) or GPT-4o (\$2.50/M input, \$10/M output) - **Capability gap:** Bonsai 2 27B retains 98.2% of the performance of Qwen3.8 27B, which itself is roughly on par with GPT-4o-mini for coding, summarization, and structured writing. For complex multi-step reasoning, creative writing, or tasks requiring deep domain knowledge, GPT-4o is still clearly better. But for daily coding, editing, and analysis work? The gap is small enough that many developers won't notice it. - **Cost:** Free, self-hosted. Requires a machine with at least 8 GB of free RAM — runs at 144 tok/s on an M5 Max MacBook Pro, or on a desktop with a mid-range GPU (RTX 3060+). No API costs, no rate limits, no data being sent to a third party. - **Practical setup path:** `brew install incoai/tap/splash && splash serve --model incoai/Qwen3.8-27B-Splash` on an M5 Mac. Or download the GGUF from Hugging Face and run it with LM Studio or Ollama ([Reddit r/LocalLLaMA](#), [Intelligent Living](#)).

Inco Splash inference engine: up to 4x faster than Ollama on Apple Silicon A new open-source inference engine called Inco Splash was released specifically for Apple Silicon (M3+). It claims up to 3x the decode speed of Ollama, 2x faster than MLX, and almost 4x faster when running agent workflows that fan out into sub-agents ([Reddit r/LocalLLaMA](#)). The entire setup: `brew install incoai/tap/splash`. This matters because Apple Silicon MacBooks are the most common developer machines, and running local models on them has been held back by inference speed. If the benchmarks hold, this could significantly accelerate local AI adoption among developers.

Laya: open-source version of the Jev voice assistant Laya, an open-source version of the Jev voice assistant, was released with 332 points and 64 comments on Hacker News as of this morning ([Hacker News](#)). The context: Jev was a commercial voice assistant that gained attention for its natural conversation flow. Laya replicates that capability as an open-source project, meaning developers can self-host a voice assistant without paying per-call API fees.

Cloudflare AI Gateway vs Portkey vs Kong: a \$50K cost gap A comparison of AI gateway tools — services that sit between your application and the AI model API, handling routing, caching, rate limiting (controlling how many requests are sent per second), and failover — found that the choice between Cloudflare AI Gateway, Portkey, and Kong can create a \$50,000 cost difference at scale ([tech-insider.org](#)). The main savings come from caching — if many users ask the same question (e.g. "what's the company policy on X?"), a gateway can serve the cached response instead of hitting the API again. Cloudflare's is cheapest if you already use their CDN; Portkey has the best routing features; Kong is most customizable but most expensive.

Warning on Clore.AI GPU renting A user on r/LocalLLaMA warned about renting GPUs on Clore.AI, a service that lets people rent out their personal GPUs to others for AI inference ([Reddit r/LocalLLaMA](#)). The concern: renters' jobs can include malicious code, and the host has limited visibility into what's being run. If someone uses your GPU to mine crypto or process illegal content, you could be liable. This is the "sharing economy" problem applied to GPU compute.

6. Watchlist: Earnings, Guidance & Movers

Notable movers from today's session (approximate, delayed):

Big gainers: - **SanDisk (SNDK):** +10.99% — no specific news in sources, but the Tom's Hardware weekly note mentions a deepening DRAM crisis ([Tom's Hardware](#)). NAND flash and DRAM prices have been rising, benefiting memory manufacturers. - **Coinbase (COIN):** +11.66%, **Robinhood (HOOD):** +9.12% — no specific news in sources; likely driven by the broader crypto/equities rebound and the Fed rate hike being "priced in." - **Applied Optoelectronics (AAOI):** +7.25%,

Coherent (COHR): +7.22% — these are optical component makers that benefit from AI data center buildout. The theme from today's Cisco article applies directly: demand for optical transceivers connecting GPUs in clusters is surging. - **KLAC (KLA Corp):** +4.74% — semiconductor capital equipment. No specific news, but the broader semi trade was strong. - **Credo (CRDO):** +4.54% — optical networking for AI. Same theme as AAOI and COHR. - **Arm Holdings (ARM):** +4.04% — chip design IP. No specific news. - **Micron (MU):** +3.92% — memory maker, benefiting from DRAM

crisis. - **Astera Labs (ALAB)**: +3.30% — data center connectivity chips.

Losers: - **Redwire (RDW)**: -7.01% — space infrastructure. No specific news in sources. - **Qualcomm (QCOM)**: -5.82% — no specific news in sources. Could be rotation out of mobile/chip stocks into AI infrastructure names. - **Netflix (NFLX)**: -4.67% — no specific news. The Paramount-Warner deal might raise concerns about consolidation among competitors. - **Fluence Energy (FLNC)**: -4.44% — battery storage. Power stocks under pressure as interest rates stay higher for longer. - **IBM**: -3.45% — no specific news. The old-tech rotation (Cisco up, IBM down) may reflect investors picking specific AI-related old-tech names.

No notable news in the sources for: TSM, NEE, RGTI, SOFI, PLTR, FLNC, AMC, TSLA, META, MSFT,

AMZN, NET, VST, IMAX, CNK, CMCSA, DLR, EQIX, SBUX, FXAIX. The moves above are likely sector rotation, not company-specific events.

Intel (INTC): -0.18% — no major move, but the company suspended its bug bounty program that previously paid up to \$100,000 per flaw. A new disclosure program through Intigriti offers no rewards ([Tom's Hardware](#)). For a company trying to rebuild its security reputation, this is a small but negative signal.

Nike (NKE): -2.34% — parted ways with Kylian Mbappé, who signed with Swiss brand On instead ([CNBC](#)). On is a smaller competitor but growing fast in running shoes. Losing Mbappé — one of the world's most marketable athletes — to a rival is a symbolic loss for Nike's soccer business.

7. Analysis: Second-Order Effects

Chain 1: Riyadh airport attack -> oil war-risk premium -> diesel prices -> inflation -> Fed policy

- **FACTS**: Smoke and flames were seen near Riyadh's main airport after an Aramco fuel tank was reportedly hit ([France 24](#)). This follows a period where the US-Iran war's cost reached \$43.6 billion ([The Edwardsville Intelligencer](#)). The Fed raised rates this week with Warsh signaling more could be coming ([CNBC](#)).
- **INFERENCE**: A direct hit on Saudi energy infrastructure near the capital would likely spike the war-risk premium in oil markets (Brent crude). Higher oil prices -> higher diesel prices (as the Detroit Free Press noted) -> higher transportation costs -> broader inflation pressure. The Fed, which targets core PCE (which excludes food and energy), might see second-round effects as higher energy costs seep into core goods through transport and production costs.
- **WEAKEST LINK**: This is LOAD-BEARING on whether the attack actually disrupts Saudi oil production or exports. If the Aramco tank fire is contained quickly and oil flows continue, the price spike could fade within days — similar to the 2019 Abqaiq attack, which moved Brent ~15% in a day but fully retraced within weeks. The attack needs to be sustained or escalate to affect production, not just fuel storage.

- **WHAT WOULD FALSIFY IT**: Brent crude prices remain flat or decline over the next 48 hours. That would suggest markets see this as a contained incident.
- **PRECEDENT**: The 2019 Abqaiq-Khuras attack on Saudi Aramco facilities temporarily cut 5.7 million barrels/day of production. Oil spiked ~15% but returned to pre-attack levels within two weeks as Saudi Arabia quickly restored output.

Chain 2: Hyperscaler capex -> GPU as asset class -> chip financing -> bank exposure -> systemic risk

- **FACTS**: Hyperscaler capex is approaching \$700 billion in 2026, with Nvidia's GPUs being used as collateral for loans ([International Finance](#)). A secondary market for used GPUs is emerging. Banks are writing loans secured by GPU clusters.
- **INFERENCE**: This is creating a new form of asset-backed lending — similar to how aircraft or ships are financed. If GPU prices crash (e.g., from a new competitor, a demand slowdown, or a tech cycle downturn), the collateral backing these loans would lose value, potentially causing margin calls (demands for additional collateral) and loan defaults. Banks could be left holding depreciating GPUs they can't easily sell.

- **WEAKEST LINK:** The assumption that GPU prices are vulnerable to a crash. Currently, demand far exceeds supply, and the lead time to build new fabs is 3-5 years. A crash is possible but not imminent. The more likely scenario is a gradual normalization of prices, not a sharp collapse.
- **WHAT WOULD FALSIFY IT:** If Nvidia's next-generation GPU (Rubin or whatever follows Blackwell) is a commercial disappointment, or if a credible competitor (AMD, custom ASICs from Google/Amazon) captures significant market share, the collateral thesis weakens.

Chain 3: Saudi escalation + Taiwan drills -> Trump-Xi summit pressure -> US-China AI talks -> export controls

- **FACTS:** Xi Jinping is visiting Washington, with AI safety, trade, and Taiwan on the agenda ([Quincy Institute](#)). Taiwan just held live-fire missile and drone drills ([Al Jazeera](#)). Saudi Arabia (a key US ally) is under attack from Houthis backed by Iran (a key Chinese partner).
- **INFERENCE:** Trump has leverage going into the summit — the US is actively engaged in the Middle East, and China needs stability there for its energy imports and Belt and Road investments. Trump could push for China to pressure Iran to restrain the Houthis, in exchange for softer terms on AI export controls or tariffs. Xi needs to show progress on Taiwan's status without appearing weak.
- **WEAKEST LINK:** This assumes Trump wants a deal. He may prefer confrontation, or he may extract concessions and then renege — a known pattern from his first term.
- **WHAT WOULD FALSIFY IT:** If Trump announces new tariffs or export restrictions on China during or

immediately after the summit, the "deal" thesis is wrong.

Chain 4: Compact local models (Bonsai 2) -> enterprise AI adoption accelerates -> less reliance on cloud API providers -> lower revenue for inference providers

- **FACTS:** Bonsai 2 27B achieves 98.2% of Qwen3.8 27B performance in 5.9 GB, runs at 144 tok/s on a laptop, and is Apache 2.0 licensed ([MarkTechPost](#)). At Dreamforce, business leaders said last year's models are already good enough for most enterprise needs ([CNBC](#)).
- **INFERENCE:** The combination of "good enough" local models and the ability to run them on cheap hardware could accelerate enterprise AI adoption in a way that reduces reliance on cloud API providers (OpenAI, Anthropic, Google). Companies worried about data privacy, latency, and API costs can now run capable models entirely on-premise. This is bearish for inference-as-a-service providers but bullish for hardware makers (Apple Silicon, consumer GPUs) and local inference software.
- **WEAKEST LINK:** The assumption that 98.2% performance retention holds across real-world enterprise tasks. Benchmarks are one thing; enterprise workflows often require consistent reliability on edge cases. Also, many enterprises lack the in-house talent to deploy and maintain local models.
- **WHAT WOULD FALSIFY IT:** If cloud API revenue continues to grow at 50%+ year-over-year despite the availability of capable local models, that would suggest enterprises still prefer managed services.

8. Building Your Knowledge

1. How "GPU collateral" works and why it matters for investors Nvidia's chips are increasingly used as collateral for loans — a bank lends a hyperscaler money, the hyperscaler pledges its GPU cluster as security. If the borrower defaults, the bank takes the GPUs. This only works if there's a liquid secondary market for the chips (buyers ready to purchase used GPUs). That secondary market is now emerging, which is a structural

confirmation that AI demand is real and long-lived. It also means GPU prices have a "floor" — if new GPU prices fall, used ones will be absorbed by secondary buyers. This is very different from past tech cycles where server hardware had near-zero resale value.

2. The difference between "hybrid warfare" and conventional war When France, Poland, and other European nations warn about Russian "hybrid" attacks,

they mean actions designed to destabilize without triggering a full military response ([NPR](#)). Think: cyberattacks on power grids, sabotage of undersea internet cables, disinformation campaigns to swing elections, arming proxy forces, assassinations of defectors abroad. The key strategic logic: NATO's Article 5 (an attack on one is an attack on all) only applies to conventional military attacks. By staying below that threshold, Russia can damage NATO members without triggering a collective response. This is why the problem is so hard to solve — the response options (diplomatic protests, sanctions, counter-cyber operations) are weaker than the attacks themselves.

3. Why "guidance matters more than the past" is even more true during rate hikes When the Fed raises rates, investors obsess over what comes next — not what just happened. That's why Warsh's three words ("substantially restrictive") mattered more than the rate hike itself ([CNBC](#)). The market prices in expectations. A rate hike that's expected barely moves markets. But a hint about the future path can move everything. This is a general principle: in investing, the price of an asset already reflects everything the market knows about the past. The only thing that can move the price is new information about the future.

4. How model compression works (ternary quantization) Bonsai 2 27B uses ternary quantization — it stores each weight (a number that determines how

the model processes input) as one of three values: -1, 0, or +1. Standard models store weights as 16-bit numbers (65,536 possible values). Ternary uses roughly 1.58 bits per weight. That's a 10x compression. The trade-off: less precision per calculation. The surprise: modern research shows that if you train the model with this constraint in mind (rather than compressing an already-trained model), the loss of quality is surprisingly small. This is an active area of research with huge practical implications — it's the difference between needing a \$30,000 GPU cluster and a \$4,000 laptop.

5. The EU's "associate member" concept and what it means for geopolitics The EU proposing Canada as its first associate member is a new category being created ([France 24](#)). An associate member would get some economic integration (access to the single market, participation in some EU programs) without full membership (no voting rights, no obligation to adopt EU laws or the euro). This is a strategic response to US "America First" policy: the EU is building a broader coalition of like-minded countries that share its regulatory and values framework. The UK, Switzerland, and Norway already have various forms of associated status. Canada would be the first non-European country. If successful, it could become a model for other countries seeking closer ties without full membership — potentially including post-Brexit UK relations deepening further.

9. Foundations & Terms

Concept Spotlight: Asset-Backed Lending (Markets/Investing)

What it is: Asset-backed lending is when a borrower gets a loan by pledging physical assets as collateral (security for the lender if the borrower fails to repay). The lender can seize and sell the assets if the borrower defaults. Think of it like a mortgage, but instead of a house, the collateral could be aircraft, shipping containers, warehouses, or — increasingly — GPU clusters.

Why it exists: Banks are more willing to lend when there's collateral because it reduces their risk. This allows companies to borrow at lower interest rates than they could with unsecured loans. It also allows

companies to raise money against assets they already own — turning idle equipment into cash.

A concrete worked example: A startup called "Neocloud" buys \$500 million worth of Nvidia H100 GPUs. It then goes to a bank and says: "We need \$300 million to build a data center. We'll pledge the GPUs as collateral." The bank values the GPUs at \$400 million (a 20% haircut, or discount, because used GPUs are worth less than new). The bank lends \$300 million at 8% interest. If Neocloud defaults, the bank seizes the GPUs and sells them on the secondary market.

How it shows up in today's news: The International Finance article about Nvidia's chips becoming "Wall Street's newest asset class" is describing exactly this phenomenon. Banks are now treating GPUs like aircraft

or real estate — they're lending against them, issuing asset-backed securities, and creating lease-back deals ([International Finance](#)).

Rough numbers worth remembering: The haircut on GPU collateral is typically 20-30% — meaning a bank will lend you 70-80% of the GPUs' estimated resale value. GPU resale values have held up unusually well: a used H100 still sells for 60-70% of its original price after a year, compared to 20-30% for a typical server CPU.

Caution: This only works as long as there's a liquid secondary market for GPUs. If AI demand collapses or a new generation makes old GPUs obsolete, the collateral could become nearly worthless. Banks learned this lesson the hard way with commercial real estate in 2008.

Concept Spotlight: Hybrid Warfare (World Affairs)

What it is: Hybrid warfare is the use of a mix of military and non-military tools to achieve strategic goals while staying below the threshold of conventional war. The tools include cyberattacks, disinformation, economic coercion, sabotage of infrastructure, proxy forces, election interference, and energy blackmail.

Why it exists: For a weaker power facing a stronger military alliance (like Russia facing NATO), hybrid warfare offers a way to inflict damage without triggering the alliance's collective defense clause. If Russia launches a conventional attack on Poland, NATO's Article 5 would require all members to respond. But if Russia launches a cyberattack on Poland's power grid and funds anti-government protests through social media, it's harder to prove and harder to respond to proportionally.

A concrete worked example: A Russian intelligence unit hacks the control systems of an undersea internet cable connecting France and the UK. They don't cut it — they introduce a fault that causes intermittent outages. France's internet slows down. French companies lose business. French citizens blame the French government for poor infrastructure. Russia denies involvement. There's no clear evidence, no smoking gun, and France can't invoke Article 5 because it's not a military attack.

How it shows up in today's news: France, Poland, and other European nations explicitly warned of escalating Russian hybrid attacks — including cyberattacks, sabotage, and disinformation ([NPR](#)). Putin

separately accused Ukraine of sabotaging peace talks ([Kyiv Post](#)) — itself a classic hybrid tactic: accusing your opponent of what you're doing.

Rough benchmarks worth remembering: The EU and NATO estimate that Russia spends roughly \$1.5-2 billion per year on hybrid operations. For comparison, the US military budget is ~\$850 billion. Hybrid warfare offers an enormous "return on investment" — a few million dollars in cyber operations can cause billions in damage.

Common misconception: Hybrid warfare is not "new." The Soviet Union ran extensive disinformation and destabilization campaigns during the Cold War. What's new is the scale and speed enabled by the internet, social media, and cheap cyber tools.

Other Terms from Today's News:

- **Collateral:** An asset pledged by a borrower to a lender as security for a loan. If the borrower defaults (fails to repay), the lender can seize and sell the collateral. In today's news, Nvidia GPUs are being used as collateral for data center loans.
 - **Secondary market:** A market where existing assets are bought and sold among investors, rather than directly from the original issuer. For GPUs, the secondary market means companies can buy used H100s instead of waiting months for new ones. A deep secondary market makes lenders more willing to accept GPUs as collateral.
 - **Associate member (EU):** A proposed new category of relationship with the European Union that offers economic integration (access to the single market) without full membership (no voting rights, no obligation to adopt EU laws or the euro). Canada would be the first non-European country proposed for this status.
 - **Fire sale:** The rapid sale of assets at heavily discounted prices, usually because the seller needs cash urgently. If GPU prices crash and lenders seize collateral, they may have to conduct fire sales, pushing prices down further. This is the risk in asset-backed lending.
 - **Apropos of something:** A phrase meaning "relevant to the topic at hand" or "speaking of which." Not a finance or geopolitics term, but useful for connecting ideas.
-

10. Bottom Line

The world this morning is defined by a dangerous convergence of conflict and markets. The Houthis strike near Riyadh airport — an Aramco fuel tank on fire, flight disruptions, the capital under alert — is the most significant escalation in the Middle East since the US-Iran war began. It raises the risk of a broader regional war that would drive oil prices sharply higher, complicating the Fed's inflation fight just as Warsh signals more rate hikes ahead. Meanwhile, Taiwan's live-fire drills ahead of the Trump-Xi summit emphasize that the US-China relationship remains the world's most consequential flashpoint.

The AI buildout continues at a frenetic pace, with hyperscaler capex approaching \$700 billion and Nvidia's GPUs becoming a new asset class for Wall Street. But the tension in the air is real: the Riyadh attack could boost oil prices and interest rates, both of which are

headwinds for the long-duration, high-valuation growth stocks that dominate the AI trade. The market is betting on an AI rebound (options activity was unusually bullish on Thursday), but the geopolitical risks are mounting faster than the earnings.

What changed since yesterday: The Riyadh airport incident is new and escalates the Middle East conflict significantly. The Pentagon's potential 40,000-troop drawdown from Europe became more concrete. Trump's ban on three media outlets from the White House is a new domestic political development. Bonsai 2 27B was released — a genuinely important model for local AI deployment. The Paramount-Warner Bros. \$110 billion merger story broke. And the AI safety debate intensified with Gemini hacking companies and a public letter demanding independent safety evaluators for Anthropic and OpenAI.

11. Sources

World & Geopolitics

- France 24 — Smoke, flames seen near Riyadh airport
- CNBC — Flames reported near Saudi capital airport
- The Edwardsville Intelligencer — Iran war's cost grows to \$43.6 billion
- Al Jazeera — Taiwan drills ahead of Trump-Xi talks
- Quincy Institute — Xi goes to Washington
- The Next Web — Last Trump-Xi AI talks produced nothing
- The Eastern Herald — Pentagon weighs pulling 40,000 troops from Europe
- NPR — France, Poland warn of Russian hybrid attacks
- Kyiv Post — Putin accuses Ukraine of sabotaging peace talks
- France 24 — Putin says Europe is preparing for war
- CNBC — Trump bans MS NOW, CNN from White House
- France 24 — Macron hosts Canadian PM Carney
- CNBC — Auto industry urges Trump to bar Chinese automakers

Markets, Money & Deals

- CNBC — Warsh's three words on rate hikes
- MarketWatch — Warsh's Fed serious about inflation
- CNBC — Japan rate hike flips script
- CNBC — Buffett steps down as Berkshire chairman
- UNN/Reuters — Paramount may buy Warner Bros. for \$110 billion
- ABC News — FCC grants Paramount request for Gulf funds
- WSJ — Brookfield acquires Australia's Reliance for \$2.9B
- CNBC — Consumer sentiment in the dumps

- CNBC — AI trade rebound sparks bullish options

AI & Infrastructure

- International Finance — How Nvidia turned chips into Wall Street's newest asset class
- Business Insider — AI capex boom masks S&P 500 earnings truth
- SiliconANGLE — Nscale files for IPO
- MarketWatch — Cisco and retro tech back for AI buildout
- Modern Diplomacy — AI's power bill becomes geopolitical

Model & Research Watch

- MarkTechPost — PrismML releases Bonsai 2 27B
- MarkTechPost — Alibaba Qwen3.8-Omni-Flash
- SCMP via Hacker News — Alibaba open-sources medical AI model
- BBC — Google's Gemini hacked three companies
- Tech Times — Aging biology benchmark finds compact AI beats frontier

What Builders Are Actually Using

- Reddit r/LocalLLaMA — Qwen3.8-27B at 144 tok/s on M5 Max MacBook Pro
- Intelligent Living — Bonsai 2 27B: 98% of frontier model in 5.9 GB
- Hacker News — Laya open source version of Jev
- tech-insider.org — Cloudflare AI Gateway vs Portkey vs Kong
- Reddit r/LocalLLaMA — Warning about Clore.AI GPU renting

Quick Check

1. Bonsai 2 27B is a compressed model that retains 98.2% of the performance of which larger model, and what makes it notable for deployment?

2. The smoke and flames reported near Riyadh airport and an Aramco fuel tank on fire are linked to escalation

with which group, and what is the primary risk to global markets if this escalates further?

3. If a bank lends \$300 million to a data center company secured by Nvidia GPUs worth \$400 million on the secondary market, and the company defaults, what happens to the GPUs and what is the risk to the bank?

Answers: 1. Bonsai 2 27B retains 98.2% of Qwen3.8 27B performance. It is notable because it fits in 5.9 GB (using ternary quantization to ~1.58 bits per weight), runs on consumer hardware like a MacBook Pro at 144 tok/s, and is Apache 2.0 licensed — making frontier-

capable AI free and local. 2. The attack is linked to Yemen's Houthi rebels, who are backed by Iran. The primary risk to global markets is a spike in oil prices if Saudi oil production or exports are disrupted, which would feed into inflation and potentially force the Fed to raise rates further. 3. The bank seizes the GPUs and sells them on the secondary market to recover its money. The risk is that if GPU prices have crashed (e.g., from a tech downturn or new generation making them obsolete), the bank may only recover a fraction of the \$300 million loan — a "fire sale" scenario that can cause losses and tighten lending to the AI sector.

Generated 2026-09-19 08:01 PDT · 130 sources · model: deepseek-ai/DeepSeek-V4-Flash via <https://api.gmi-serving.com/v1>